

# Overview of Machine Learning (II)

## Regression and Gradient Descent

*Concepts, Supervised Learning, Unsupervised Learning,  
Regression, Classification*

# The Essential Elements of (most) ML

To use most ML methods, we will need to conceptualize our problem into the following:

## Features (inputs, descriptors)

$$\{\mathbf{x}_i\} \leftrightarrow \mathbf{X}$$

$\mathbf{x}_i$  *feature vector of sample i*

- a numerical description of (ideally) characteristics that distinguish one sample from another
- may (or may not) have direct implications on the modeling outputs

## Labels (outputs)

$$\{y_i\} \leftrightarrow \mathbf{y}; \{\mathbf{y}_i\} \leftrightarrow \mathbf{Y}$$

$y_i$  OR  $\mathbf{y}_i$  *scalar or vector  
label of sample i*

- also a numerical (integer or real) description of sample i
- usually reserved for some special quantity or property of interest

## Labeled Data

$$\{(\mathbf{x}, \mathbf{y})_i\}$$

*a set of tuples where  
features and labels  
are known*

## Unlabeled Data

$$\{\mathbf{x}_i\}$$

*labels are not  
necessarily known or  
provided with features*

## Model

*a function that operates on features*

$$f(\mathbf{x}) \text{ or } \hat{\mathbf{y}}$$

- often defines a mapping from feature space to label space

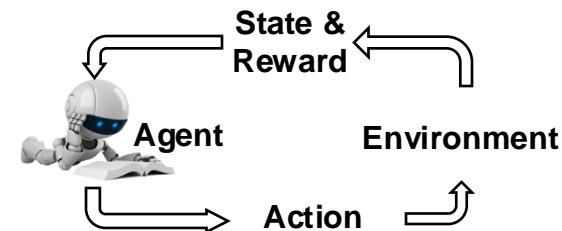
## Predictions

*the function output or  
predicted labels*

$$\hat{\mathbf{y}} = f(\mathbf{x}) \text{ or } \hat{\mathbf{y}} = \hat{\mathbf{y}}$$

# Classes of Machine Learning

Machine learning is deployed in three main modes:



## Supervised Learning

- In ***supervised learning***, we aim to create a model that can predict  $y$  as a function of  $x$ .
- The optimization/learning of our model is ***supervised*** because the algorithm will exploit knowledge of labels over the dataset

Supervised learning can be used for either

- ***Regression*** – predict a *continuous* label. This is likely to be true for QSPR problems in physical science.  
e.g., conductivity, melting point, band gaps
- ***Classification*** – predict *categorical* labels or class membership. This can be useful for characterizing discrete outcomes  
e.g., (in)soluble, (un)synthesizable, (in)activity, hazardous

## Unsupervised Learning

- In ***unsupervised learning***, we aim to create a model that identifies patterns in  $x$ .
- The optimization/learning of our model is ***unsupervised*** because the algorithm will not exploit knowledge of labels over the dataset

Unsupervised learning is usually used for

- ***Clustering*** – partition features into a set of different classes/groups, which is the  $y$ .  
e.g., chemical classes
- ***Signal processing*** – Uncover the underlying signal within a set of features. This is often a part of representation learning.  
e.g., protein folding pathways
- ***Generating*** – create a model distribution over  $x$  such that we can generate new samples

## Reinforcement Learning

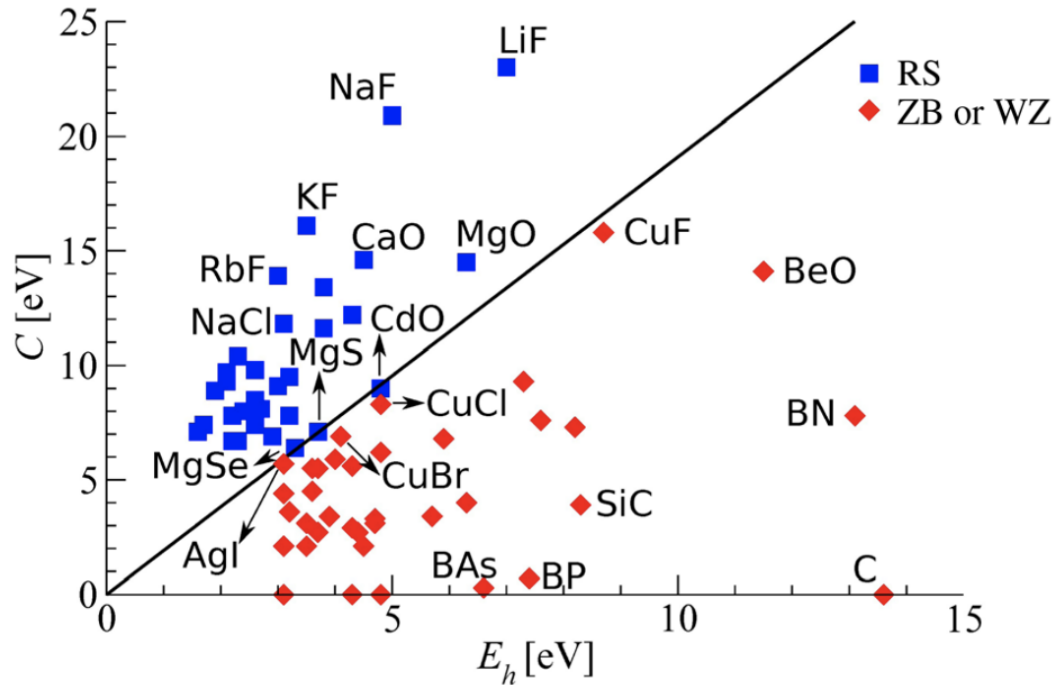
- In reinforcement learning, an “agent” learns how to interact with its environment based on feedback via cumulative rewards/penalties
- Many things that people think are reinforcement learning are probably not reinforcement learning
- Usually about planning and scheduling

e.g., automated process synthesis, process control

- In ***semi-supervised learning***, we want a model that can predict  $y$  as a function of  $x$ , just as in supervised learning
- ***Both labeled and unlabeled data*** are used in modes like *co-training*, *pseudo-labeling*, and *label propagation*

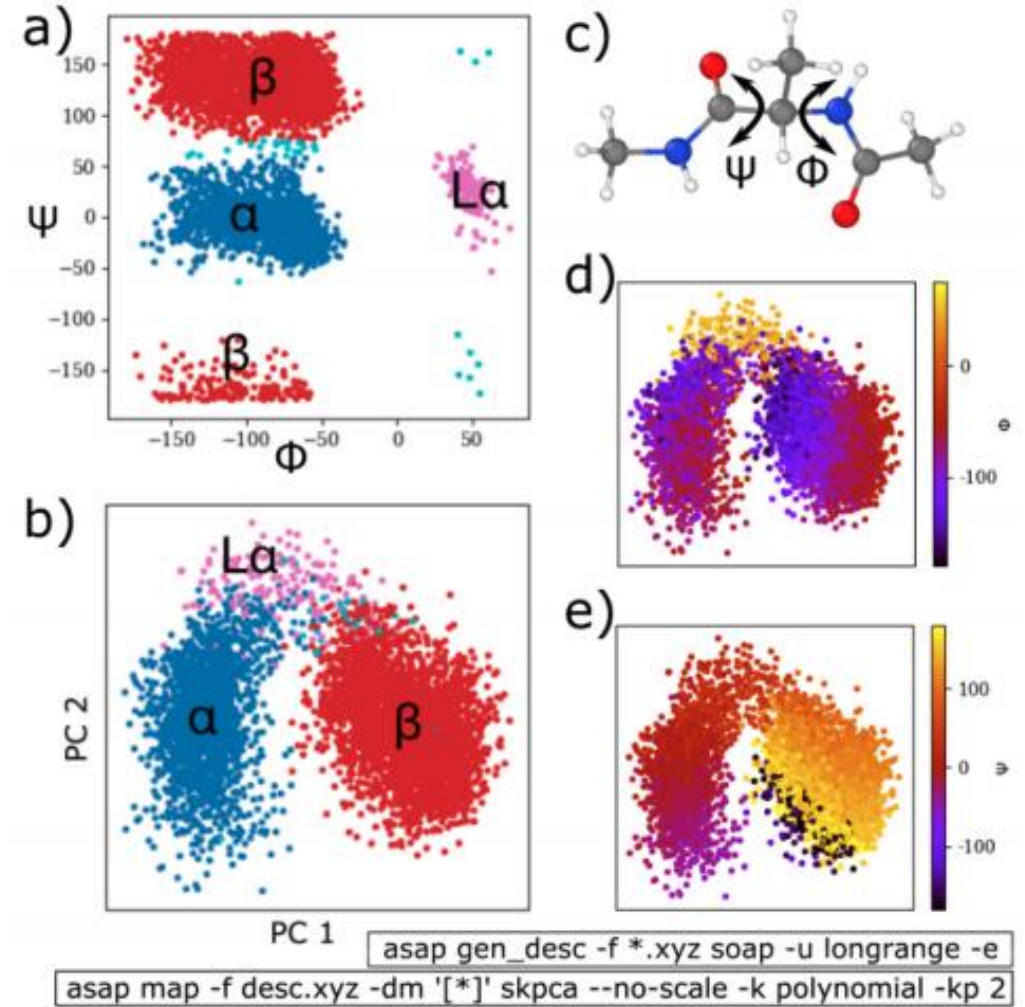
- In ***self-supervised learning***, we eventually want a model that can predict  $y$  as a function of  $x$ ,
- ***Only unlabeled data*** are used during training; one form is *contrastive learning*

# Example: supervised & unsupervised Learning



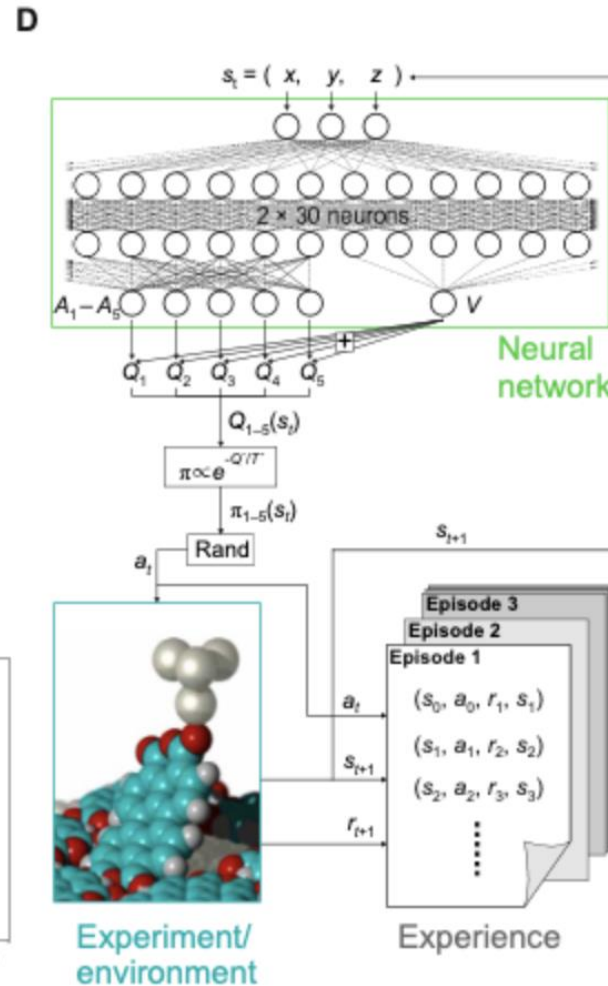
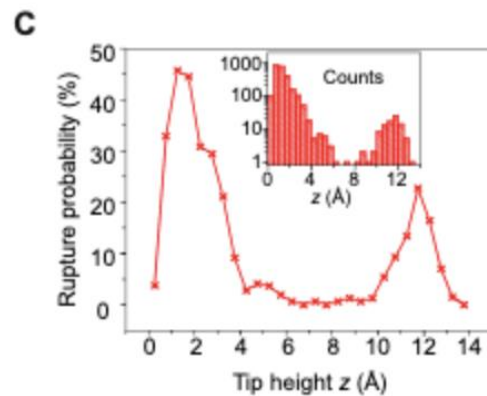
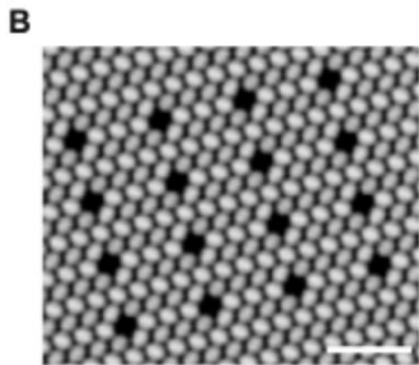
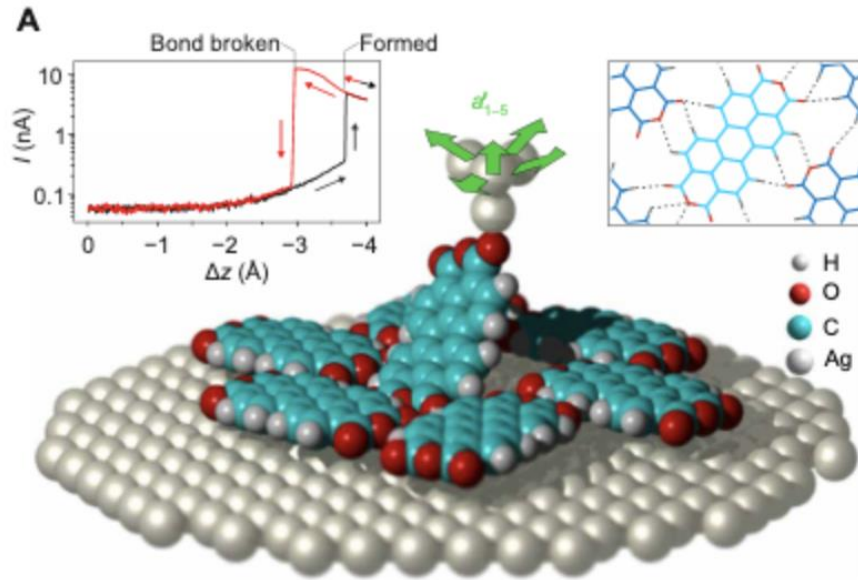
Ghiringhelli et al. *PRL* 114 (2015)

Two simple "descriptors" (related to nearest neighbor distance and dielectric constant) define a function that serves as a decision boundary that distinguishes between rocksalt and zinc blende or wurtzite crystal structures.



This illustrates a typical Ramachandran plot of alanine dipeptide by comparison to a unsupervised learning over molecular configurations.

# Example: Reinforcement Learning



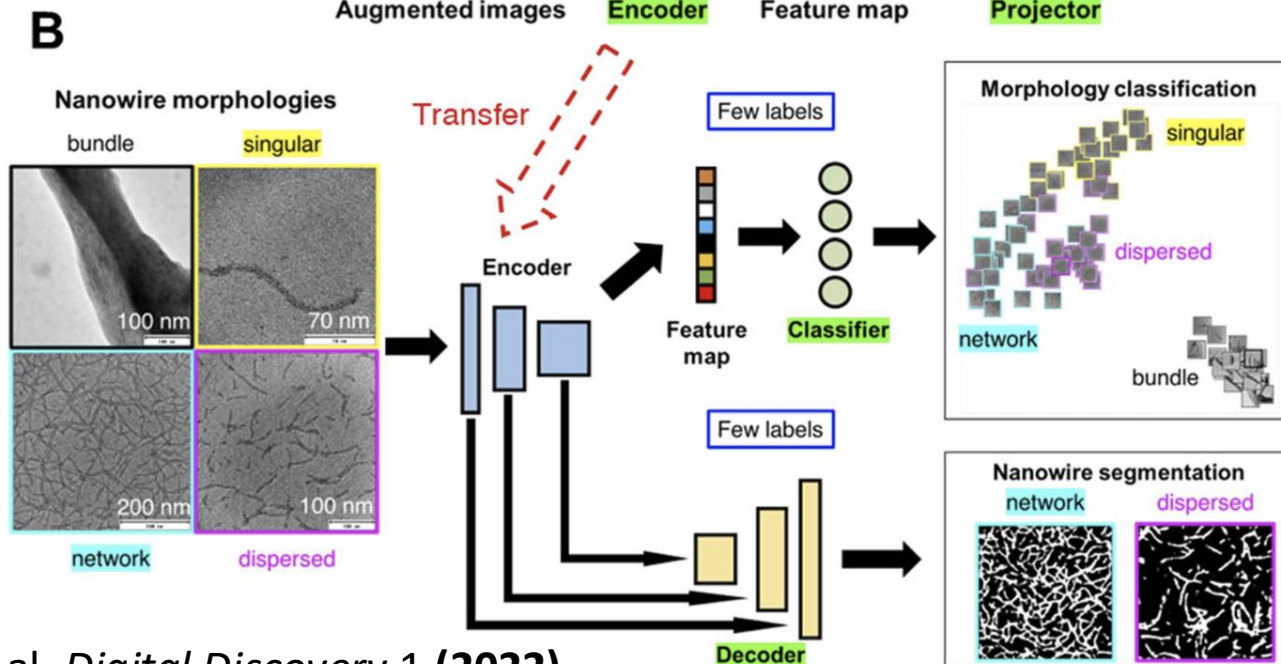
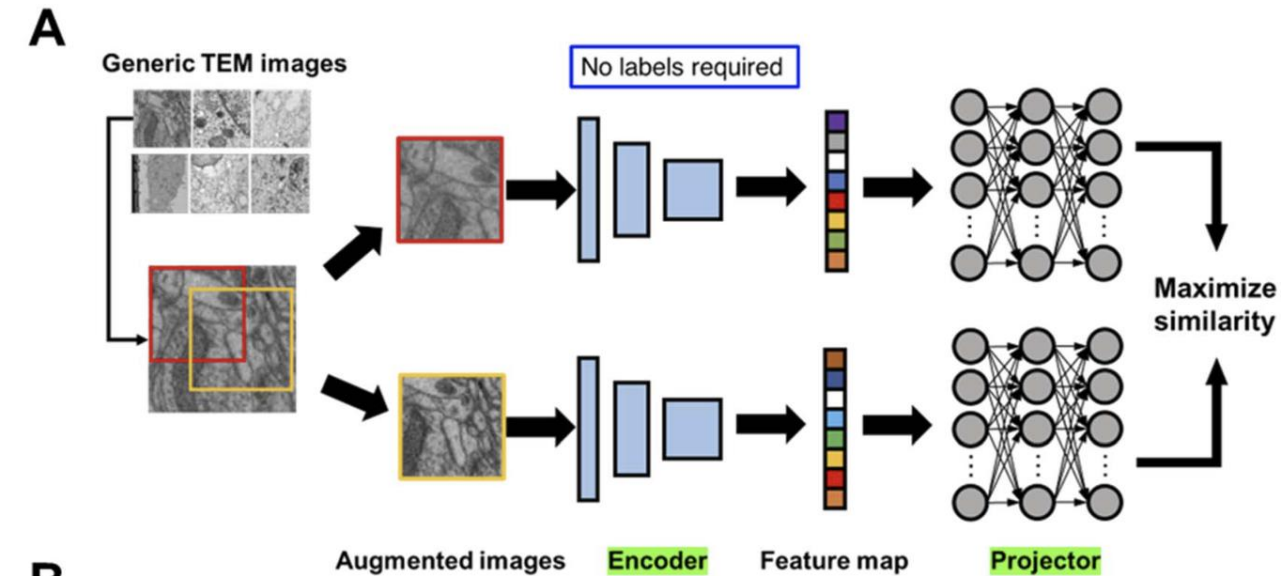
## Guiding nanofabrication with single-molecule manipulation

- Scanning probe microscope can remove molecules from supramolecular assembly, but apparently this is non-trivial manual task
- Reinforcement learning is used to develop a protocol to move the tip in a manner that enables effective molecule lifting

We design the reward system as follows: If the environment transitions to a nonterminal state, we assign a default reward of  $r_{t+1} = 0.01$  (see Materials and Methods for a discussion). If transitioning into a state in which the SPM tip loses contact with the molecule, the agent is penalized with  $r_{t+1} = -1$ , and the current episode stops. Last, if transitioning into a state where the molecule has been lifted successfully, we assign a reward of  $r_{t+1} = +1$ , and the episode also stops. After each failed episode, the molecule, by virtue of



# Example: Self- + Semi- Supervised Learning



## “Fancy” ML workflow for microscopy segmentation & classification

- Combines many “advanced” architecture concepts with semi-supervised approach in a “transfer learning” paradigm.
- Self-supervised learning component comes from matching an image to itself! (they must come from the same class... probably?)
- Overall goal is efficient labeling of TEM/data efficiency

# **Regression**

# **Gradient Descent**

# The basic problem of curve-fitting

Machine learning Regression is often characterized as “fancy curve-fitting”; to understand the (un)fairness of that statement, we will first describe good ole regular curve-fitting

## Linear Least-Squares Regression

$$f(x) = \theta_0 + \theta_1 x$$

Objective:

$$\begin{aligned} \min_{\theta} \mathcal{E}(f) &= \min_{\theta} \sum_{k=1}^n |e_k|^2 \\ &= \min_{\theta} \sum_{k=1}^n (\theta_0 + \theta_1 x - y_k)^2 \end{aligned}$$

Given  $\{(\mathbf{x}_i, y_i)\}$  produce “optimal”  $f$   $\hat{y} = f(\mathbf{x}, \theta)$   
that minimizes some error metric (“loss”)  $\mathcal{E}(\{y_k, \hat{y}_k\})$

$$e_k = \hat{y}_k - y_k$$

### Some possible loss functions

$$\begin{aligned} \mathcal{E}_{\infty}(f) &= \max_k |e_k| & \mathcal{E}_2(f) &= \sqrt{\frac{1}{n} \sum_{k=1}^n |e_k|^2} \\ \mathcal{E}_1(f) &= \sqrt{\frac{1}{n} \sum_{k=1}^n |e_k|} & \mathcal{E}_p(f) &= \sqrt[p]{\frac{1}{n} \sum_{k=1}^n |e_k|^p} \end{aligned}$$

*note that the “optimal”  $f$  depends on the loss function*



# The basic problem of curve-fitting

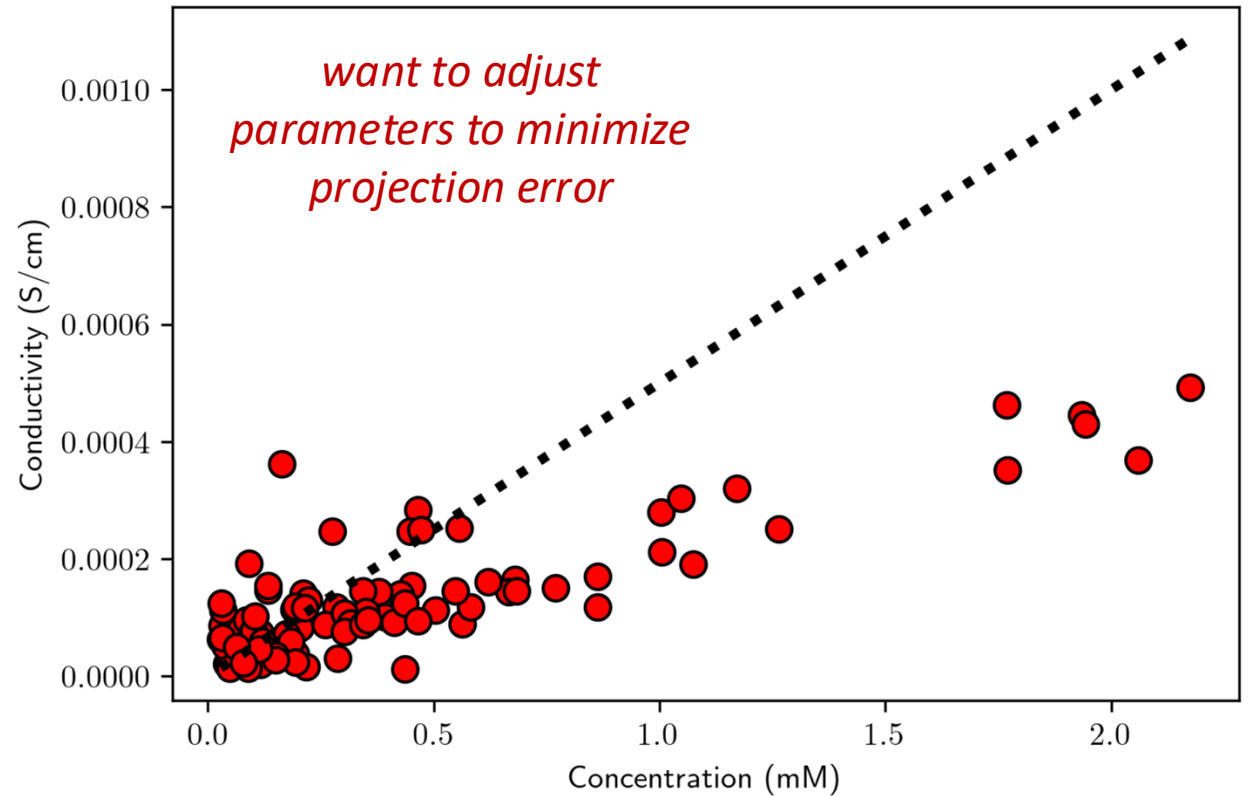
Machine learning Regression is often characterized as “fancy curve-fitting”; to understand the (un)fairness of that statement, we will first describe good ole regular curve-fitting

## Linear Least-Squares Regression

$$f(x) = \theta_0 + \theta_1 x$$

**Objective:**

$$\begin{aligned} \min_{\theta} \mathcal{E}(f) &= \min_{\theta} \sum_{k=1}^n |e_k|^2 \\ &= \min_{\theta} \sum_{k=1}^n (\theta_0 + \theta_1 x_k - y_k)^2 \end{aligned}$$



We will explore using gradient descent for this problem, but it can be approached easily/exactly. How?

# The basic problem of curve-fitting

Machine learning Regression is often characterized as “fancy curve-fitting”; to understand the (un)fairness of that statement, we will first describe good ole regular curve-fitting

## Linear Least-Squares Regression

$$f(x) = \theta_0 + \theta_1 x$$

**Objective:**

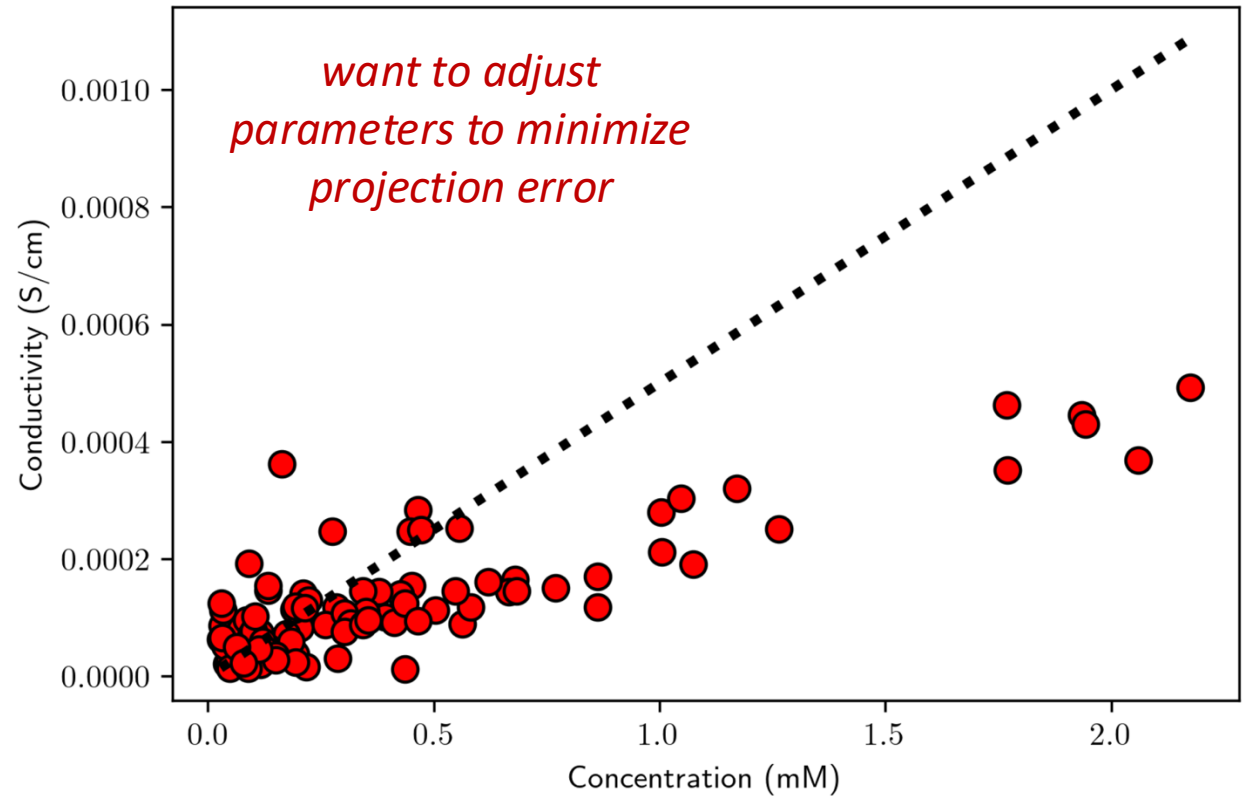
$$\min_{\theta} \mathcal{E}(f) = \min_{\theta} \sum_{k=1}^n |e_k|^2$$

*set derivatives to zero and solve*

$$= \min_{\theta} \sum_{k=1}^n (\theta_0 + \theta_1 x_k - y_k)^2$$

$$\begin{pmatrix} n & \sum_{k=1}^n x_k \\ \sum_{k=1}^n x_k & \sum_{k=1}^n x_k^2 \end{pmatrix} \begin{pmatrix} \theta_0 \\ \theta_1 \end{pmatrix} = \begin{pmatrix} \sum_{k=1}^n y_k \\ \sum_{k=1}^n x_k y_k \end{pmatrix}$$

*this only possible if our function is linear in all its parameters*



# Non-linear Regression

Machine learning Regression is often characterized as “fancy curve-fitting”; to understand the (un)fairness of that statement, we will first describe good ole regular curve-fitting

## Non-linear regression

$f(x, \boldsymbol{\theta})$  *now just some general function, which is not necessarily linear in its parameters*

e.g.,

$$f(x, \boldsymbol{\theta}) = \theta_0 \cos(\theta_1 x + \theta_2) + \theta_3$$

If we consider a loss related to  $l_2$  - norm, then

$$\mathcal{E}(\boldsymbol{\theta}) = \sum_{k=1}^n (f(x_k, \boldsymbol{\theta}) - y_k)^2; \quad \frac{\partial \mathcal{E}}{\partial \theta_i} = 0 \quad \forall i$$

$$\Rightarrow \sum_{k=1}^n (f(x_k, \boldsymbol{\theta}) - y_k) \frac{\partial f}{\partial \theta_i} = 0 \quad \forall i$$

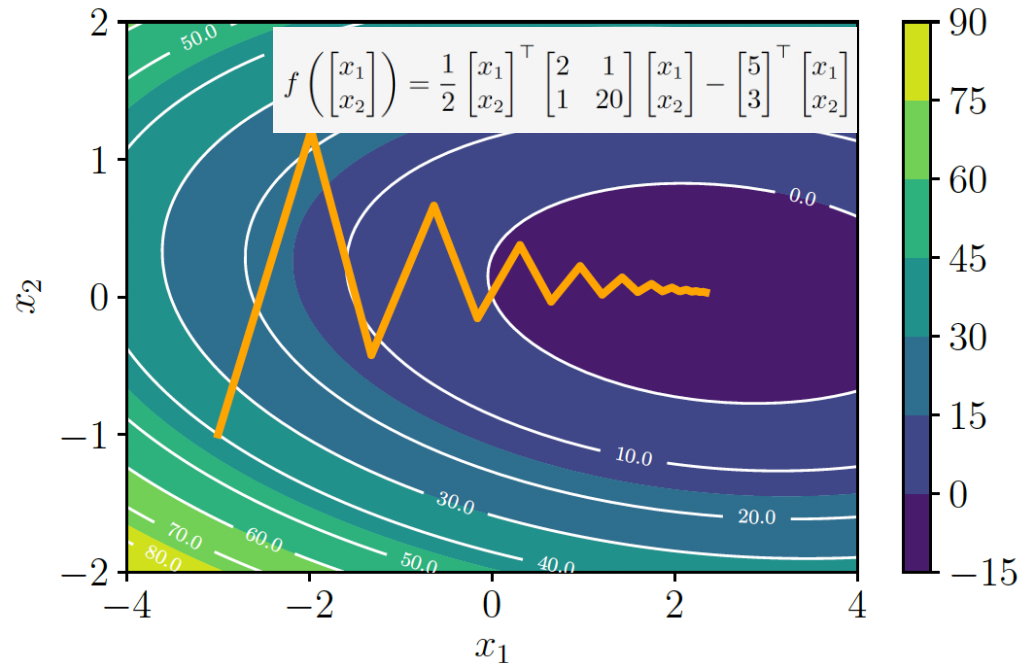
$$\sum_{k=1}^n e_k \frac{\partial f}{\partial \theta_i} = 0 \quad \forall i$$

# Gradient Descent

Essential task:  $f : \mathbb{R}^n \rightarrow \mathbb{R}, \mathbf{x} \mapsto f(\mathbf{x})$   $\min_{\mathbf{x}} f(\mathbf{x})$

## Gradient Descent

$$\mathbf{x}_{i+1} = \mathbf{x}_i - \gamma_i [\nabla f(\mathbf{x}_i)]^T$$



- *very simple method*
- *depends on (adaptive) stepsize*
- *slowly convergent to closest minima*

## with momentum

$$\mathbf{x}_{i+1} = \mathbf{x}_i - \gamma_i [\nabla f(\mathbf{x}_i)]^T + \alpha \Delta \mathbf{x}_i$$

$$\Delta \mathbf{x}_i = \alpha \Delta \mathbf{x}_{i-1} - \gamma_{i-1} [\nabla f(\mathbf{x}_{i-1})]^T$$

- *uses “memory” to reduce jitter*

## stochastic

$$f(\mathbf{x}) = \sum_{k=1}^N f_k(\mathbf{x})$$

$$\mathbf{x}_{i+1} = \mathbf{x}_i - \gamma_i \sum_{k \in \mathbb{N}: k \leq N} [\nabla f_k(\mathbf{x}_i)]^T$$

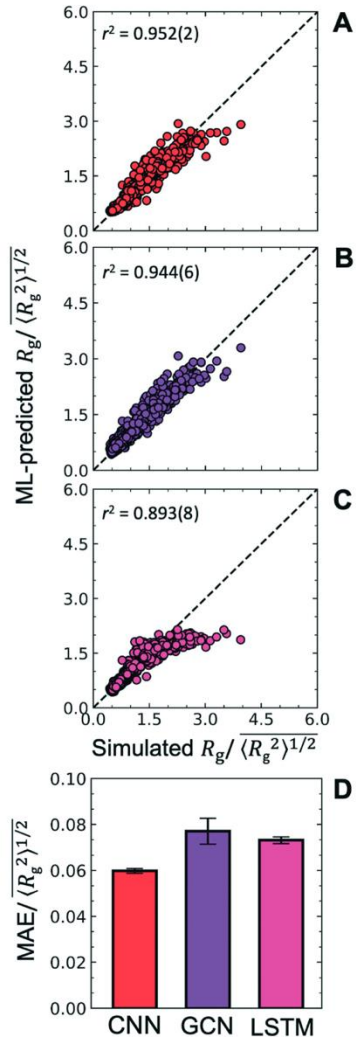
- *useful for large  $N$ , which may not be atypical in machine learning applications*

# **Notebook Exercise**



# Activity: Premise and Objective

We will understand basic essence of parameter optimization using the example of linear regression and gradient descent; our problem of study relates to polymer physics



MSDE

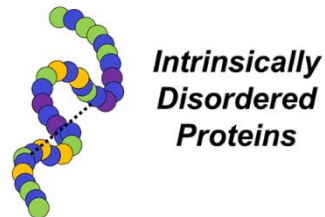
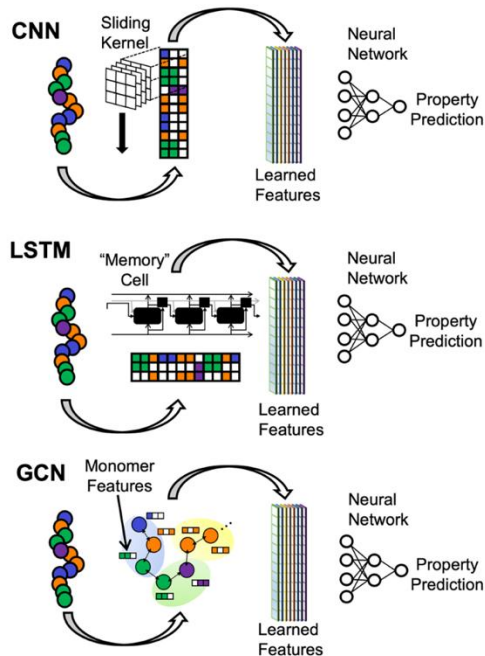
PAPER



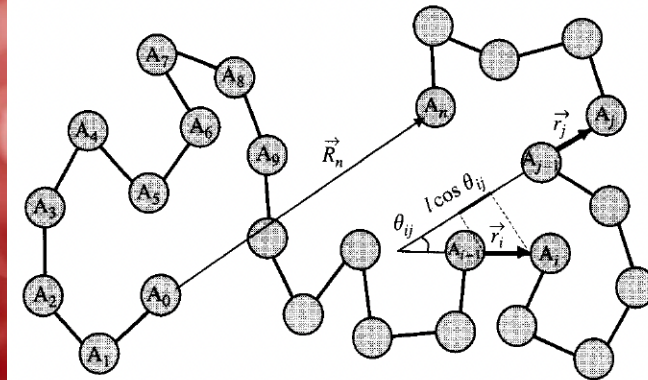
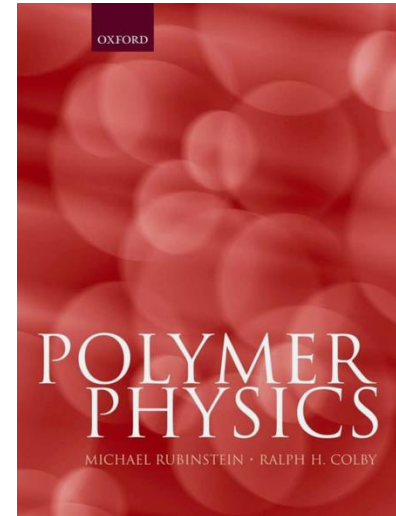
Cite this: *Mol. Syst. Des. Eng.*, 2022, 7, 661

**Featurization strategies for polymer sequence or composition design by machine learning†**

Roshan A. Patel, Carlos H. Borca and Michael A. Webb \*



[View Article Online](#)  
[View Journal](#) | [View Issue](#)



$$\begin{aligned}\langle R^2 \rangle &\equiv \langle \vec{R}_n^2 \rangle = \langle \vec{R}_n \cdot \vec{R}_n \rangle = \left\langle \left( \sum_{i=1}^n \vec{r}_i \right) \cdot \left( \sum_{j=1}^n \vec{r}_j \right) \right\rangle \\ &= \sum_{i=1}^n \sum_{j=1}^n \langle \vec{r}_i \cdot \vec{r}_j \rangle = l^2 \sum_{i=1}^n \sum_{j=1}^n \langle \cos \theta_{ij} \rangle \\ \langle R^2 \rangle &= n l^2.\end{aligned}$$

# Initialization and data inspection

## Plotting the data

```
# Modules used by Prof. Webb
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import urllib.request
import random
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

```
url_for_labels = "https://raw.githubusercontent.com/webbtheosim/featurization/main/Dataset_A/labels.csv"
url_for_sequences = "https://raw.githubusercontent.com/webbtheosim/featurization/main/Dataset_A/sequences.txt"
idpdata = pd.read_csv(
    url_for_labels
)

y = idpdata['ROG (A)'].to_numpy()/10. # these are now labels
seqs = [line.strip().split() for line in urllib.request.urlopen(url_for_sequences)]
X = np.array([len(seq) for seq in seqs])**0.5 # these are features

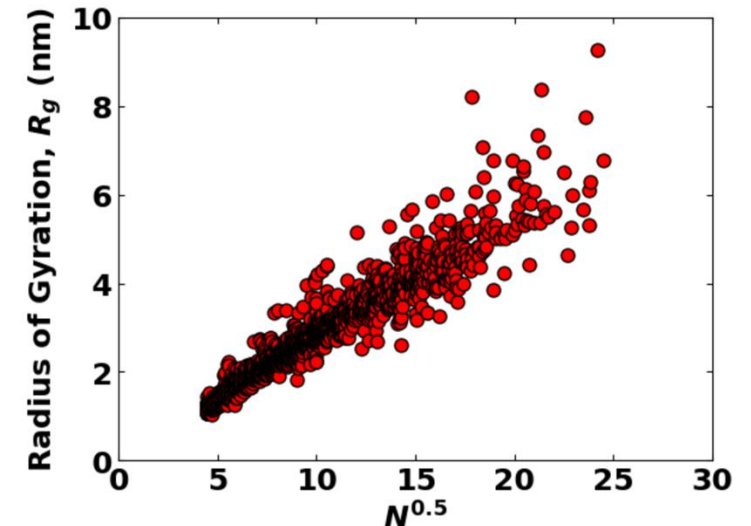
idpdata.head()
```

	ROG (A)	CV (J/K)	TAUS (fs)
0	11.725914	0.444604	36585.162
1	11.912079	0.370302	40234.011
2	11.375047	0.399939	38123.675
3	11.457038	0.407542	34174.561
4	11.509964	0.449730	34279.740

```
# global specifications on plots
plt.rcParams.update({'font.size': 18,
                    'font.weight': 'bold',
                    'axes.labelweight': 'bold'})

def plot_raw_data(x,y):
    plt.plot(x, y, marker='o', linestyle="", markersize=8, \
             color='r', markeredgecolor='k')
    plt.ylabel("Radius of Gyration,  $R_g$  (nm)")
    plt.xlabel(" $N^{\{0.5\}}$ ")
    plt.xlim(0,30)
    plt.ylim(0,10)
    ax = plt.gca()
    ax.tick_params(direction='in')
    ax.yaxis.set_ticks_position('both')
    ax.xaxis.set_ticks_position('both')
    return ax

ax = plot_raw_data(X,y)
```



# Human Hypothesis to the Data

$$R_g = \theta_0 + \theta_1 N^{0.5}$$

## sklearn.metrics.mean\_squared\_error

```
sklearn.metrics.mean_squared_error(y_true, y_pred, *, sample_weight=None,
multioutput='uniform_average', squared=True) \[source\]
```

Mean squared error regression loss.

Read more in the [User Guide](#).

**Parameters:**

- y\_true** : array-like of shape (n\_samples,) or (n\_samples, n\_outputs)  
Ground truth (correct) target values.
- y\_pred** : array-like of shape (n\_samples,) or (n\_samples, n\_outputs)  
Estimated target values.
- sample\_weight** : array-like of shape (n\_samples,), default=None  
Sample weights.
- multioutput** : {'raw\_values', 'uniform\_average'} or array-like of shape (n\_outputs,), default='uniform\_average'  
Defines aggregating of multiple output values. Array-like value defines weights used to average errors.
  - 'raw\_values' :  
Returns a full set of errors in case of multioutput input.
  - 'uniform\_average' :  
Errors of all outputs are averaged with uniform weight.
- squared** : bool, default=True  
If True returns MSE value, if False returns RMSE value.

**Returns:**

- loss** : float or ndarray of floats  
A non-negative floating point value (the best value is 0.0), or an array of floating point values, one for each individual target.

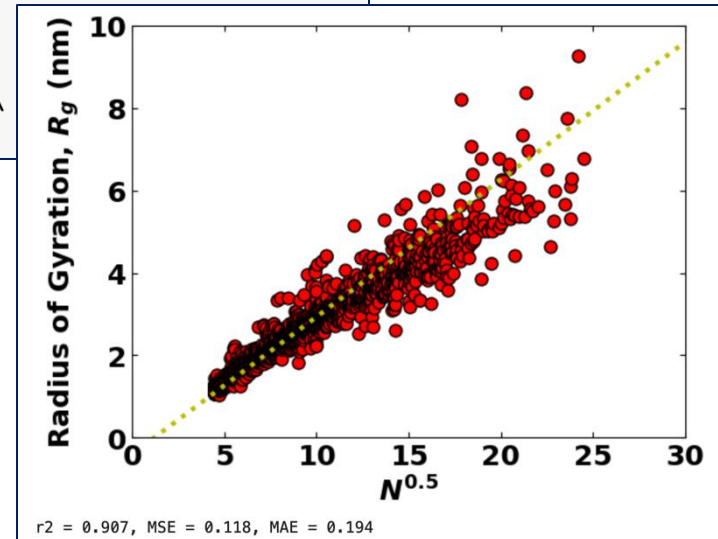
```
# basic set up
Nmax = 900
xline= np.array(range(Nmax+1))*0.5
f = lambda x, th: th[0] + th[1]*x

# fill in parameters
thetas = XXXX # you want thetas to be a 2x1 array in shape!

# make predictions using function
yline = f(xline,thetas)

# examine hypothesis
ax = plot_raw_data(X,y)
ax.plot(xline,yline,color='y',linewidth=3,linestyle=':')
plt.show()

# make predictions from features and compute evaluation metrics
yhat = f(X,thetas) # this is a vector of predictions at the X values given
r2 = r2_score(XXXXX)
rmse = mean_squared_error(XXXXX)
mae = mean_absolute_error(XXXXX)
print("r2 = {:>5.3f}, MSE = {:>5.3f}, MAE = {:>5.3f}"\
      .format(r2,rmse,mae))
```



# Human Hypothesis to the Data

$$R_g = \theta_0 + \theta_1 N^{0.5}$$

## sklearn.metrics.mean\_squared\_error

```
sklearn.metrics.mean_squared_error(y_true, y_pred, *, sample_weight=None,
multioutput='uniform_average', squared=True) \[source\]
```

Mean squared error regression loss.

Read more in the [User Guide](#).

**Parameters:**

- y\_true** : array-like of shape (n\_samples,) or (n\_samples, n\_outputs)  
Ground truth (correct) target values.
- y\_pred** : array-like of shape (n\_samples,) or (n\_samples, n\_outputs)  
Estimated target values.
- sample\_weight** : array-like of shape (n\_samples,), default=None  
Sample weights.
- multioutput** : {'raw\_values', 'uniform\_average'} or array-like of shape (n\_outputs,), default='uniform\_average'  
Defines aggregating of multiple output values. Array-like value defines weights used to average errors.
  - 'raw\_values' :  
Returns a full set of errors in case of multioutput input.
  - 'uniform\_average' :  
Errors of all outputs are averaged with uniform weight.
- squared** : bool, default=True  
If True returns MSE value, if False returns RMSE value.

**Returns:**

- loss** : float or ndarray of floats  
A non-negative floating point value (the best value is 0.0), or an array of floating point values, one for each individual target.

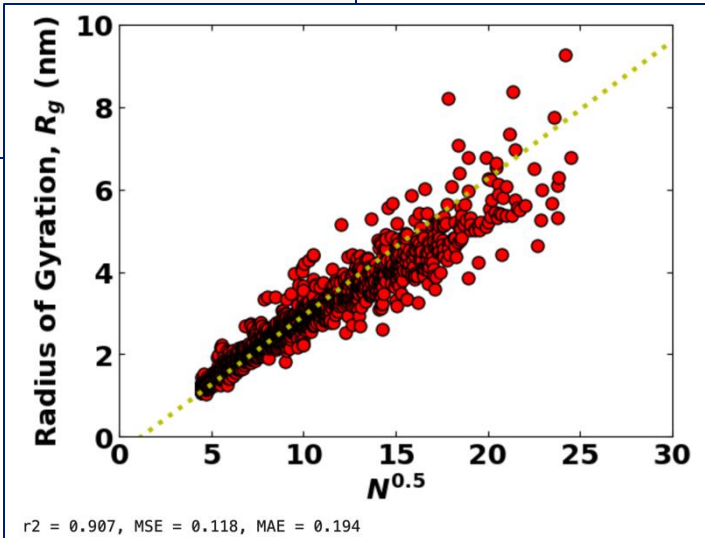
```
# basic set up
Nmax = 900
xline= np.array(range(Nmax+1))*0.5
f = lambda x, th: th[0] + th[1]*x

# fill in parameters
thetas = XXXX # you want thetas to be a 2x1 array in shape!

# make predictions using function
yline = f(xline,thetas)

# examine hypothesis
ax = plot_raw_data(X,y)
ax.plot(xline,yline,color='y',linewidth=3,linestyle=':')
plt.show()

# make predictions from features and compute evaluation metrics
yhat = f(X,thetas) # this is a vector of predictions at the X values given
r2 = r2_score(XXXXX)
rmse = mean_squared_error(XXXXX)
mae = mean_absolute_error(XXXXX)
print("r2 = {:>5.3f}, MSE = {:>5.3f}, MAE = {:>5.3f}"\
      .format(r2,rmse,mae))
```





# Linear Algebraic Solution

Because our model is linear in all its parameters, we can find an exact solution using linear algebra

$$R_g = \theta_0 + \theta_1 N^{0.5}$$

$$Ax = b \Leftrightarrow A^T Ax = A^T b \Leftrightarrow x = \underbrace{(A^T A)^{-1} A^T}_{\text{Moore-Penrose pseudo-inverse}} b$$

NumPy

Linear algebra (numpy.linalg) > numpy.linalg.pinv

## numpy.linalg.pinv

`linalg.pinv(a, rcond=1e-15, hermitian=False)` [\[source\]](#)

Compute the (Moore-Penrose) pseudo-inverse of a matrix.

Calculate the generalized inverse of a matrix using its singular-value decomposition (SVD) and including all *large* singular values.

**Changed in version 1.14:** Can now operate on stacks of matrices

**Parameters:**

- a** : (...) *M, N* array\_like  
Matrix or stack of matrices to be pseudo-inverted.
- rcond** : (...) array\_like of float  
Cutoff for small singular values. Singular values less than or equal to `rcond * largest_singular_value` are set to zero. Broadcasts against the stack of matrices.
- hermitian** : bool, optional  
If True, *a* is assumed to be Hermitian (symmetric if real-valued), enabling a more efficient method for finding singular values. Defaults to False.

**New in version 1.17.0.**

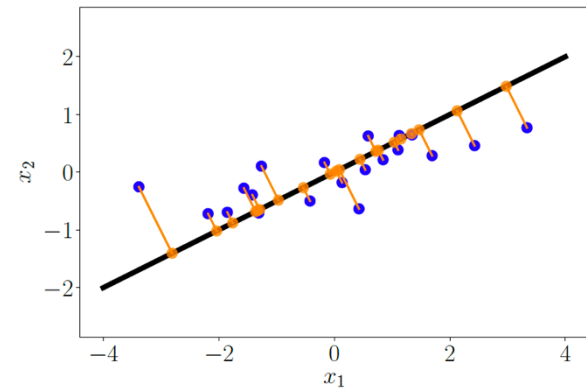
```
N = len(y)
M = 2
A = np.ones((N,M))
A[:,1] = X[:,1]
theta0pt = XXX # use np.linalg.pinv
yhat = f(X,theta0pt)
r2 = r2_score(XXX)
mse = mean_squared_error(XXX)
mae = mean_absolute_error(XXX)
print("theta_0 = {:>8.4f}".format(theta0pt[0]))
print("theta_1 = {:>8.4f}".format(theta0pt[1]))
print("r2 = {:>5.3f}, MSE = {:>8.5f}, MAE = {:>5.3f}"\
      .format(r2,rmse,mae))

theta_0 = -0.0384
theta_1 = 0.2827
r2 = 0.941, MSE = 0.11762, MAE = 0.152
```

*Moore-Penrose  
pseudo-inverse*

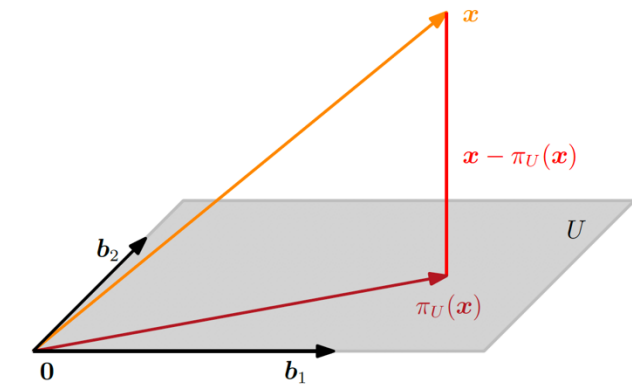


# Reminder: least-squares is minimizing projection error



We want  $\pi_U(\mathbf{x}) = \sum_{i=1}^M \lambda_i \mathbf{b}_i = \mathbf{B}\boldsymbol{\lambda}$  given  $\mathcal{B}_U = (\mathbf{b}_1, \dots, \mathbf{b}_M)$

such that  $\pi_U(\mathbf{x}) - \mathbf{x}$  is **orthogonal** to  $U$  and its **distance** is **minimized**



Assuming the dot product as the inner product...

$$\implies \mathbf{b}_i^T (\mathbf{x} - \mathbf{B}\boldsymbol{\lambda}) = 0, \quad i = 1, \dots, M$$

$$\begin{aligned} &\iff \\ &\mathbf{B}^T (\mathbf{x} - \mathbf{B}\boldsymbol{\lambda}) = \mathbf{0} \iff \mathbf{B}^T \mathbf{B}\boldsymbol{\lambda} = \mathbf{B}^T \mathbf{x} \\ &\text{normal equation} \end{aligned}$$

$$\boldsymbol{\lambda} = (\mathbf{B}^T \mathbf{B})^{-1} \mathbf{B}^T \mathbf{x}$$

if  $\mathbf{B}$  describes an orthonormal basis??  $\implies \mathbf{P}_\pi = \mathbf{B}(\mathbf{B}^T \mathbf{B})^{-1} \mathbf{B}^T$

# Optimization with a Loss Function

Training (optimizing parameters for) supervised ML models requires specification of a loss function and means to navigate it; let's take a look at a simple loss function

```
def loss(x,y,theta):
    ''' Function to calculate cost function assuming a hypothesis of form
     $y^{\wedge} = X * \theta$ 
    Inputs:
    x = array of dependent variable
    y = array of training examples
    theta = array of parameters for hypothesis

    Returns:
    E = cost function
    '''
    n = len(y) #number of training examples
    features = np.ones((n,len(theta))) # X
    features[:,1] = x[:]
    ypred = features@theta # predictions with current hypothesis
    E = np.sum((ypred[:,0]-y[:])**2)/n #Cost function
    return E

def plot_loss(t0,t1):
    #Initialize E as a matrix to store cost function values
    E = np.zeros((len(t0),len(t1)))

    # Populate matrix
    for i,theta0 in enumerate(theta0s):
        for j,theta1 in enumerate(theta1s):
            theta_ij = np.array([[theta0,theta1]])
            E[i,j] = loss(X,y,theta_ij)
    t0g,t1g = np.meshgrid(t0,t1)
    fig = plt.figure(figsize=(15,4))
    ax1 = fig.add_subplot(1,2,1,projection='3d')
    surf = ax1.plot_surface(t0g, t1g, E, linewidth=0, \
                           antialiased=False, cmap='coolwarm')
    ax1.set_xlabel(r"$\theta_1$")
    ax1.set_ylabel(r"$\theta_0$")
    ax1.set_zlabel(r"$E$")
    ax2 = fig.add_subplot(1,2,2)
    CS = ax2.contour(t0g,t1g,E.T,np.logspace(-3,2,25),cmap='coolwarm')
    ax2.set_xlabel(r"$\theta_0$")
    ax2.set_ylabel(r"$\theta_1$")

    return fig,ax1,ax2
```

#Define grid over which to calculate the loss function

N = 50

theta0Rng = [-5,5]

theta1Rng = [-0.5,1.5]

theta0s = np.linspace(theta0Rng[0],theta0Rng[1],N)

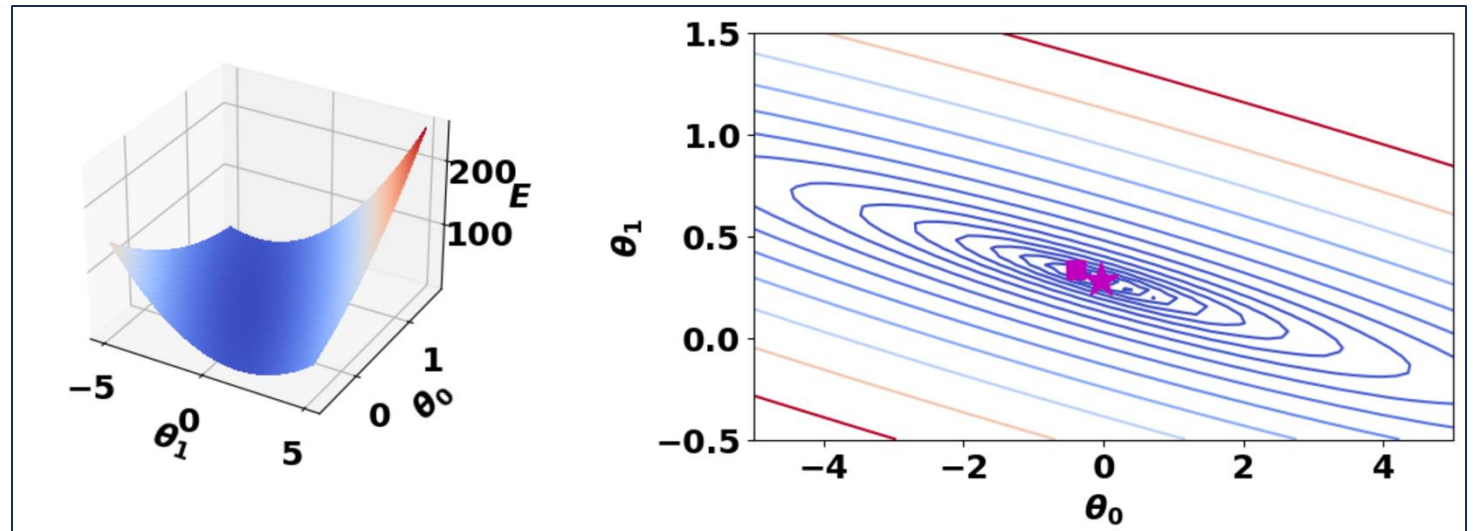
theta1s = np.linspace(theta1Rng[0],theta1Rng[1],N)

fig,ax1,ax2 = plot\_loss(theta0s,theta1s)

ax2.plot(thetas[0],thetas[1],marker='s',color='m',markersize=10)

ax2.plot(theta0pt[0],theta0pt[1],marker='\*',color='m',markersize=20)

plt.show()



# Results with Gradient Descent Optimization

Most optimization methods make use of information regarding the gradients of the loss function with respect to the parameters; these guide selection of the next parameters

```
def E2loss(yhat,y):
    return np.sum((np.squeeze(yhat)[:]-y[:])**2)/len(y)

def Grad_Descent(x,y,theta,alpha,nIters,x_te=None,y_te=None):
    '''Gradient descent algorithm
    Inputs:
    x = dependent variable
    y = training data
    theta = parameters
    alpha = learning rate
    iters = number of iterations
    Output:
    theta = final parameters
    E = array of cost as a function of iterations
    '''
    n = len(y) #number of training examples
    features = np.ones((n,len(theta)))
    features[:,1] = x[:]
    yhat = features@theta # predictions with current hypothesis
    E_hist = [E2loss(yhat,y)]

    if x_te is not None:
        E_hist_te = [E2loss(f(x_te,theta),y_te)]

    for i in range(nIters):
        e = yhat[:,0] - y[:]
        theta = theta - (alpha*e[:,np.newaxis].T@features).T #
        yhat = features@theta # predictions with current hypothesis
        E_hist.append(E2loss(yhat,y))
        if x_te is not None:
            E_hist_te.append(E2loss(f(x_te,theta),y_te))

    if x_te is not None:
        return theta,E_hist,E_hist_te
    else:
        return theta,E_hist
```

## Gradient Descent

$$\mathbf{x}_{i+1} = \mathbf{x}_i - \gamma_i [\nabla f(\mathbf{x}_i)]^T$$

```
th0 = XXX
alpha = 8e-6
nIters = 5000
thetaGD, EGD = Grad_Descent(X,y,th0,alpha,nIters)
print(XXX)
print(XXX)

theta_0 = -0.0384
theta_1 = 0.2827
```

```
fig,ax = plt.subplots()
ax.plot(np.array(range(nIters+1))+1,np.array(EGD),\
        linestyle='-',color = 'k',linewidth=3)
plt.xscale("log")
plt.yscale("log")
ax.set_xlabel("Iterations")
ax.set_ylabel("Loss")
plt.show()

# examine solution
ax = plot_raw_data(X,y)
ax.plot(xline,f(xline,thetaGD),color='k',linewidth=3,linestyle='--')
plt.show()

r2 = r2_score(XXX)
mse = mean_squared_error(XXX)
mae = mean_absolute_error(y,XXX)
print("r2 = {:.5f}, MSE = {:.5f}, MAE = {:.5f}"\
      .format(r2,rmse,mae))

fig,ax1,ax2 = plot_loss(theta0s,theta1s)
ax2.plot(thetas[0],thetas[1],marker='s',color='m',markersize=10)
ax2.plot(thetaOpt[0],thetaOpt[1],marker='*',color='m',markersize=10)
ax2.plot(thetaGD[0],thetaGD[1],marker='*',color='y',markersize=10)
plt.show()
```

