

Multivariate Regression & Feature Scaling

From Lectures 3/4

ML structures problems around identifying inputs (features) and outputs (labels); the rest is algorithmic details.

- **supervised learning**
 - make predictions on labels from features
 - (semi- and self-) supervised address data scarcity by principled cheating
- **unsupervised learning**
 - learn patterns and relationships in data from features
- **reinforcement learning**
 - agent learns to operate (actions) within environment (state) subject to a reward system

Supervised ML is indeed fancy curve-fitting

- *the loss function (user-defined) determines what is optimal*
- *most familiar with l_2 -loss (MSE)*
- *optimization frequently proceeds by taking gradients of loss function with respect to model parameters*

To understand basic mechanics

- **linear regression**
 - a model that is linear in its parameters (not necessarily linear in output-input relationship)
- **gradient descent**
 - *workhorse in parameter optimization with most other algorithms correcting deficiencies*

Unpacking “Linear Least-Squares” Regression

Linear model

$$f(x) = \sum_{i=0}^m \theta_i g_i(x)$$

A linear model is linear in its parameters

Example:

$$m = 2; g_0 = 1; g_1 = x; g_2 = x^2$$

$$f(x) = \theta_0 + \theta_1 x + \theta_2 x^2$$

Least-squares

General notion of optimization

Given $\{(\mathbf{x}_i, y_i)\}$ produce “optimal” f $\hat{y} = f(\mathbf{x}, \boldsymbol{\theta})$
that minimizes some error metric (“loss”) $\mathcal{E}(\{y_k, \hat{y}_k\})$

*The “least squares” aspect specifies the loss
as related to L2-norm*

$$\min_{\boldsymbol{\theta}} \mathcal{E}(f) = \min_{\boldsymbol{\theta}} \sum_{k=1}^n |e_k|^2 \quad \mathcal{E}_2(f) = \sqrt{\frac{1}{n} \sum_{k=1}^n |e_k|^2}$$

- **Linear least-squares regression** can be exactly solved in the framework of linear algebra using techniques for matrix inversion/diagonalization (see notes)
- **Machine learning** is distinguished by having *non-linear parametric dependence* and (possibly) *different loss functions*; this necessitates alternatives to pure linear algebra.

Unpacking “non-Linear Least-Squares” Regression

non-Linear model

$$f(x) = \sum_{i=0}^m g_i(x, \theta)$$

Example:

$$f(x, \theta) = \theta_0 \cos(\theta_1 x + \theta_2) + \theta_3$$

not important that the x is in the cosine, but rather the non-linear relationship amongst thetas is the sticking point

Least-squares

$$\mathcal{E}(\theta) = \sum_{k=1}^n (f(x_k, \theta) - y_k)^2; \quad \frac{\partial \mathcal{E}}{\partial \theta_i} = 0 \quad \forall i$$

$$\Rightarrow \sum_{k=1}^n (f(x_k, \theta) - y_k) \frac{\partial f}{\partial \theta_i} = 0 \quad \forall i$$


$$\sum_{k=1}^n e_k \frac{\partial f}{\partial \theta_i} = 0 \quad \forall i$$

chain rule is your friend, and this strategy can be applied to other losses

- **Linear least-squares regression** can be exactly solved in the framework of linear algebra using techniques for matrix inversion/diagonalization (see notes)
- **Machine learning** is distinguished by having *non-linear parametric dependence* and (possibly) *different loss functions*; this necessitates alternatives to pure linear algebra. Machine learning models have multivariable (high-dimensional) inputs!

Multivariate Regression

Previous example illustrates working with simple bivariate data; there is only one feature per label



While we may often work with one label as output, the labels will be modeled as a function of many variables

$$f(x) = \theta_0 + \theta_1 x$$

$$\theta^{(i+1)} = \theta^{(i)} - \gamma^{(i)} \left[\nabla \mathcal{E}(f(\theta^{(i)})) \right]^T$$

Extension to multiple variables is straightforward

$$f(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

Size (ft ²)	No. Units	Capital Cost (\$)
2104	3	399900
1600	3	329900
2400	3	369000
1416	2	232000
3000	4	539900

More generally...

$$f(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots = \sum_{i=1}^m \theta_i x_i; x_0 = 1$$
$$= \theta^T x$$

Notice that nothing related to cost function or gradient descent becomes any more complicated!

Multivariate Regression

$$\{(\mathbf{x}_k, y_k)\} \text{ for } k = 1, \dots, n \quad f(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x} \quad \boldsymbol{\theta} \in \mathbb{R}^m$$


$$\mathbf{X} = [\mathbf{x}_1 \ \mathbf{x}_2 \ \dots \ \mathbf{x}_n]; \ \mathbf{y} = [y_1, y_2, \dots, y_n]^T$$

$$\mathcal{E} = \sum_{k=1}^m (\boldsymbol{\theta}^T \mathbf{x}_k - y_k)^2$$

$$\mathcal{E} = (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

$$\nabla_{\boldsymbol{\theta}} \mathcal{E} = 0$$

$$\Rightarrow \boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$



Practice: Multivariate Regression

ex1data2.csv

Size (ft ²)	No. Units	Capital Cost (\$)
2104	3	399900
1600	3	329900
2400	3	369000
1416	2	232000
3000	4	539900

⋮

Task:

1. Build a regression model with the data provided to estimate capital cost as a function of both size and number of units.
2. Use the model to predict cost for space of 1650 ft² and 3 units.
3. Solve the same problem with the normal equation to verify the numerical result

Feature Scaling

ex1data2.csv

Size (ft ²)	No. Units	Capital Cost (\$)
2104	3	399900
1600	3	329900
2400	3	369000
1416	2	232000
3000	4	539900

⋮

Task:

1. Build a regression model with the data provided to estimate capital cost as a function of both size and number of units.
2. Use the model to predict cost for space of 1650 ft² and 3 units.
3. Solve the same problem with the normal equation to verify the numerical result

Spin:

We are going to introduce here the concept of ***feature scaling***, which is a transformation of data (mapping) such that all values are on similar scales (*usually in the range of -1 to 1 or 0 to 1*):

- is typically good practice for constructing machine learning models but is essential for any algorithms that compute ***distances***; it is not necessary for algorithms that rely on “rules” (**decision trees**)
- facilitates faster convergence during model training (*why?*)
- provides better initial choice over hyperparameters
- avoids unintentional bias of some inputs

Feature Scaling

Common feature scaling approaches

linear mappings

- Min Max Scaling
- Standard Scaling
- Max Abs Scaling
- Robust Transformer

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

$$x' = \frac{x - \mu}{\sigma}$$

$$x' = \frac{x}{\max(|x|)}$$

$$x' = \frac{x - q_2}{q_3 - q_1}$$

nonlinear mappings

- Quantile/Rank Transformer
- Power Transformer
- Norm Scaling

can map data to uniform or Gaussian distribution
(can distort correlations/distances)


usually for mapping data to Gaussian distribution
(Box-Cox, Yeo-Johnson are different flavors)

$$x' = \frac{x}{||x||}$$


Feature Scaling: Comparing Non-Linear Transforms

original data


Lognormal




Chi-squared




Weibull




Gaussian




Uniform




Bimodal




After Box-Cox
 $\lambda = 0.02$




After Box-Cox
 $\lambda = 0.28$




After Box-Cox
 $\lambda = 12.12$




After Box-Cox
 $\lambda = 6.99$




After Box-Cox
 $\lambda = 0.68$




After Box-Cox
 $\lambda = 1.61$




After Yeo-Johnson
 $\lambda = -0.8$




After Yeo-Johnson
 $\lambda = -0.11$




After Yeo-Johnson
 $\lambda = 23.52$




After Yeo-Johnson
 $\lambda = 7.05$



After Yeo-Johnson
 $\lambda = 0.64$




After Yeo-Johnson
 $\lambda = 1.62$




non-parametric


After Quantile transform




After Quantile transform




After Quantile transform




After Quantile transform



After Quantile transform



After Quantile transform



Feature Scaling: Scikit-learn

Basic Functionality

$X \leftarrow (n\text{samples}, m\text{features})$
 $y \leftarrow (n\text{samples}, 1)$

```
from sklearn import preprocessing

Xscaler = preprocessing.TheScaler() .fit(X)
X_sc     = Xscaler.transform(X)

yscaler = preprocessing.TheScaler() .fit(y)
y_sc     = yscaler.transform(y)

Xagain  = Xscaler.inverse_transform(X_sc)
yagain  = yscaler.inverse_transform(y_sc)
```

Notebook Exercise

Activity: Premise and Objective

Machine learning excels (over standard methods and human intuition) when confronted with **high-dimensional** data. We will explore multivariate regression in the notebook to see how this does not remarkably change problem structure or present major challenges

```
# modules used by Prof. Webb
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn import preprocessing
from scipy.optimize import minimize
```

Progression of the notebook:

- Look at the data
- Perform feature scaling
- Optimize using gradient descent (from scratch)
- Compare to linear-algebraic solutions
- Compare to python optimization

linear multivariate model: $c(s, n) = \theta_0 + \theta_1 s + \theta_2 n$

capital
cost, c

equipment
size, s

number of
units, n

Size, Units, Capital

2104, 3, 399900

1600, 3, 329900

2400, 3, 369000

1416, 2, 232000

3000, 4, 539900


1985, 4, 299900

1534, 3, 314900

1427, 3, 198999

1380, 3, 212000

:



Activity: Data, Scaling, Gradient Descent

```
# extract the data
url_for_data = "https://raw.githubusercontent.com/
data = pd.read_csv(url_for_data)
print(data.head(5))

# partition into features (X)/ labels (y)
X = np.array(data.iloc[:, :2])
y = np.array(data.iloc[:, 2])
```

```
def E2loss(ypred,y):
    return np.sum(ypred[:]-y[:])**2/len(y)

def Grad_Descent(X,y,theta,alpha,nIters):
    '''Gradient descent algorithm
    Inputs:
    X = dependent variables, comes in at Nx2
    y = training data, comes in at Nx1
    theta = parameters, comes in as 3x1
    alpha = learning rate
    iters = number of iterations
    Output:
    theta = final parameters
    E = array of cost as a function of iterations
    '''
    n = len(y) #number of training examples
    features = np.ones((n,len(theta)))
    features[:,1:] = X[:, :]
    ypred = features@theta # predictions with current hypothesis

    E_hist = [E2loss(ypred,y)]
    for i in range(nIters):
        e = ypred[:,0] - y[:,0]
        theta = theta - (alpha*e[:,np.newaxis].T@features).T #
        ypred = features@theta # predictions with current hypothesis
        E_hist.append(E2loss(ypred,y))

    return theta,E_hist
```


Using scikit-learn to perform variable scaling/transformation

```
# example for scikit learn
X_stdscaler = preprocessing.StandardScaler().fit(X)
X_std = X_stdscaler.transform(X)
y_stdscaler = preprocessing.StandardScaler().fit(y.reshape(-1,1))
y_std = y_stdscaler.transform(y.reshape(-1,1))
```

```
###BEGIN SOLUTION
theta0 = np.array([[0.01],[0.01],[0.01]])
theta_GD, E_GD = Grad_Descent(X_std,y_std,theta0,0.01, 100)
plt.plot(E_GD[:])

###END SOLUTION
Xtest = np.array([[1650,3]]) # input
Xtest_sc = X_stdscaler.transform(Xtest)
Ytest_sc = theta_GD[0] + Xtest_sc@theta_GD[1:,0]
Ytest = y_stdscaler.inverse_transform(Ytest_sc.reshape(-1,1))
print("The optimal parameters found are \n", theta_GD)
print("The prediction from ", Xtest, "is ", Ytest)
```

```
The optimal parameters found are
[[-7.52436308e-17]
 [ 8.84765988e-01]
 [-5.31788196e-02]]
The prediction from [[1650 3]] is [[293081.464336]]
```



Activity: Comparison to Normal Equations

$$c(s, n) = \theta_0 + \theta_1 s + \theta_2 n$$

$$\mathcal{E} = \sum_{k=1}^m (\boldsymbol{\theta}^T \mathbf{x}_k - y_k)^2$$

$$\mathcal{E} = (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

$$\nabla_{\boldsymbol{\theta}} \mathcal{E} = 0$$

$$\Rightarrow \boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

```
##Normal Equations
```

```
Xtest    = np.array([[1,1650,3]]) # input
```

```
###BEGIN SOLUTION
```

```
#Set up feature matrix
```

```
n        = len(y) #number of training examples
```

```
features = np.ones((n,3))
```

```
features[:,1:] = X[:,1:]
```

```
XXinv = np.linalg.inv(features.T @ features)
```

```
theta = XXinv @ features.T @ y
```

```
cost = np.dot(Xtest,theta)
```

```
print("The optimal parameters found are \n", theta)
```

```
print("The prediction from ", Xtest, "is ", cost)
```

```
###END SOLUTION
```

```
The optimal parameters found are
```

```
[89597.9095428    139.21067402 -8738.01911233]
```

```
The prediction from [[ 1 1650   3]] is [293081.4643349]
```

Activity: Comparison to Package Optimization

scipy.optimize.minimize

```
scipy.optimize.minimize(fun, x0, args=(), method=None, jac=None, hess=None,
                        hessp=None, bounds=None, constraints=(), tol=None, callback=None, options=None)
```

Minimization of scalar function of one or more variables. [\[source\]](#)

Parameters: **fun** : callable

The objective function to be minimized.

`fun(x, *args) -> float`

where `x` is a 1-D array with shape `(n,)` and `args` is a tuple of the fixed parameters needed to completely specify the function.

x0 : ndarray, shape (n,)

Initial guess. Array of real elements of size `(n,)`, where `n` is the number of independent variables.

args : tuple, optional

Extra arguments passed to the objective function and its derivatives (`fun`, `jac` and `hess` functions).

method : str or callable, optional

Type of solver. Should be one of

- 'Nelder-Mead' ([see here](#))
- 'Powell' ([see here](#))
- 'CG' ([see here](#))
- 'BFGS' ([see here](#))
- 'Newton-CG' ([see here](#))
- 'L-BFGS-B' ([see here](#))
- 'TNC' ([see here](#))
- 'COBYLA' ([see here](#))
- 'SLSQP' ([see here](#))
- 'trust-constr' ([see here](#))
- 'dogleg' ([see here](#))
- 'trust-ncg' ([see here](#))
- 'trust-exact' ([see here](#))
- 'trust-krylov' ([see here](#))
- custom - a callable object, see below for description.

If not given, chosen to be one of **BFGS**, **L-BFGS-B**, **SLSQP**, depending on whether or not the problem has constraints or bounds.

```
def E2lossv2(theta,X,y):
    thetaVec = theta.T
    ypred     = theta[0] + X@thetaVec[1:]
    return np.sum(ypred[:]-y[:])**2/len(y)

loss = lambda theta: E2lossv2(theta,X_std,y_std)
theta0 = np.array([8.11502877e-05,8.26550700e-01,4.63616701e-03])
res = minimize(loss,theta0,tol=1e-12,method='BFGS')
theta_py = res.x[:]
Ytest_sc = theta_py[0] + Xtest_sc@theta_py[1:].T
Ytest     = y_stdscaler.inverse_transform(Ytest_sc.reshape(-1,1))
print("The optimal parameters found are \n", theta_py)
print("The prediction from ", Xtest, "is ", Ytest)
```

The optimal parameters found are

`[-7.44660311e-09 8.26550700e-01 4.63616701e-03]`

The prediction from `[[1650 3]]` is `[[294676.60596926]]`