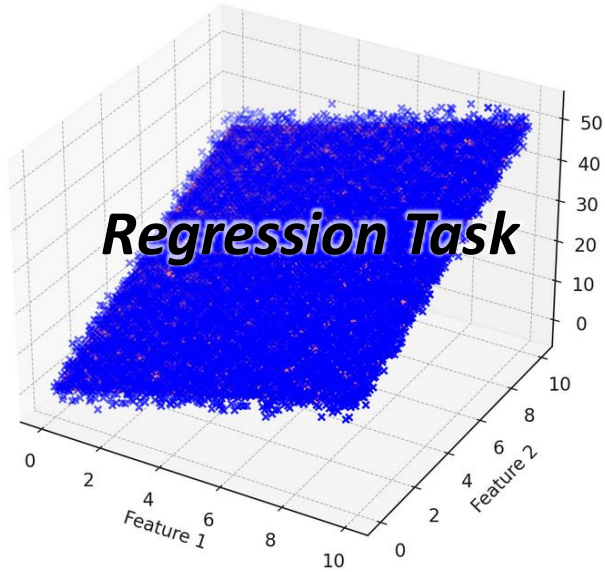
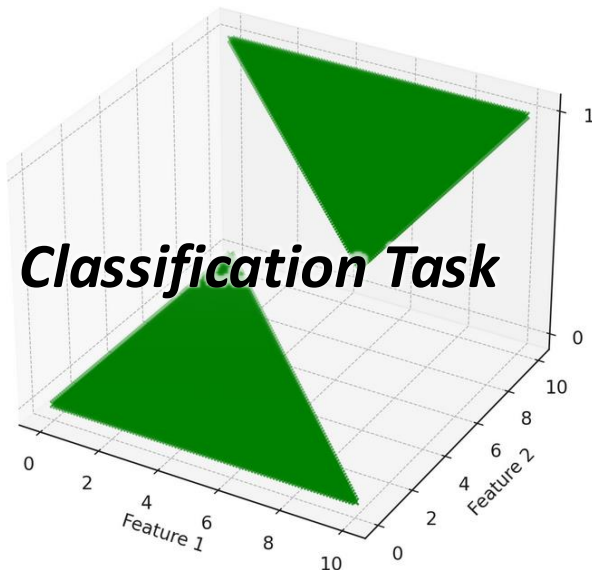


# Classification versus Regression

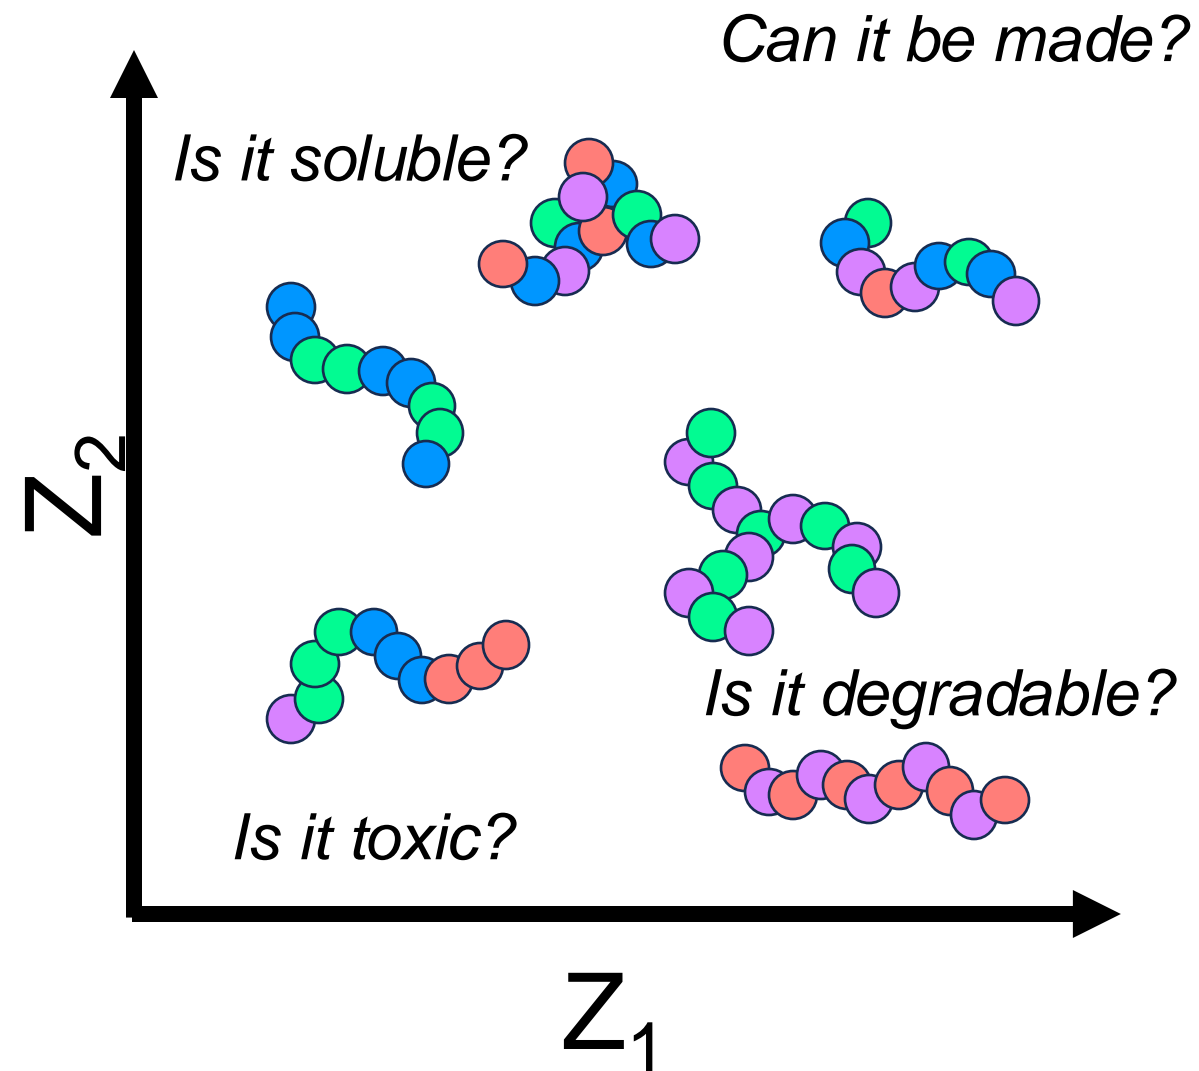
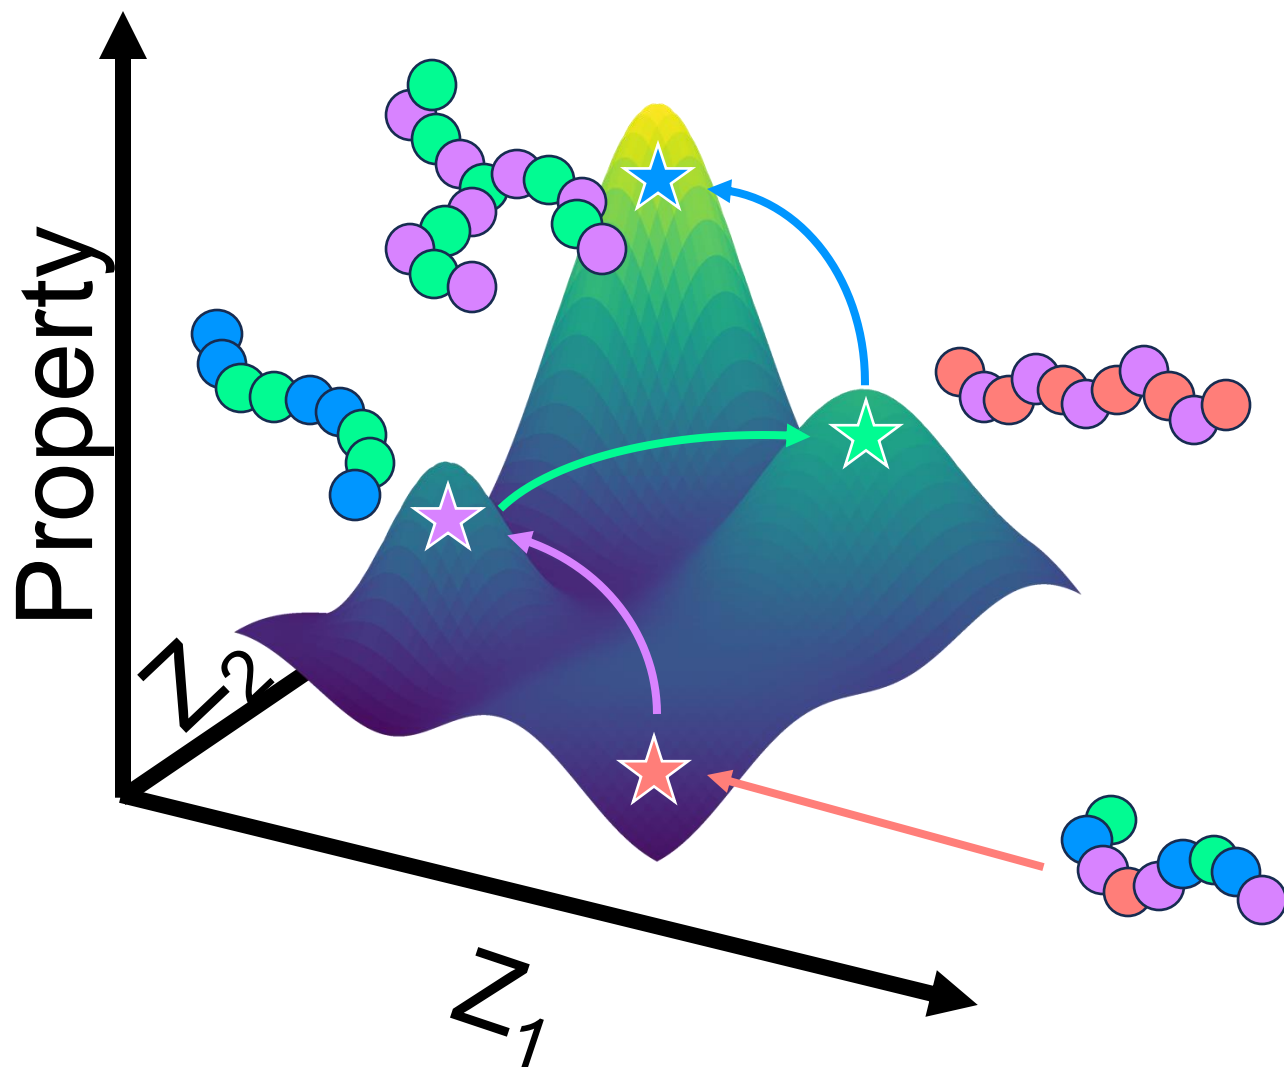


- **Output:** Predicts **continuous values** related to chemical or materials properties.
- **Goal:** Estimate a numerical value based on input features in a chemical or engineering context.
- **Examples:**
  - Predicting the **yield strength** of a new polymer based on its molecular structure.
  - Estimating the **reaction rate** of a chemical process at different temperatures. Predicting the **lifetime** or **degradation rate** of a battery under various operating conditions.

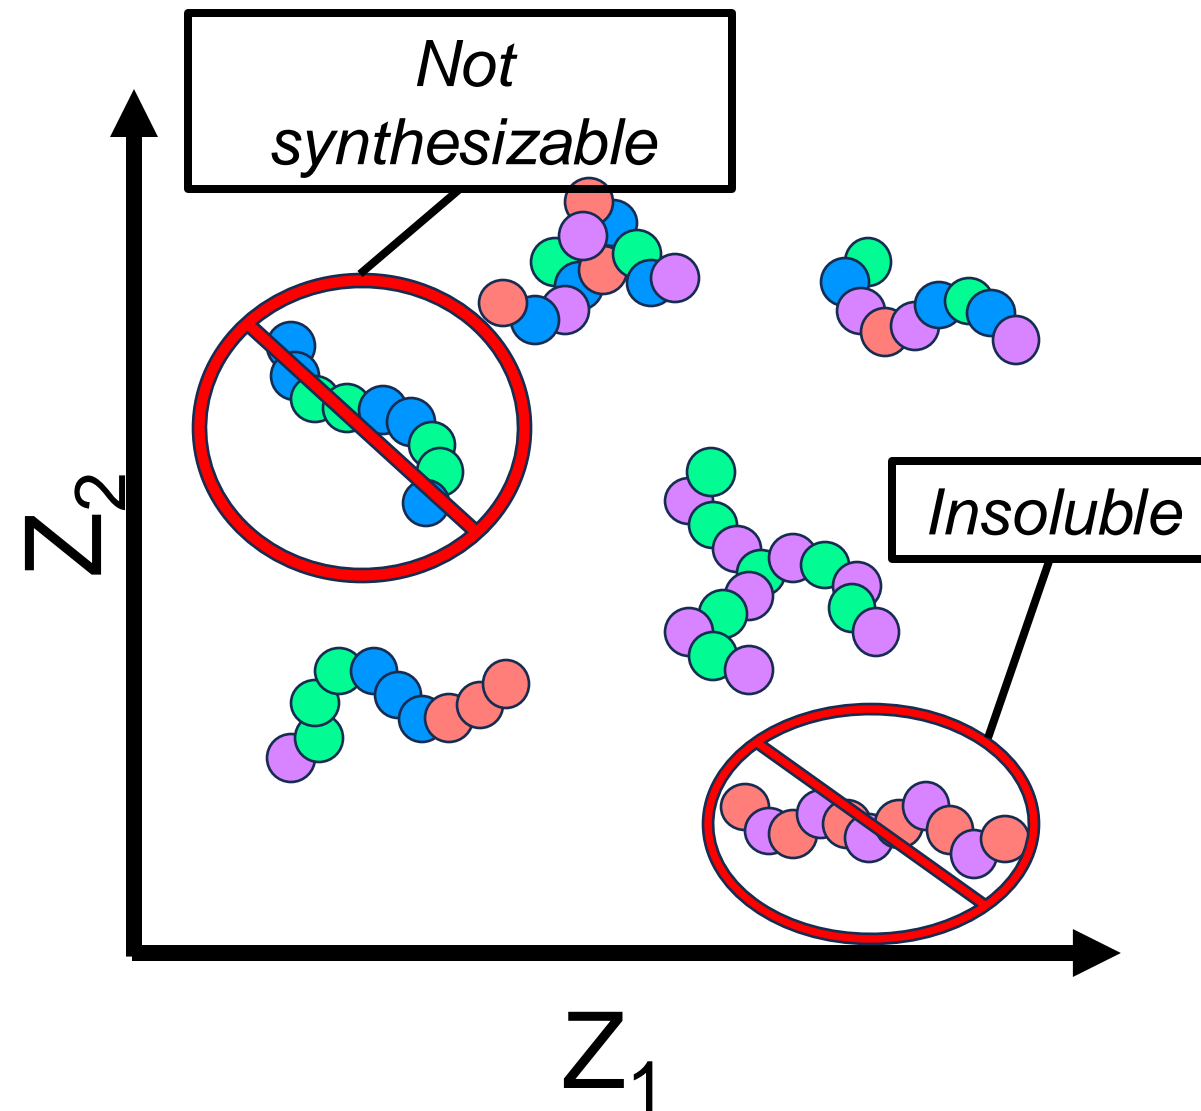
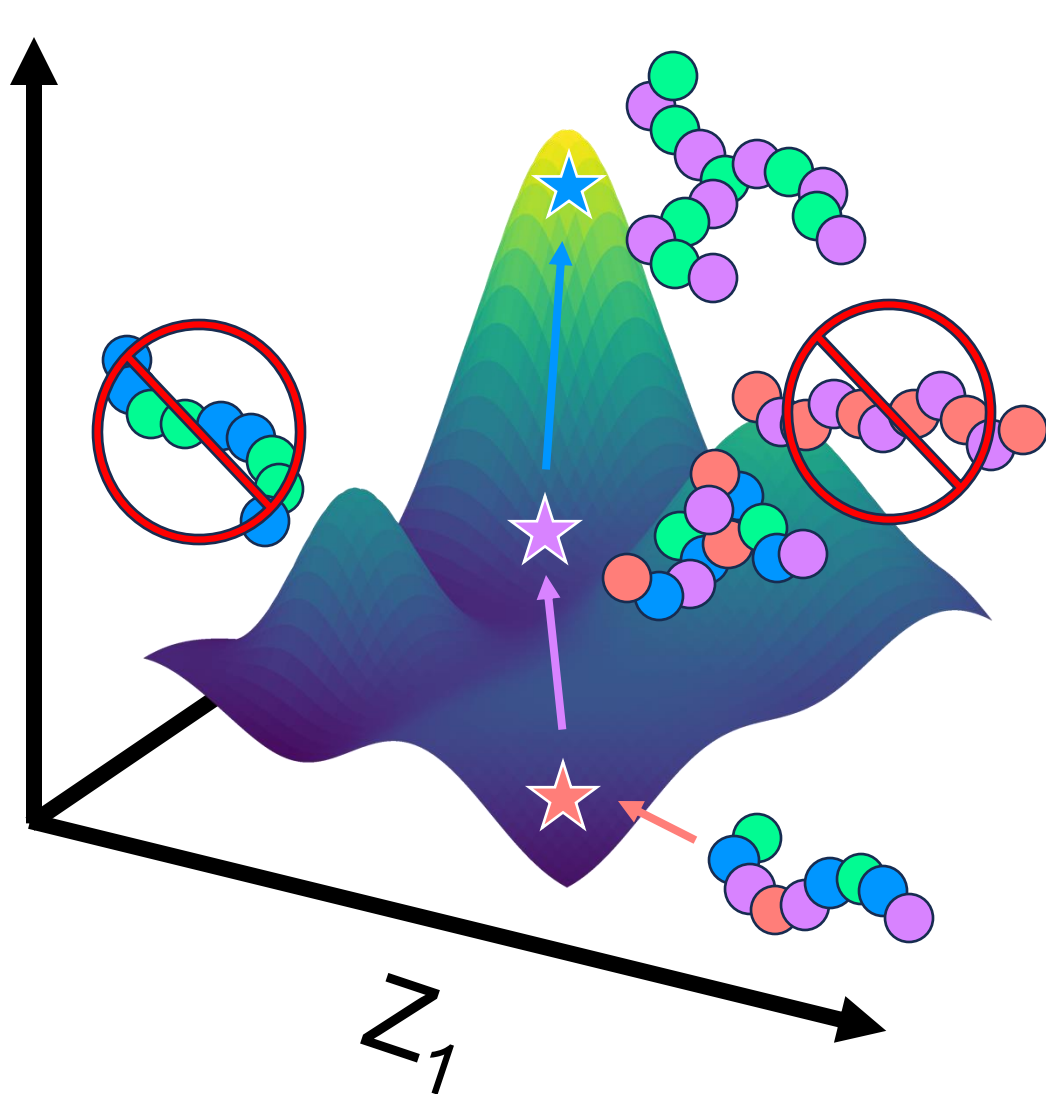


- **Output:** Predicts **discrete labels** or **categories** related to chemical or materials performance.
- **Goal:** Assign input data to one of several predefined categories in the field of chemistry or materials science.
- **Examples:**
  - Determining whether a chemical reaction will be **exothermic** or **endothermic** given certain reactants and conditions.
  - Classifying polymers as **soluble** or **insoluble** in a particular solvent.
  - Predicting whether a material will be **brittle** or **ductile** based on its microstructure.

# Classification in Chemical/Materials Optimization



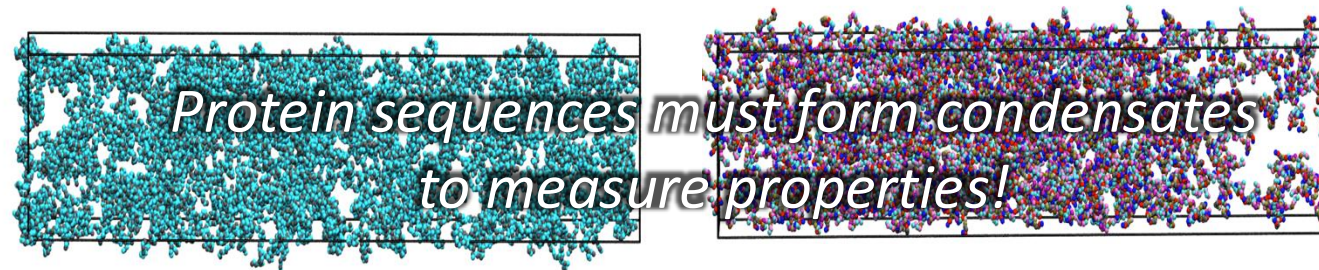
# Classification in Chemical/Materials Optimization



# composition design



*Polymers must remain soluble  
for enzyme assay and  
characterization*



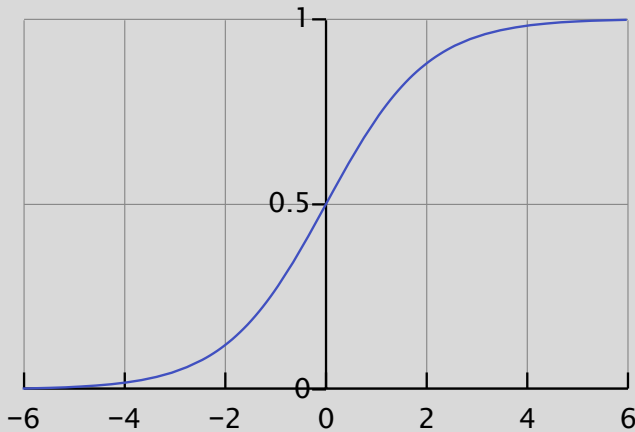


# Logistic Regression: Pathway towards Classification

In **logistic regression**, we want to restrict our predictions to be on the interval  $[0,1]$  to represent probabilities of a class

## Logistic Function

$$f(x) = \frac{L}{1 + e^{-k(x-x_0)}}$$



$$L = 1, k = 1, x_0 = 0$$

other Sigmoid shapes can be used for analogous purpose

The essential premise of a **logistic model** is to represent the **log-odds** of a label as a linear combination of the features

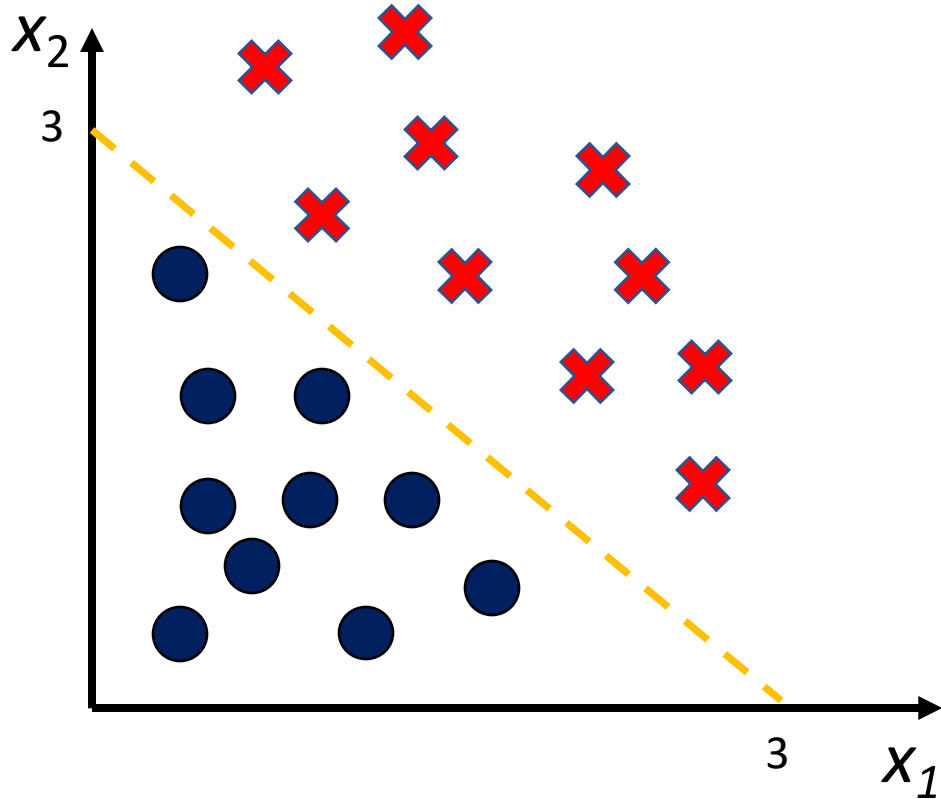
$$\ell = \log_b \frac{p}{1-p} = \mathbf{x}^T \boldsymbol{\theta}$$

$$\Rightarrow p = \frac{1}{1 + b^{-\mathbf{x}^T \boldsymbol{\theta}}} \xrightarrow{b=e} \frac{1}{1 + e^{-\mathbf{x}^T \boldsymbol{\theta}}}$$

Model predictions:  $\hat{y} \leftarrow f(x) = p(y = 1 | \mathbf{x}, \boldsymbol{\theta})$

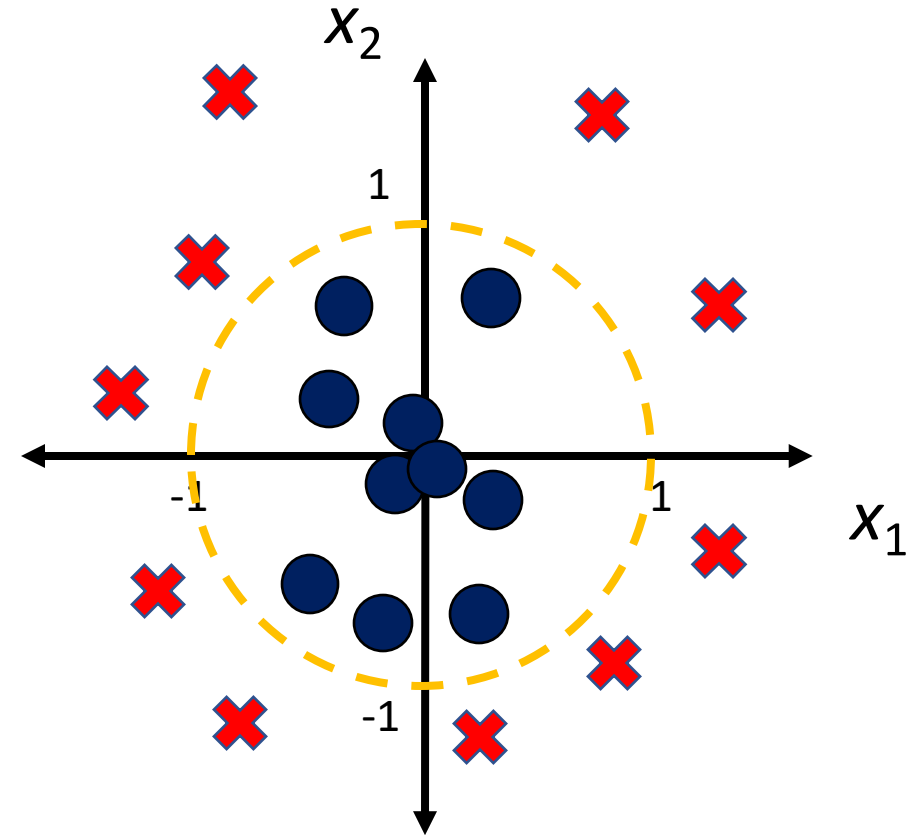
for  $f(x) = 0.7$ , we interpret that to mean a 70% chance that  $y = 1$

# Parameterizing Decision Boundaries



$$f(x_1, x_2) = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \theta_2 x_2)}}$$

*what would be a good set of thetas?*



$$f(x_1, x_2) = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)}}$$

*what about in this case?*

# Approaching a Cost Function

As in linear regression, we will identify optimal parameters via minimization of an appropriate cost function; here, we have something to think about

Suppose model predictions are supplied via

$$\hat{y} \leftarrow f(x) = p(y = 1 | \mathbf{x}, \boldsymbol{\theta}) = \frac{1}{1 + e^{-\mathbf{x}^T \boldsymbol{\theta}}}$$

*Can you anticipate any potential issues with our previous mean-squared error metric?*

$$\mathcal{E}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

*As an alternative, we might consider*

$$\mathcal{E}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \varepsilon(\hat{y}_i, y_i)$$

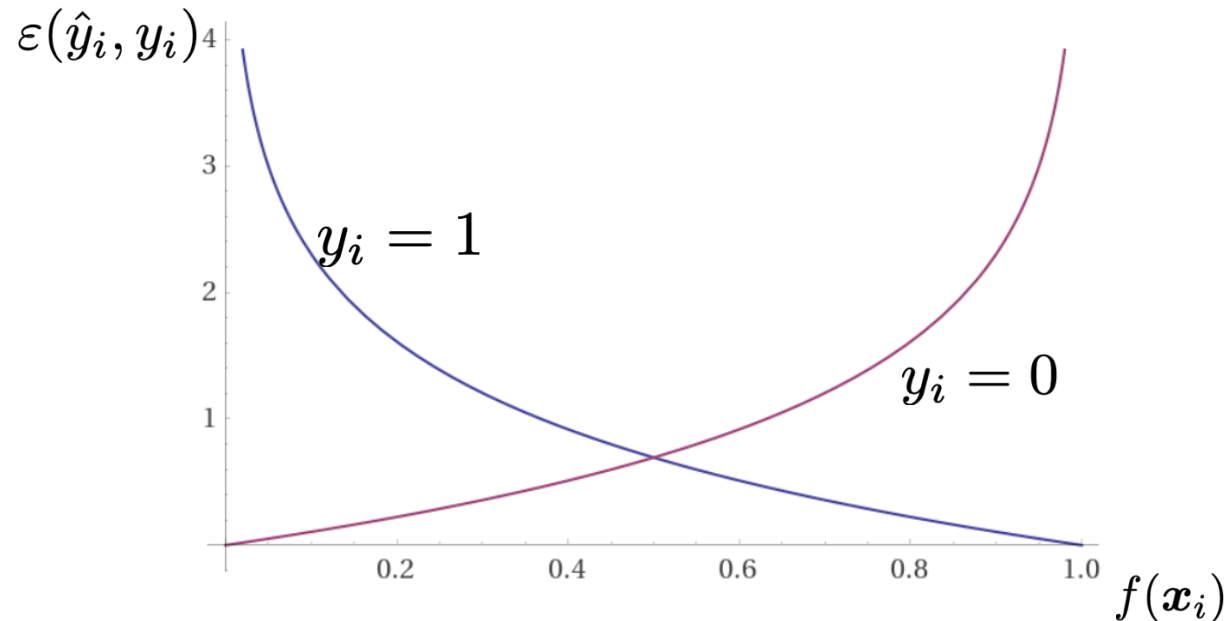
$$\varepsilon(\hat{y}_i, y_i) = \begin{cases} -\log[f(\mathbf{x}_i)], & \text{if } y_i = 1 \\ -\log[1 - f(\mathbf{x}_i)], & \text{if } y_i = 0 \end{cases}$$

# Approaching a Cost Function

$$\hat{y} \leftarrow f(x) = p(y = 1 | \mathbf{x}, \boldsymbol{\theta}) = \frac{1}{1 + e^{-\mathbf{x}^T \boldsymbol{\theta}}}$$

$$\mathcal{E}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \varepsilon(\hat{y}_i, y_i)$$

$$\varepsilon(\hat{y}_i, y_i) = \begin{cases} -\log[f(\mathbf{x}_i)], & \text{if } y_i = 1 \\ -\log[1 - f(\mathbf{x}_i)], & \text{if } y_i = 0 \end{cases}$$





# Minimizing the Cost Function

$$\varepsilon(\hat{y}_i, y_i) = -y_i \log f(\mathbf{x}_i) - (1 - y_i) \log [1 - f(\mathbf{x}_i)]$$

$$\mathcal{E}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \varepsilon(\hat{y}_i, y_i) \quad \hat{y} \leftarrow f(\mathbf{x}) = p(y = 1 | \mathbf{x}, \boldsymbol{\theta}) = \frac{1}{1 + e^{-\mathbf{x}^T \boldsymbol{\theta}}}$$

*Suppose we were to use gradient descent*

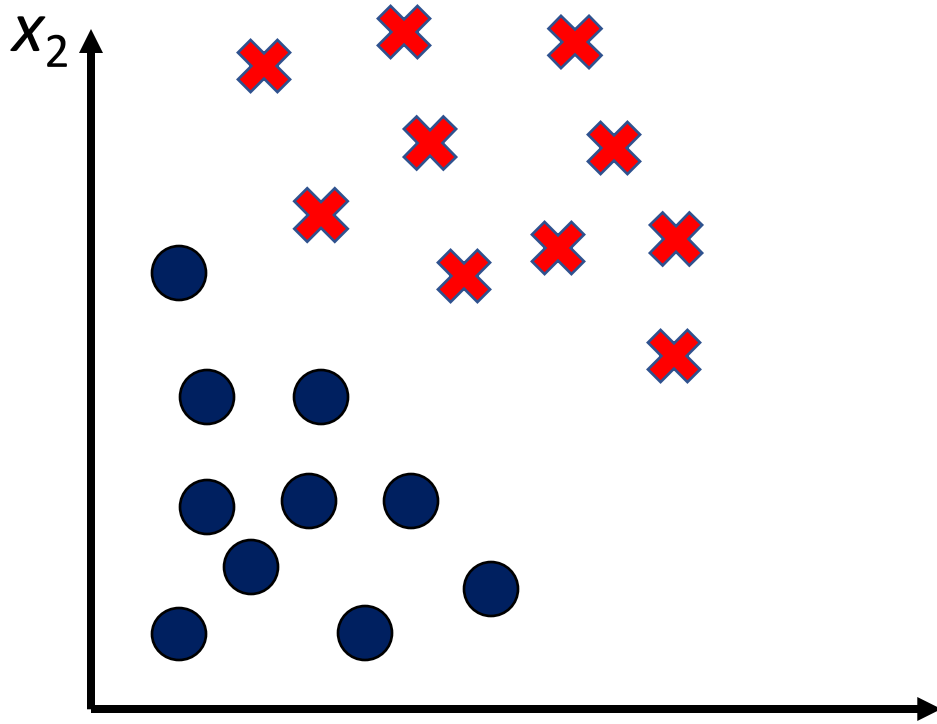
$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \gamma_i [\nabla_{\boldsymbol{\theta}} \mathcal{E}(\boldsymbol{\theta})]$$

$\vdots$

$$\boldsymbol{\theta}_j \leftarrow \boldsymbol{\theta}_j - \gamma_i \sum_{i=1}^n [f(\mathbf{x}_i) - y_i] (\mathbf{x}_i)_j$$

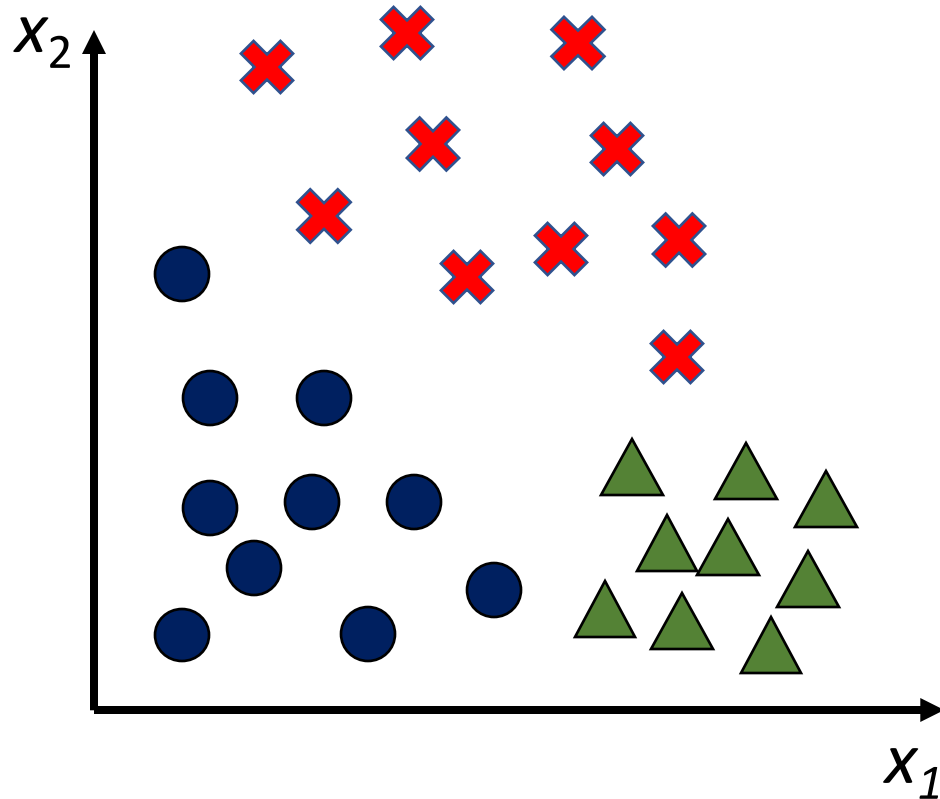
*This is the same result  
as for linear regression!*

# Application now to classification



- **Binary case:** we just need to find the optimal decision boundary that partitions these classes

# Application now to classification: One vs. All



- **Binary case:** we just need to find the optimal decision boundary that partitions these classes

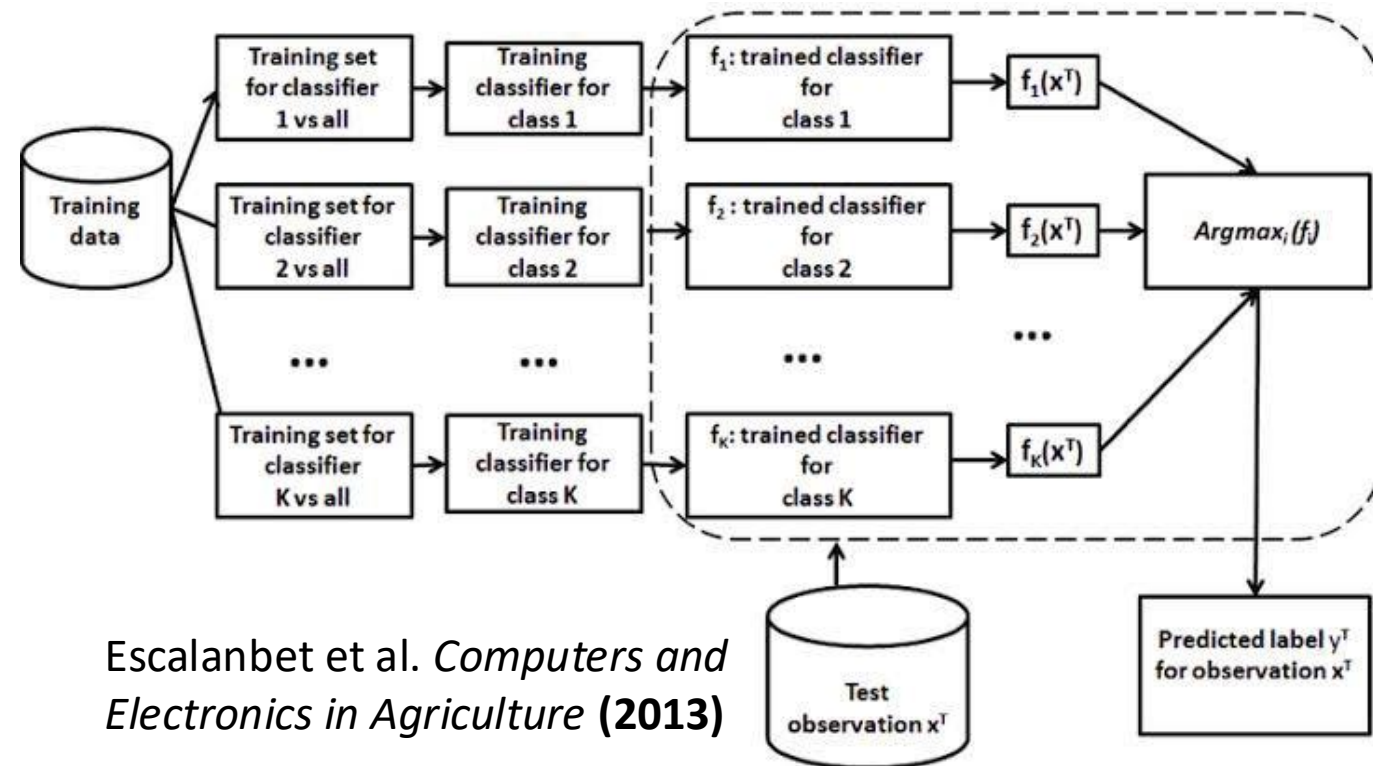
*what do we do for multiple classes?*

- **Multiclass case:** there are multiple strategies

Consider our data to have  $N$  classes...

## One vs. All (Rest)

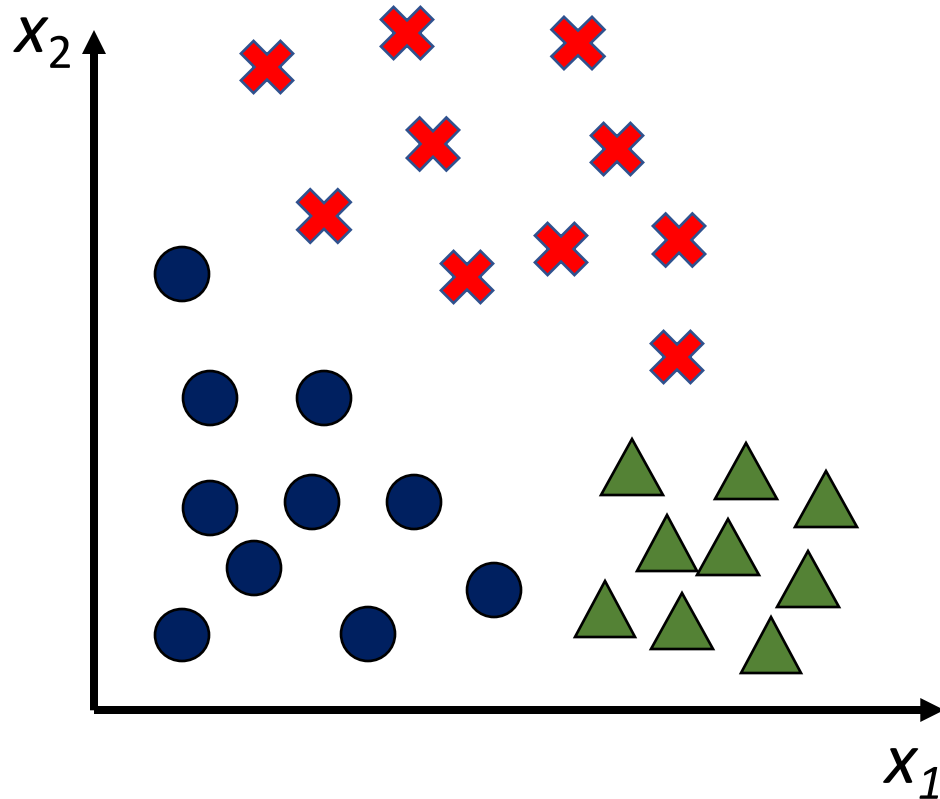
- formulate  $N$  binary classifier models



Escalanbet et al. *Computers and Electronics in Agriculture* (2013)

- pick the class that exhibits the highest score

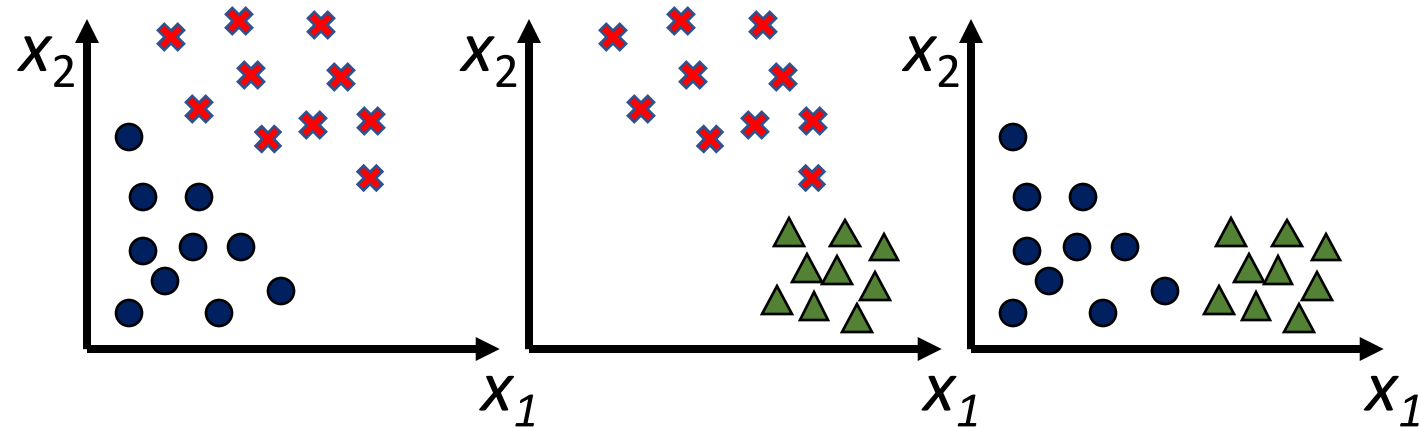
# Application now to classification: One vs. One



Consider our data to have  $N$  classes...

## One vs. One

- formulate  $N(N-1)/2$  binary classifier models



- pick the class that receives the most positive identifications

- **Binary case:** we just need to find the optimal decision boundary that partitions these classes

*what do we do for multiple classes?*

- **Multiclass case:** there are multiple strategies

# Logistic Regression Classification in scikit\_learn

```
1 from sklearn.datasets import make_classification
2 from sklearn.linear_model import LogisticRegression
3
4 # define problem
5 n      = 500 # number of data points
6 m      = 10  # size of feature vector
7 Nclass = 5   # number of classes
8 model_type = 'ovr' # ovr = one versus rest (examine other options)
9
10 # construct dataset
11 X_train, y_train = make_classification(n_samples=ndata,
12                                     n_features=m,
13                                     n_classes=Nclass,
14                                     n_redundant=m/2,
15                                     n_informative=m/2)
16
17 # define a model
18 myModel = LogisticRegression(multi_class='ovr')
19
20 # train the model
21 myModel.fit(X_train, y_train)
22
23 # check outcome of trained model
24 y_pred = model.predict(X_train)
```

# Support Vector Machine (SVM)

“Support-vector networks”, Corinna Cortes & Vladimir Vapnik, 1995

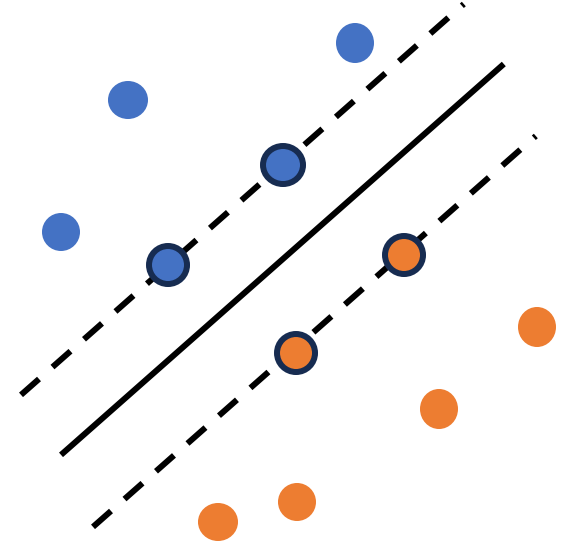


## Four Key SVM Ideas:

1. Margin
2. Support Vectors
3. Kernel
4. Regularization parameter

**Input:** 2D array (feature matrix)

**Output:** 1D array (vector of classes)



## When to use SVM?

1. **Number of features > number of samples:** gene expression, process system monitoring.
2. **Small to medium-size data.**
3. **Features are all numerical.**



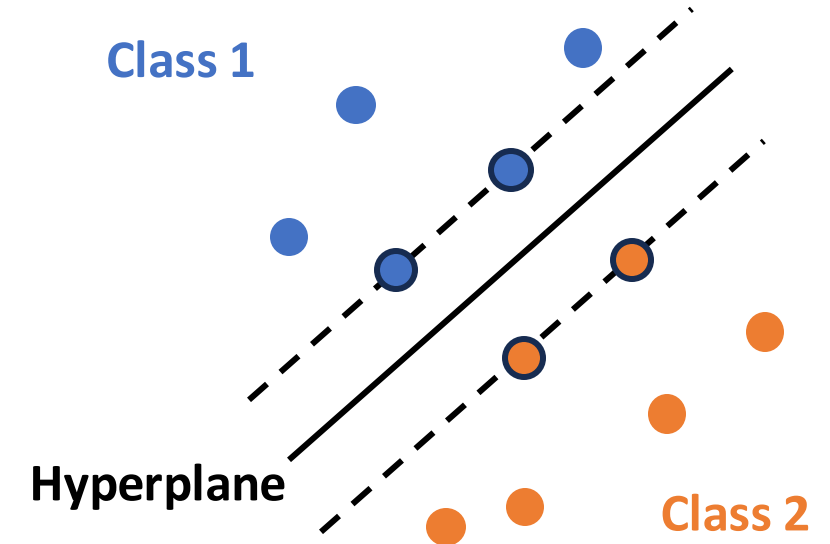
# Linear Separator Learning

**Definition:** Linear separator learning is a supervised learning task in which the goal is to find a **hyperplane** (or **a linear decision boundary**) in a feature space that can separate **two classes** of data points.

The **decision function** is

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

- $f(\mathbf{x}) > 0$  for all  $\mathbf{x}$  belong to class 1 (e.g., positive).
- $f(\mathbf{x}) < 0$  for all  $\mathbf{x}$  belong to class 2 (e.g., negative).
- $f(\mathbf{x}) = 0$  for all  $\mathbf{x}$  on the hyperplane.



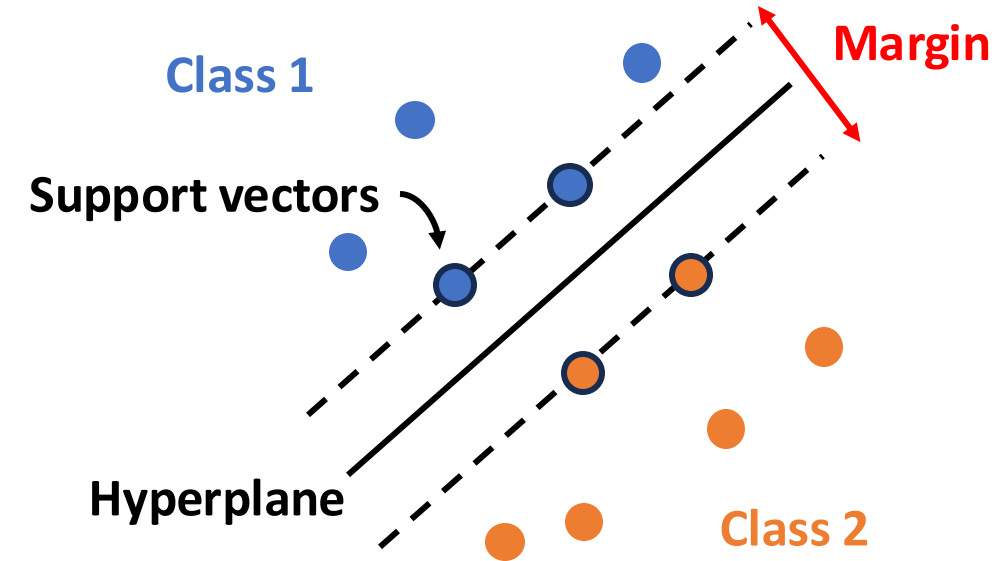
A data point  $(\mathbf{x}, y)$  is classified **correctly** if

$$y \cdot (\mathbf{w}^T \mathbf{x} + b) > 0$$

# Large Margin Classification

**Definition:** Large margin classification is a supervised learning method where the objective is to find a hyperplane that:

1. separates the classes of data, and
2. maximizes the margin between the hyperplane and the closest data points from either class.



## Why?

Given a data set that is linearly separable, there are **infinitely many hyperplanes** that can separate the two classes.

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) > 0, \quad i = 1, 2, \dots, m$$

$$\alpha \mathbf{w}, \alpha b \text{ for any } \alpha > 0$$

In SVM, we aim to find the hyperplane that maximizes the **margin**.

**Definition:** **Margin** is the distance between the hyperplane and the closest data points from either class.

**Definition:** **Support vectors** are the closest data points.

# Large Margin Classification

What is the geometric interpretation of  $w$ ?

$$\mathbf{w}^T (\mathbf{x}_2 - \mathbf{x}_1) = 0$$

$w$  is **orthogonal** to any vector that lies on the hyperplane.

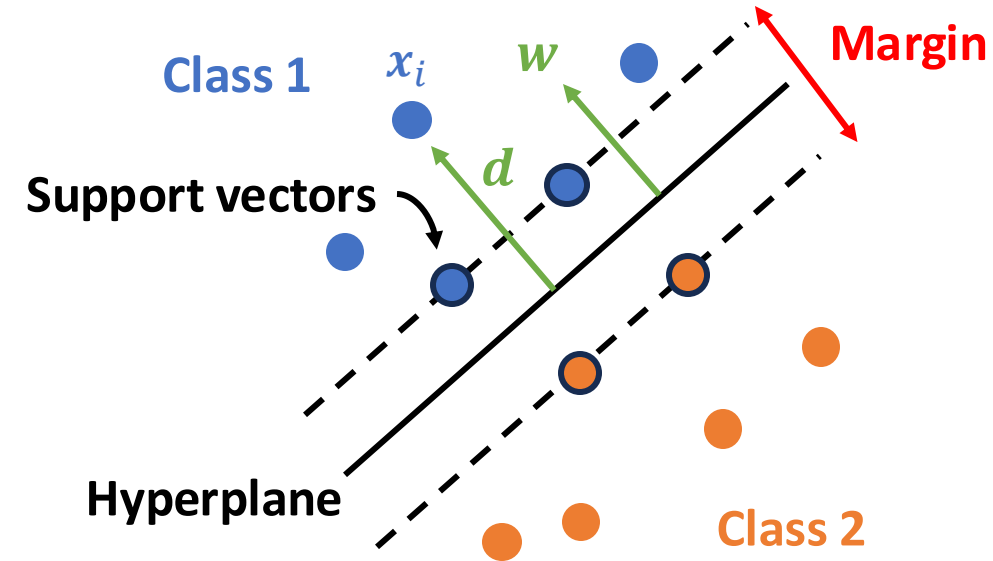
The distance  $d$  of a point  $x_i$  to the hyperplane is

$$\mathbf{w}^T (\mathbf{x}_i - \mathbf{d}) + b = 0$$

$$\mathbf{w}^T (\mathbf{x}_i - \alpha \mathbf{w}) + b = 0$$

$$\alpha = \frac{\mathbf{w}^T \mathbf{x}_i + b}{\mathbf{w}^T \mathbf{w}}$$

$$\|\mathbf{d}\|_2 = \sqrt{\alpha^2 \mathbf{w}^T \mathbf{w}} = \frac{|\mathbf{w}^T \mathbf{x}_i + b|}{\|\mathbf{w}\|_2}$$



The hyperplane must lie right in the middle of the two classes, with an equal distance from the positive and negative support vectors.

Because the hyperplane is **scale invariant**

$$\text{Margin} = \frac{2}{\|\mathbf{w}\|_2} \quad \text{Why?}$$

# The Hard-Margin SVM

**Definition:** Hard-margin SVM is a type of SVM used for binary classification when the data is linearly separable.

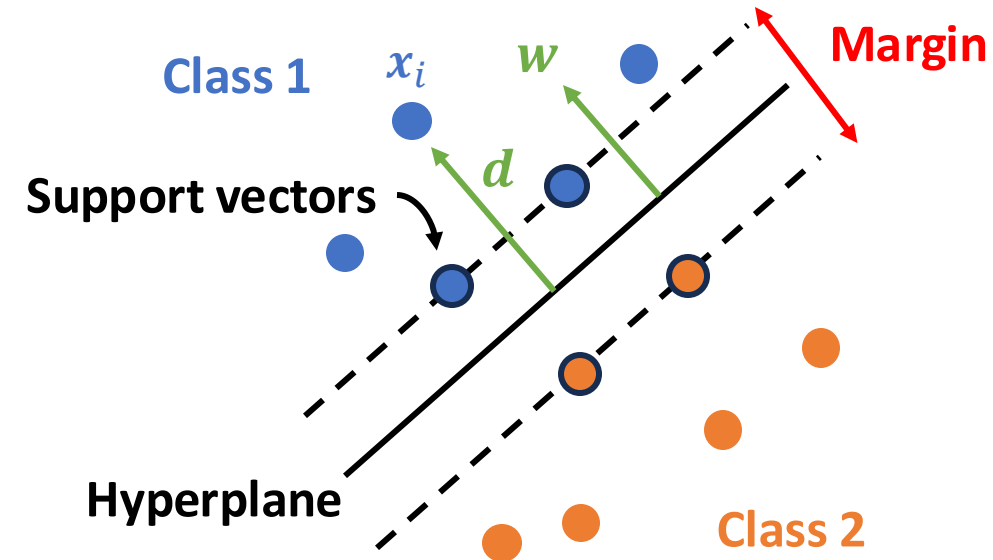
**Optimization:**

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|_2$$

subject to :

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1, \quad \forall i = 1, 2, \dots, m$$

- The constraint ensures that all points are correctly classified and lie on the correct side of the margin.
- This is a **convex quadratic programming problem**, which can be efficiently solved using standard optimization techniques.



**What if the data points are not linearly separable?**

# The Soft-Margin SVM

**Definition:** Soft-margin SVM is a type of SVM that allows some points to violate the margin constraints by introducing **slack variables** that permit misclassification or margin violations, while still attempting to maximize the margin.

**Optimization:**

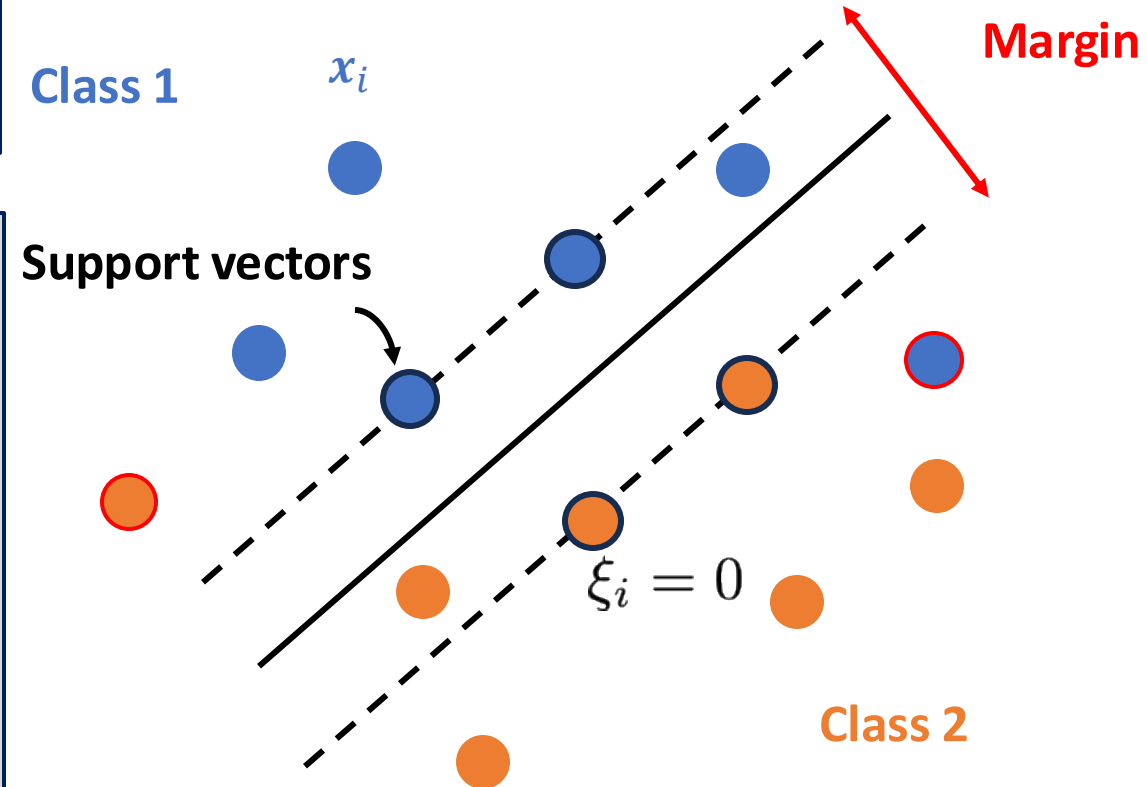
$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|_2 + C \sum_{i=1}^n \xi_i$$

subject to :

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \forall i = 1, 2, \dots, m$$

$$\xi_i \geq 0, \quad \forall i = 1, 2, \dots, m$$

- $\xi_i$ : slack variable.
- $C$ : regularization parameter.



# The Soft-Margin SVM

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_{i=1}^n \xi_i$$

subject to :

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, \forall i$$

$$\xi_i \geq 0, \forall i$$

**Class 1**

$\mathbf{x}_i$

**Support vectors**

$$\mathbf{w}^T \mathbf{x} + b = -1$$

$$\mathbf{w}^T \mathbf{x} + b = 0$$

$$\mathbf{w}^T \mathbf{x} + b = +1$$

$$\xi_i > 1$$

$$0 < \xi_i < 1$$

$$\xi_i = 0$$

**Class 2**

- Smaller  $C$ : Allows more margin violations but tries to maximize the margin.
- Larger  $C$ : A smaller margin but fewer violations



# The Soft-Margin SVM

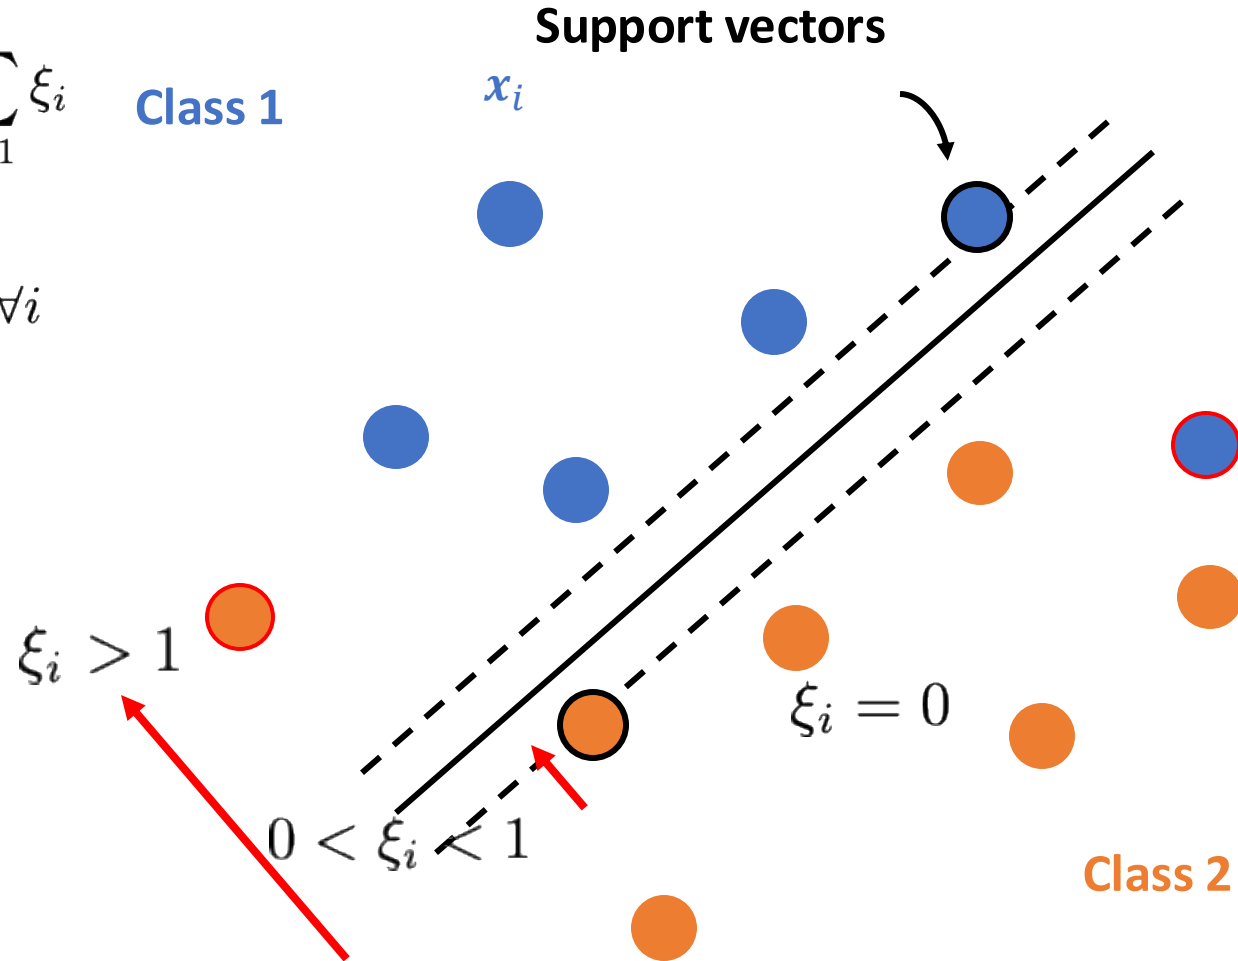
$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|_2 + C \sum_{i=1}^n \xi_i$$

Class 1

subject to :

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, \forall i$$

$$\xi_i \geq 0, \forall i$$



- Smaller  $C$ : Allows more margin violations but tries to maximize the margin.
- Larger  $C$ : A smaller margin but fewer violations

# Non-Linear SVM

**Definition: Non-linear SVM** extends linear SVM by using a kernel function to transform data into a higher-dimensional space, making it separable.

Non-linear SVM **hyperplane** is

$$\mathbf{w}^T \phi(\mathbf{x}) + b = 0$$

$$\phi : \mathbb{R}^n \rightarrow \mathbb{R}^d, \quad d > n$$

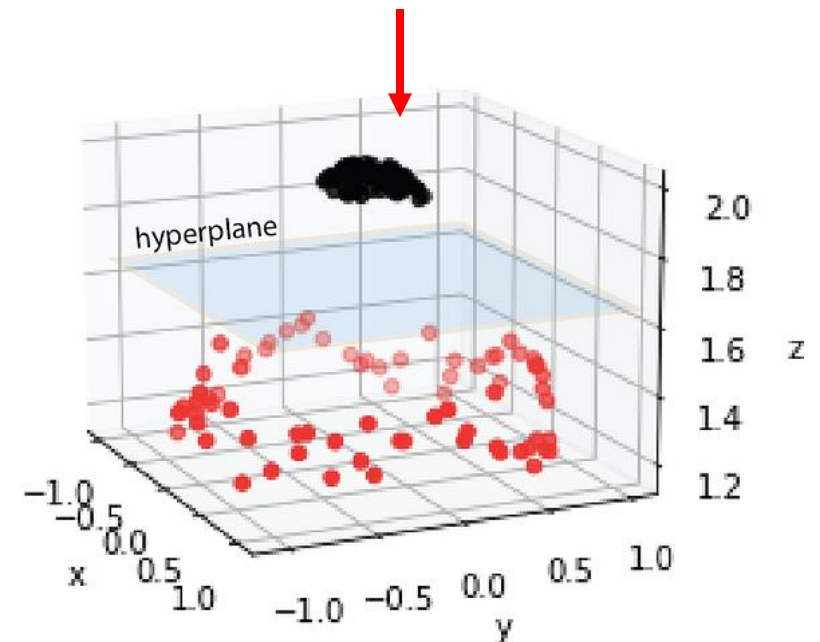
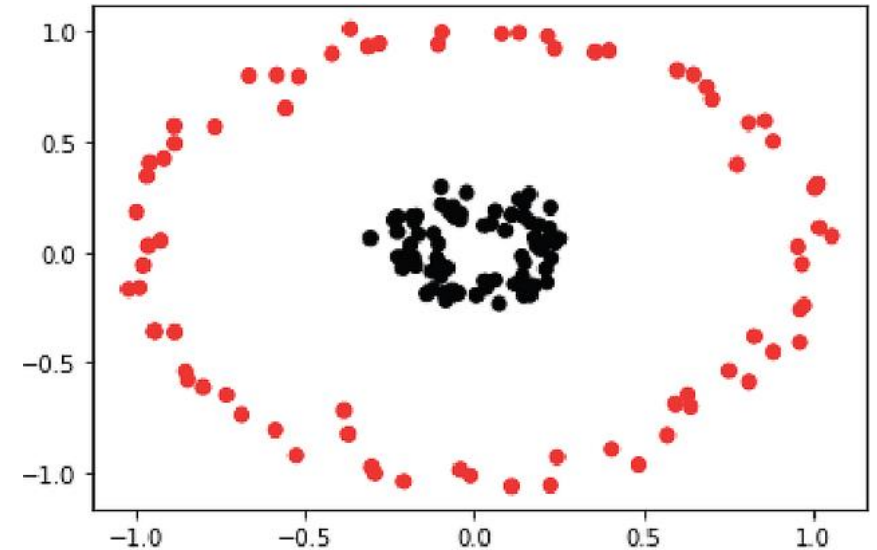
- $\phi(\mathbf{x})$  is a non-linear mapping function.

- Radial basis function (RBF):

$$k(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2)$$

- Polynomial:

$$k(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i^T \mathbf{x}_j + c)^d$$



# Kernel Trick

The **kernel trick** allows us to avoid explicitly computing  $\phi(\mathbf{x})$  while still achieving the benefits of mapping the data into a higher-dimensional space.

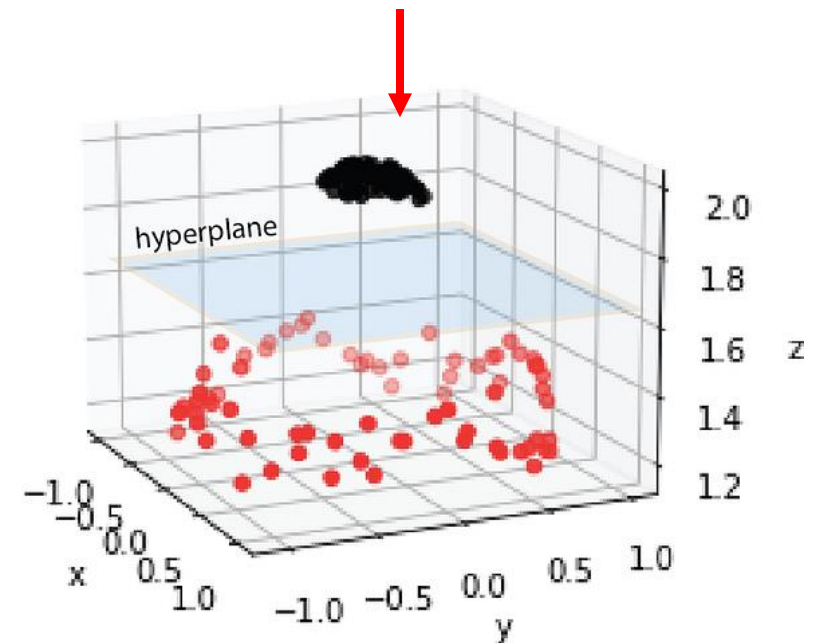
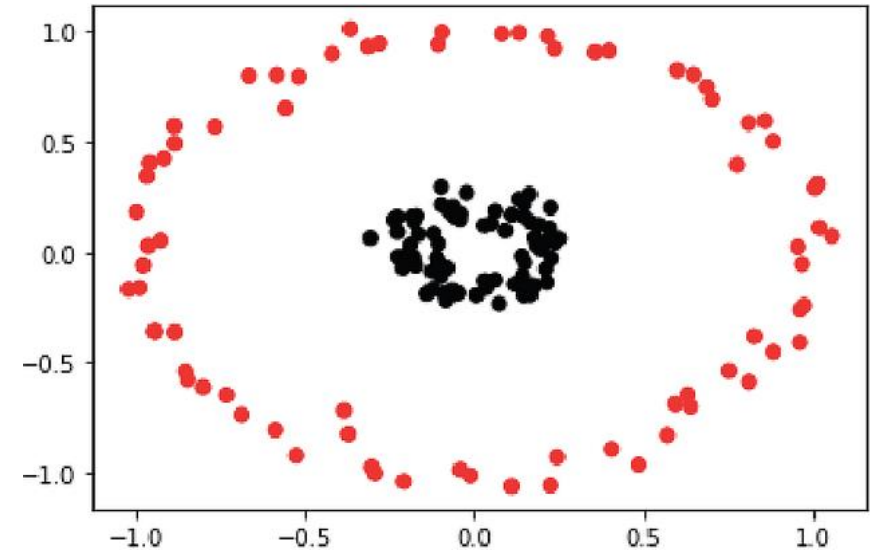
A kernel function  $K(\mathbf{x}_i, \mathbf{x}_j)$  computes the **dot product** of two transformed data points  $\phi(\mathbf{x}_i)$  and  $\phi(\mathbf{x}_j)$ :

$$k(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2)$$

↓ Taylor expansion

$$k(\mathbf{x}_i, \mathbf{x}_j) = C \left( \sum_{n=0}^{\infty} \frac{(\mathbf{x}_i^T \mathbf{x}_j)^n}{n!} \right)$$

- Infinite dimension dot product for RBF kernel.
- The kernel measures the **similarity or distance** between points, helping the SVM model understand how they relate to each other in the new space.



# SVM Classification in scikit\_learn

## sklearn.svm.SVC

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0,
shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None,
verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False,
random_state=None)
```

[\[source\]](#)

### Parameters:

**C : float, default=1.0**

Regularization parameter. The strength of the regularization is inversely proportional to C. Must be strictly positive. The penalty is a squared l2 penalty.

**kernel : {'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'} or callable, default='rbf'**

Specifies the kernel type to be used in the algorithm. If none is given, 'rbf' will be used. If a callable is given it is used to pre-compute the kernel matrix from data matrices; that matrix should be an array of shape (n\_samples, n\_samples). For an intuitive visualization of different kernel types see [Plot classification boundaries with different SVM Kernels](#).

# SVM Classification in scikit\_learn

## 1. Simple linear SVC and decision boundary

```
[38] # load the iris dataset
iris = datasets.load_iris()
X = iris.data[:100, :2] # take only the first two features
# sepal length and width
y = iris.target[:100]

# split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

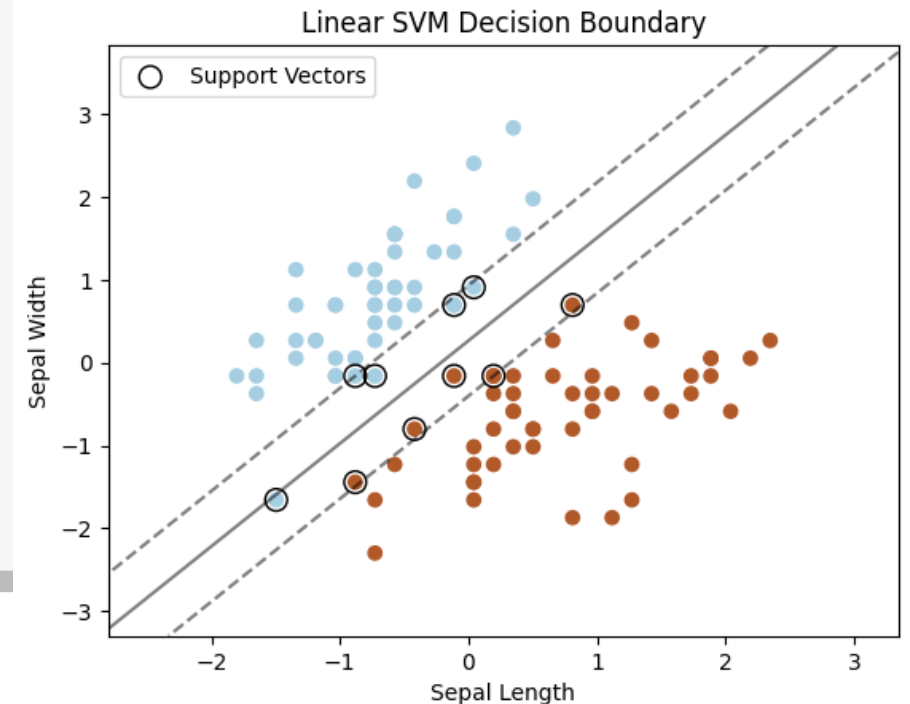
# standardize features by removing the mean and scaling to unit variance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# create and train a linear SVM model
clf = SVC(kernel="linear", C=1) # vary C, see impact
clf.fit(X_train, y_train)

# prediction on the test set
y_pred = clf.predict(X_test)
acc = accuracy_score(y_test, y_pred)

print(f"Accuracy: {acc*100:0.2f}%")
```

Accuracy: 100.00%



# SVM Classification in scikit\_learn

## 2. SVC with non-linear kernels

```
[41] # load the iris dataset
iris = datasets.load_iris()
X = iris.data[:100, :2] # take only the first two features
# sepal length and width
y = iris.target[:100]

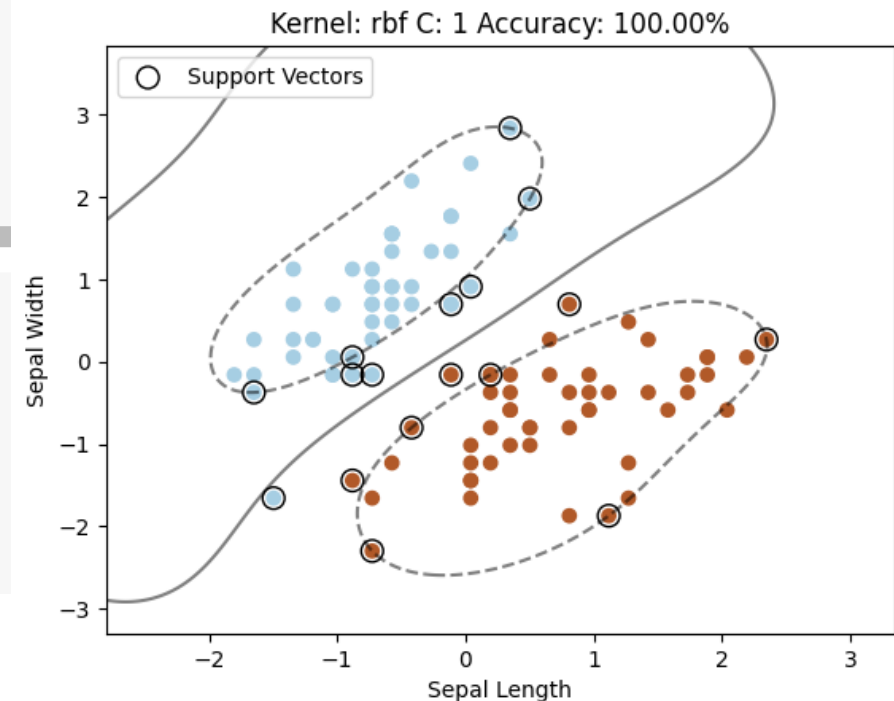
# split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# standardize features by removing the mean and scaling to unit variance
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

X_std = scaler.transform(X)
```

```
[43] # define a function to plot svc decision boundary
def plot_svc_boundary(X_train, y_train, X_test, y_test, kernel, C=1, degree=3):
    clf = SVC(kernel=kernel, C=C, degree=degree)
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
```





# Hyperparameter Tuning

**Definition:** Grid search is an **exhaustive** search algorithm used to find the optimal hyperparameters for a given model by evaluating all possible combinations of a predefined set of hyperparameter values.

`sklearn.svm.SVC`

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0,
shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None,
verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False,
random_state=None)
```

[\[source\]](#)

Validation Accuracy		C			
		0.1	1	10	100
kernel	linear	0.5	0.6	0.7	0.6
	rbf	0.6	0.7	0.8	0.7
	sigmoid	0.7	0.8	0.9	0.8
	poly	0.6	0.7	0.8	0.7

Training

Validation

Test

Grid Search

Best Model

# Grid Search in scikit\_learn

## `sklearn.model_selection.GridSearchCV`

```
class sklearn.model_selection.GridSearchCV(estimator, param_grid, *, scoring=None, n_jobs=None, refit=True, cv=None, verbose=0, pre_dispatch='2*n_jobs', error_score=nan, return_train_score=False)
```

[\[source\]](#)

### Parameters:

#### **estimator : estimator object**

This is assumed to implement the scikit-learn estimator interface. Either estimator needs to provide a `score` function, or `scoring` must be passed.

#### **param\_grid : dict or list of dictionaries**

Dictionary with parameters names (`str`) as keys and lists of parameter settings to try as values, or a list of such dictionaries, in which case the grids spanned by each dictionary in the list are explored. This enables searching over any sequence of parameter settings.

#### **scoring : str, callable, list, tuple or dict, default=None**

Strategy to evaluate the performance of the cross-validated model on the test set.

If `scoring` represents a single score, one can use:

- a single string (see [The scoring parameter: defining model evaluation rules](#));
- a callable (see [Defining your scoring strategy from metric functions](#)) that returns a single value.

If `scoring` represents multiple scores, one can use:

- a list or tuple of unique strings;
- a callable returning a dictionary where the keys are the metric names and the values are the metric scores;
- a dictionary with metric names as keys and callables as values.

See [Specifying multiple metrics for evaluation](#) for an example.

# Grid Search in scikit\_learn

## 4. SVC hyperparameter tuning

```
[70] iris = datasets.load_iris()
X = iris.data[:, :2]
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
[74] # define the model
model = SVC()

# define the hyperparameters to tune
param_grid = {
    'C': [0.1, 1, 10, 100],
    'kernel': ['linear', 'rbf', 'poly', 'sigmoid']
}

# set up GridSearchCV with 5-fold cross-validation
# if you have a very large search space, you can use RandomizedSearchCV
grid_search = GridSearchCV(model, param_grid, cv=5,
                           scoring='accuracy', return_train_score=True)

# fit the model
grid_search.fit(X_train, y_train)

# get the best hyperparameters
best_params = grid_search.best_params_

# get the validation accuracies for all combinations of hyperparameters
cv_results = pd.DataFrame(grid_search.cv_results_)
```

Validation Accuracies for Hyperparameter Combinations:

	param_C	param_kernel	mean_val_score
0	0.1	linear	0.775000
1	0.1	rbf	0.716667
2	0.1	poly	0.650000
3	0.1	sigmoid	0.783333
4	1	linear	0.766667
5	1	rbf	0.766667
6	1	poly	0.675000
7	1	sigmoid	0.783333
8	10	linear	0.775000
9	10	rbf	0.741667
10	10	poly	0.716667
11	10	sigmoid	0.691667
12	100	linear	0.775000
13	100	rbf	0.725000
14	100	poly	0.725000
15	100	sigmoid	0.675000

Best Hyperparameters: {'C': 0.1, 'kernel': 'sigmoid'}

Val Accuracy with Best Hyperparameters: 0.9333333333333333