


The background of the slide is a complex 3D visualization. It features a dense cloud of small, multi-colored spheres (red, blue, green, yellow, purple) scattered throughout a 3D coordinate system. A grid of lines is visible, suggesting a high-dimensional space. In the foreground, there are several grey, rectangular blocks or cubes arranged in a structured manner. A large, dark grey arrow points from the right side of the image towards the central text. The overall aesthetic is technical and data-driven.

Dimensionality Reduction: *learning to ignore things*

Why should we use dimensionality reduction?



*Do we need all of these features to accurately predict the output?
Can we get away with only using a subset of the feature space?*

Considerations

- Model trains faster since it has fewer dimensions.
- Make the model simpler for researchers to interpret/visualize.
- Improve model accuracy due to less misleading or noisy data.

Types of Dimensionality Reduction Techniques


Most DR techniques fall into a few categories

- **Feature Elimination and Extraction**

*we have been working with this already a bit
in general, the goal is to systematically remove
uninformative/redundant variables*


- **Linear Transformations/Components**

*here the strategy is to find a linear
transformation of your feature space that
results in more “useful” set of coordinates*



- **Non-linear Transformations & manifolds**

*in this case, we usually perform some non-
linear transformation and project the features
onto a lower-dimensional manifold*



Not linearly separable!

Feature Elimination and Extraction

From a feature set, we are looking for a subset of features to use:

- Remove features with too many missing values (**Missing Value Ratio**)
- Remove features that exhibit small variance (**Low-variance Filter**)
- Remove highly correlated features (e.g., using Pearson's r) (**High Correlation filter**)
- Assess **feature importance** to model predictions
 - by using e.g., **Random Forest, SHAP**
 - by systematically removing features (**Backward Feature Elimination**)
 - by systematically adding features (**Forward Feature Selection**)

These are all very simple, powerful strategies that may be complementary to other methods → always worth considering!


Principal Component Analysis (PCA)

PCA is the “canonical” and most common DR technique


Important features of PCA

- *PCA assumes linear relationships between variables.*
- *PCA is scale dependent (features with larger values look more important)*
- *PCA looks at variance in the feature data (it can be important to preprocess the data via normalization, mean-centering, or scaling).*
- *PCA is one of the central applications of SVD*
- *PCA can be used for preprocessing for other DR techniques*

original data space



PCA



The basic goal:

$$\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_p; \mathbf{x}_i \in \mathbb{R}^d$$

$$\downarrow \mathbf{W}$$

$$\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_p; \mathbf{y}_i \in \mathbb{R}^k$$

$$k \ll d$$


Principal Component Analysis (PCA)

The big idea:

1. We want to identify an orthonormal basis to represent $\mathbf{x}_i - \boldsymbol{\mu} \rightarrow d$ -dimensional basis
2. We will only pick k of the d basis vectors
3. All of the points will then be represented in a k -dimensional subspace spanned by $\mathbf{w}_1, \dots, \mathbf{w}_k$

$$\mathbf{x}_i - \boldsymbol{\mu} = \sum_{j=1}^d \mathbf{y}_i[j] \mathbf{w}_j \quad \text{exact reconstruction}$$

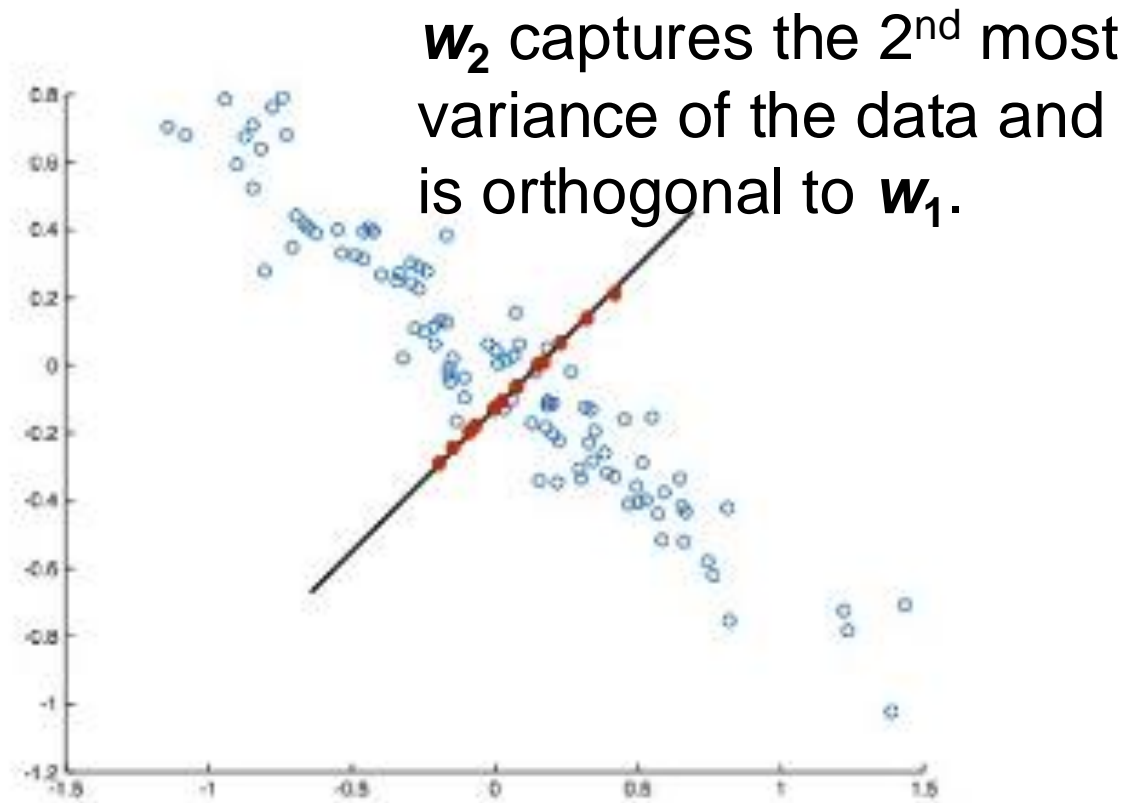
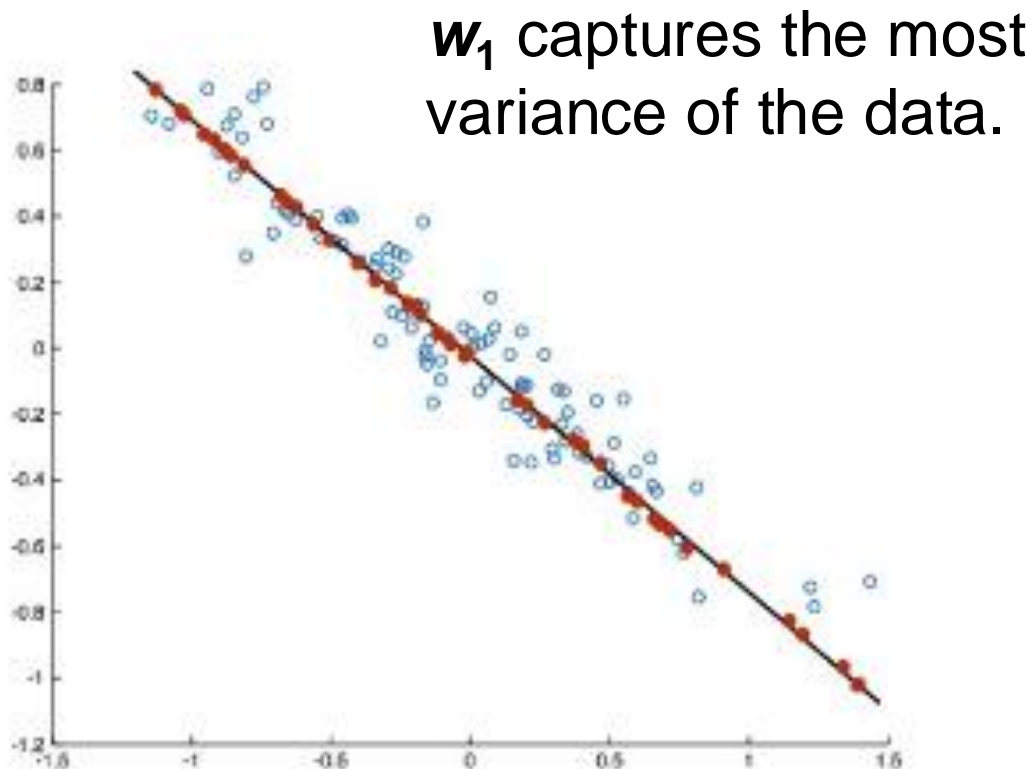
$$\hat{\mathbf{x}}_i - \boldsymbol{\mu} = \sum_{j=1}^k \mathbf{y}_i[j] \mathbf{w}_j \quad \text{approximate reconstruction}$$



Algorithm \rightarrow how to choose \mathbf{W} and which to keep

PCA: How to pick basis vectors

Idea: Find the directions along which the data is maximally spread (highest variance).



PCA: How to pick basis vectors

Idea: Find the directions along which the data is maximally spread (highest variance).

$$\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_p] \quad \mu = \left[\frac{1}{p} \sum_{i=1}^p \mathbf{x}_i[1], \dots, \frac{1}{p} \sum_{i=1}^p \mathbf{x}_i[d] \right]^T$$

$$\mathbf{B} = \mathbf{X} - \mu \mathbf{1}_d^T$$

- mean-subtracted data

$$\mathbf{C} = \frac{1}{n-1} \mathbf{B}^T \mathbf{B}$$

- covariance matrix

$$\mathbf{w}_1 = \arg \max_{\|\mathbf{w}_1\|=1} \mathbf{w}_1^T \mathbf{B}^T \mathbf{B} \mathbf{w}_1$$

- first principal component

subsequent components can be determined by orthogonalization and repetition

PCA: How to pick basis vectors

Idea: Find the directions along which the data is maximally spread (highest variance).

$$\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_p] \quad \mu = \left[\frac{1}{p} \sum_{i=1}^p \mathbf{x}_i[1], \dots, \frac{1}{p} \sum_{i=1}^p \mathbf{x}_i[d] \right]^T$$

$$\mathbf{B} = \mathbf{X} - \mu \mathbf{1}_d^T$$

- mean-subtracted data

$$\mathbf{C} = \frac{1}{n-1} \mathbf{B}^T \mathbf{B}$$

- covariance matrix

$$\mathbf{X} = \mathbf{U} \mathbf{\Sigma} \mathbf{W}^T$$

Or you can get them all in one go by doing the SVD or some eigendecomposition

PCA: additional details/realizations

- We want to maximize this function subject to the constraint that the norm of \mathbf{w} is 1. We can do this with Lagrange multipliers

$$w_1 = \arg \max w^T \Sigma w - \lambda \|\mathbf{w}\|_2^2$$

- We now take the derivative of this with respect to \mathbf{w} equate to 0, and arrive at

$$\Sigma \mathbf{w}_1 - \lambda \mathbf{w}_1 = 0$$

- This is an eigenvalue equation! The direction \mathbf{w}_1 that we obtain by maximizing the variance in the direction is some unit vector that satisfies this eigenvalue equation. We want to maximize $\mathbf{w}^T \Sigma \mathbf{w}$, so if we plug in the eigenvalue solution we get that

$$\mathbf{w}^T \Sigma \mathbf{w} = \lambda \|\mathbf{w}_1\|_2^2 = \lambda$$

So the eigenvector \mathbf{w}_1 that maximizes the variance is the one with the largest eigenvalue!

- We can proceed in similar fashion to pick eigenvectors with the second, third, etc largest eigenvalues to form this basis.

Principal components analysis then amounts to simply finding the eigenvectors of the covariance matrix Σ !

PCA: basic Python implementation

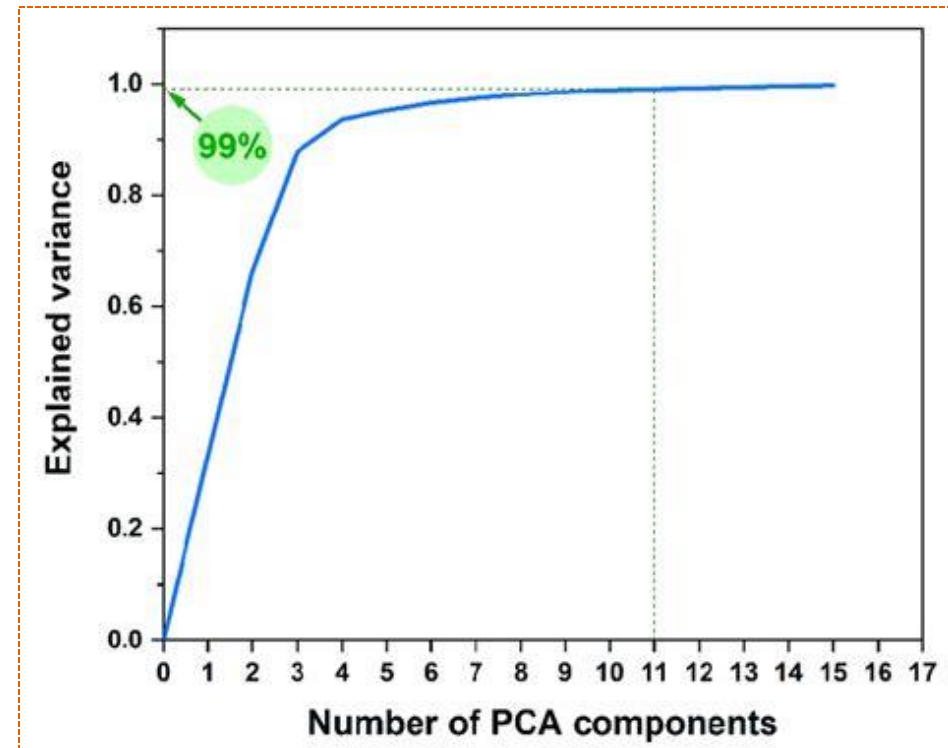
```
def pca(x,n_components):  
    #Vanilla PCA  
    #subtract off the mean of the data  
    x_mean = x - np.mean(x, axis=0)  
  
    #Calculate the covariance matrix  
    #Numpy as this calculation automated  
    cov_mat = np.cov(x_mean, rowvar = False)  
  
    #Solve for eigensystem of covariance matrix  
    eval,evalvec = eigh(cov_mat)  
  
    #Sort eigenvalues and eigenvectors  
    sort_i = np.argsort(eval)[::-1]  
    sort_eval = eval[sort_i]  
    sort_evalvec = evalvec[:,sort_i]  
  
    #Now take the first "n_components" principal components  
    eval_subset = sort_eval[0:n_components]  
    evalvec_subset = sort_evalvec[:,0:n_components]  
  
    #Project dataset onto principal components  
    x_transform = np.dot(evalvec_subset.transpose(), x_mean.transpose()).transpose()  
  
    return x_transform, evalvec_subset, eval_subset
```

PCA: how much of the data is explained?

It is natural to ask how much of the information in a data set is lost by projecting the observations onto a small subset of principal components.

The amount of variance explained can be analytically related to the eigenvalues of the covariance matrix via:

$$\text{Explained Variance Ratio} = \frac{\lambda_j}{\sum_{j=1}^d \lambda_j}$$



PCA: Application to Visualizing Chemical Space

Open Macromolecular Genome: Generative Design of Synthetically Accessible Polymers

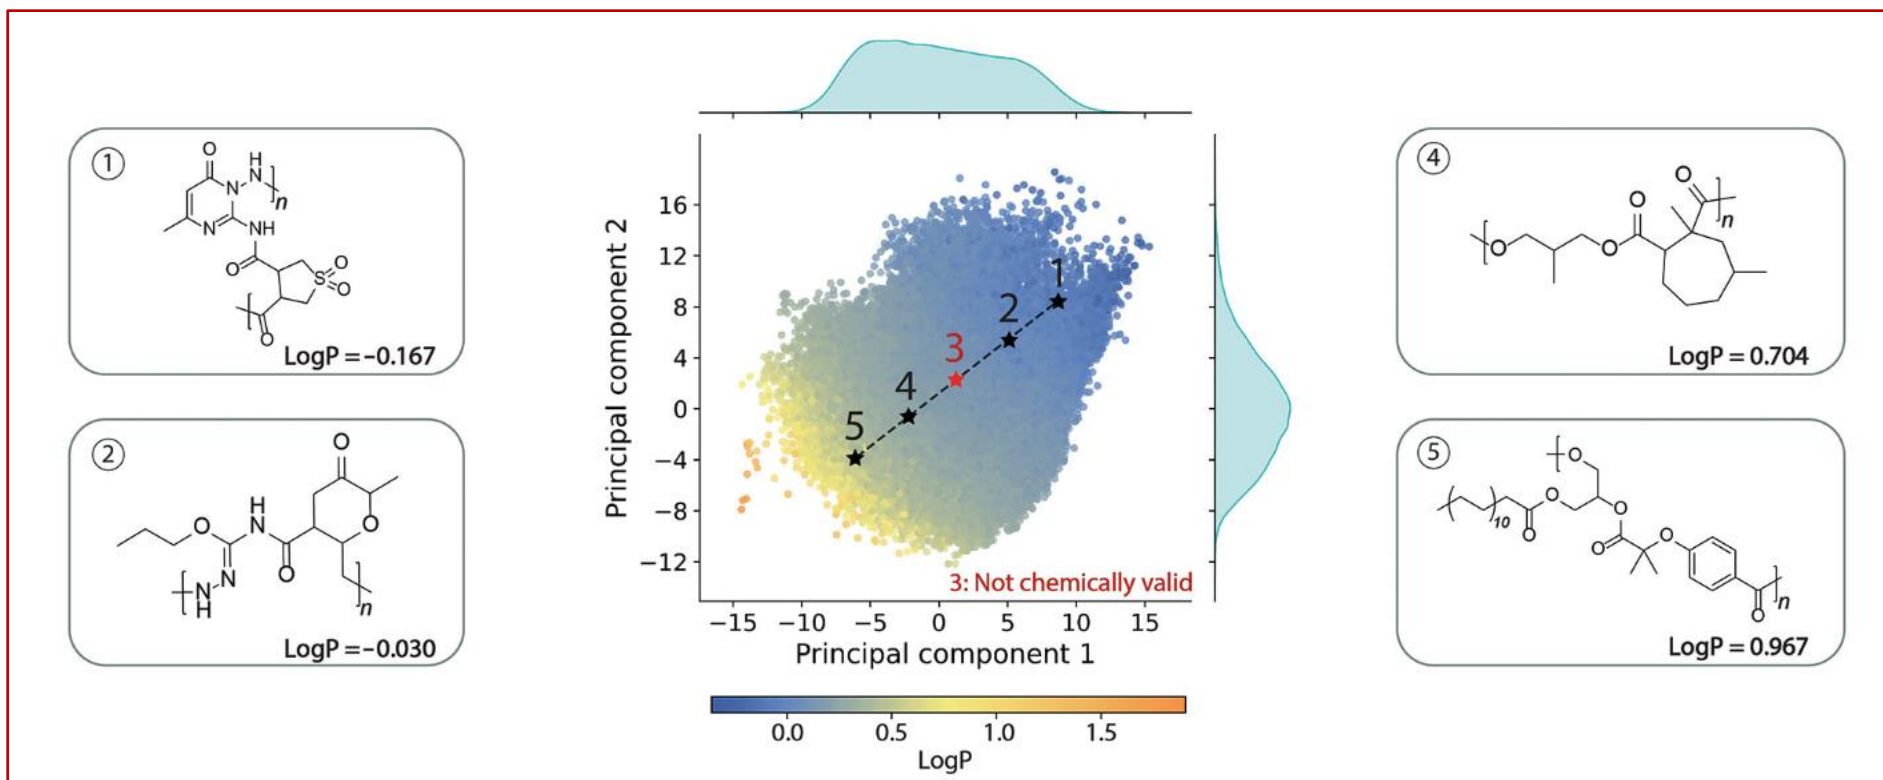
Seonghwan Kim, Charles M. Schroeder, and Nicholas E. Jackson*



Cite This: *ACS Polym. Au* 2023, 3, 318–330




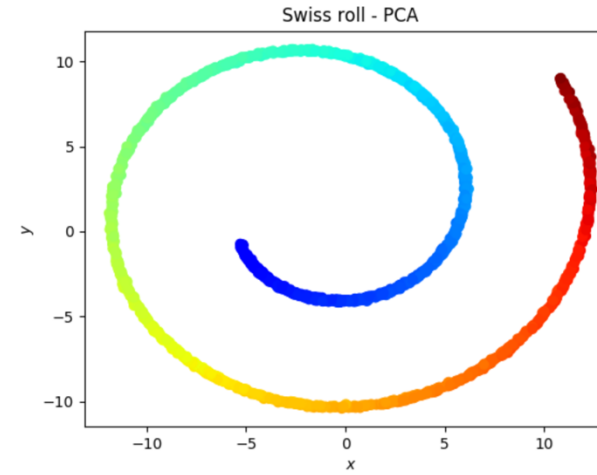

Read Online



Correcting for failures of PCA

PCA generally fails to detect low-dimensional manifolds


Swiss roll



Tenenbaum, J. B.; de Silva, V. & Langford, J. C. A Global Geometric Framework for Nonlinear Dimensionality Reduction *Science*, **2000**, 290, 2319-2323

A (relatively) Simple Modification: Kernel PCA

Consider the following...



What happens if we
do PCA here?


The Kernel Trick

- effectively project the data into a *higher*-dimensional space
- this usually enables some further resolution of data
- the kernel trick provides an efficient approach to the data transformation
- specifically, it allows for computations of distances in the higher-dimensional feature space but never actually performs the coordinate transformation

Let $\phi : \mathcal{V}_{\text{old}} \rightarrow \mathcal{V}_{\text{new}}$


Then, $k(\mathbf{x}, \mathbf{y}) := \langle \phi(\mathbf{x}), \phi(\mathbf{y}) \rangle_{\mathcal{V}_{\text{new}}}$

this is a non-trivial but
somewhat arbitrary function



A (relatively) Simple Modification: Kernel PCA

Consider the following...



What happens if we
do PCA here?

The Kernel Trick

- effectively project the data into a *higher*-dimensional space
- this usually enables some further resolution of data
- the kernel trick provides an efficient approach to the data transformation
- specifically, it allows for computations of distances in the higher-dimensional feature space but never actually performs the coordinate transformation

Common kernels:


$$k(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^T \mathbf{y} + 1)^d$$

$$k(\mathbf{x}, \mathbf{y}) = e^{-\gamma ||\mathbf{x} - \mathbf{y}||^2}$$

Note: Care must be taken to use zero-mean data in KPCA

A (relatively) Simple Modification: Kernel PCA

Consider the following...



**What happens if we
do PCA here?**

Algorithm to implement Kernel PCA:

1. Compute the Kernel matrix, K .

$$K = \begin{bmatrix} \kappa(x^{(1)}, x^{(1)}) & \cdots & \kappa(x^{(1)}, x^{(n)}) \\ \vdots & \ddots & \vdots \\ \kappa(x^{(n)}, x^{(1)}) & \cdots & \kappa(x^{(n)}, x^{(n)}) \end{bmatrix}$$

2. Center the kernel matrix (like normal feature scaling, but in kernel space) using

$$K' = K - 1_n K - K 1_n + 1_n K 1_n$$

$1_n = n \times n$ matrix with all elements equal to $1/n$.

3. Solve $K'v = \lambda v$ and extract the top k eigenvectors.

Other Greatest Hits of Dimensionality Reduction

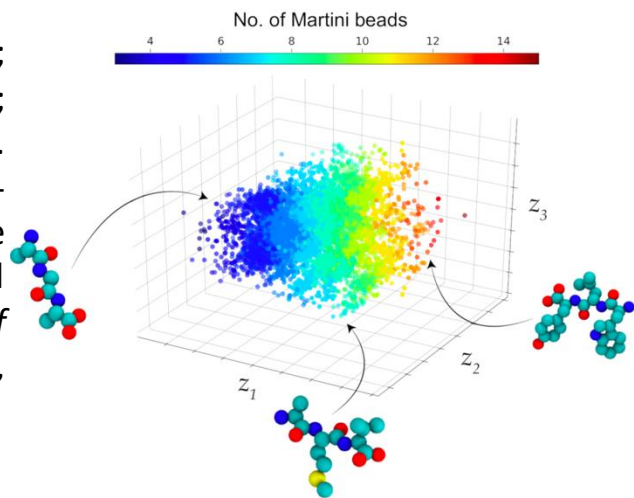
- Linear Discriminant Analysis
- Generalized Discrimination Analysis
- **UMAP**
- Non-negative Matrix Factorization
- Classical Scaling
- Maximum Variance Unfolding
- **Diffusion Maps**
- Locally Linear Embedding
- **t-distributed stochastic neighbor embedding (t-SNE)**
- Laplacian Eigenmaps
- Hessian LLE
- Local Tangent Space Analysis
- Sammon Mapping
- Multilayer Autoencoders
- Locally Linear Coordination
- Manifold Charting
- **ISOMAP**
- Kernel PCA

There are many more beyond this!

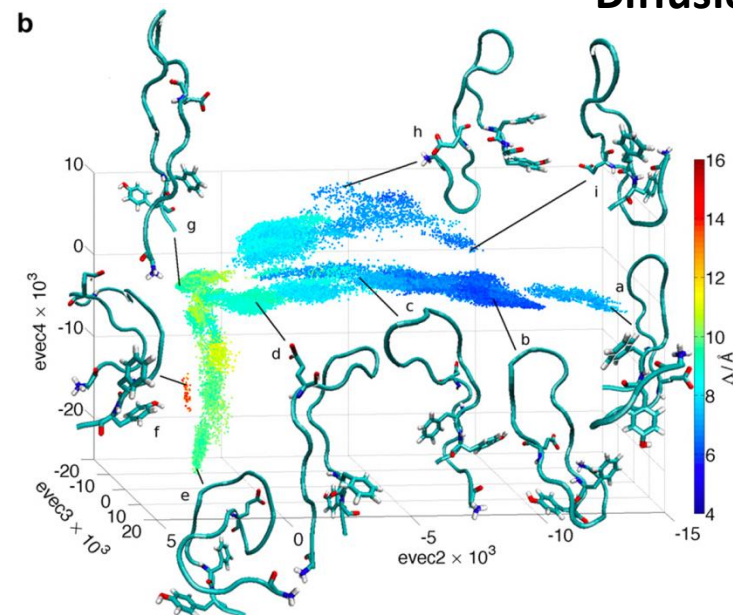
Other Greatest Hits of Dimensionality Reduction

Variational Autoencoders

Shmilovich, K.; Mansbach, R. A.; Sidky, H.; Dunne, O. E.; Panda, S. S.; Tovar, J. D. & Ferguson, A. L. Discovery of Self-Assembling pi-Conjugated Peptides by Active Learning-Directed Coarse-Grained Molecular Simulation. *The Journal of Physical Chemistry B* **2020**, 124, 3873-3891

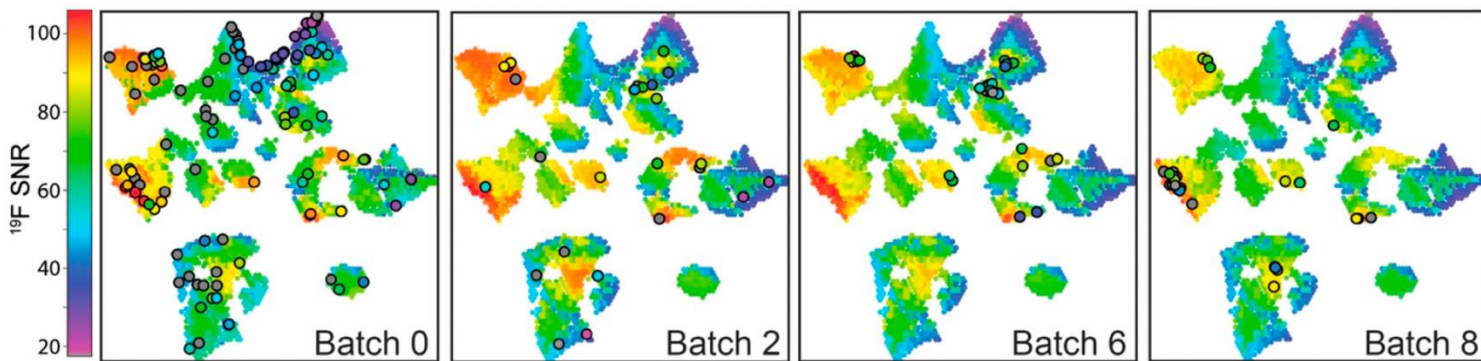


Diffusion Maps



Ferguson, A. L.; Panagiotopoulos, A. Z.; Kevrekidis, I. G. & Debenedetti, P. G. Nonlinear dimensionality reduction in molecular simulation: The diffusion map approach *Chemical Physics Letters*, **2011**, 509, 1-11

UMAP




Reis, M.; Gusev, F.; Taylor, N. G.; Chung, S. H.; Verber, M. D.; Lee, Y. Z.; Isayev, O. & Leibfarth, F. A. Machine-Learning-Guided Discovery of ¹⁹F MRI Agents Enabled by Automated Copolymer Synthesis *Journal of the American Chemical Society*, **2021**, 143, 17677-17689

Descriptor

t-SNE


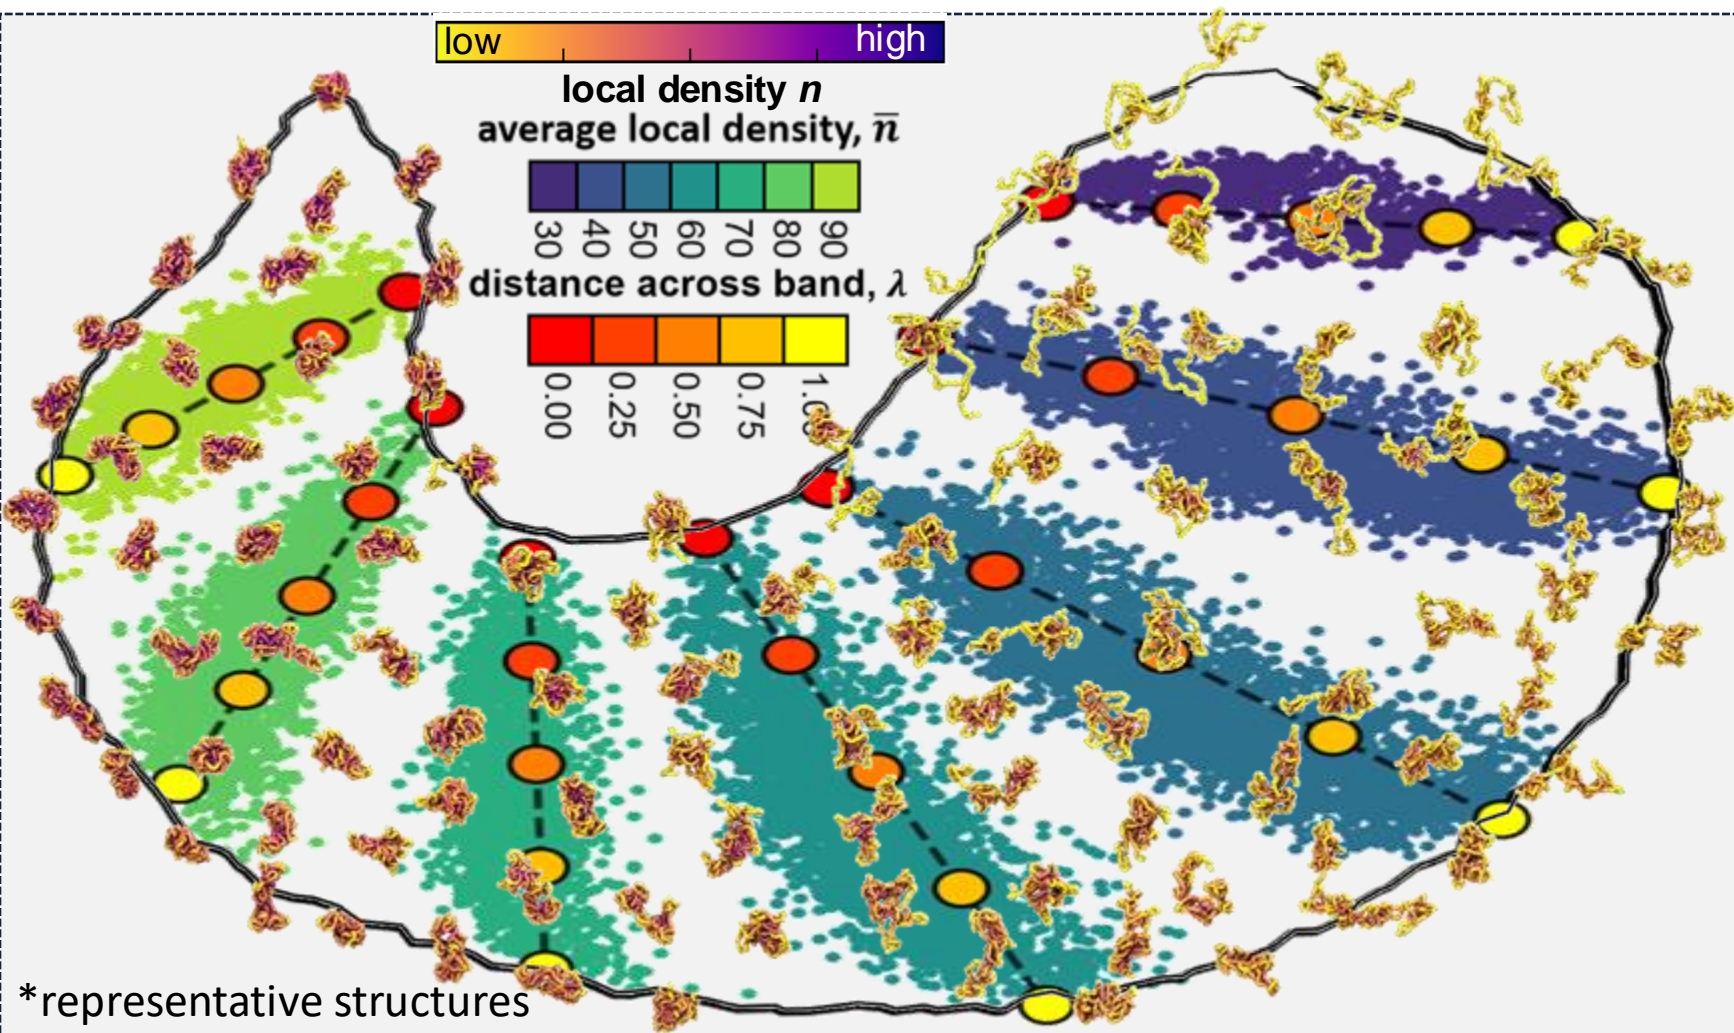
Fingerprint



Tao, L.; Chen, G. & Li, Y. Machine learning discovery of high-temperature polymers *Patterns*, **2021**, 2, 100225

Other Greatest Hits of Dimensionality Reduction

- 1) Structures are **primarily** differentiated by the **average** local bead density
- 2) Structures are **secondarily** differentiated by the **variance** in local bead density



Patel, Colmenares, Webb. ACS
Polymers Au. 2023

Standard DR Techniques in Scikit-learn

Sklearn.decomposition

<https://scikit-learn.org/stable/modules/decomposition.html>

- Kernel PCA
- Independent Components Analysis
- Linear Factor Analysis
- Non-negative Matrix Factorization
- Truncated SVD

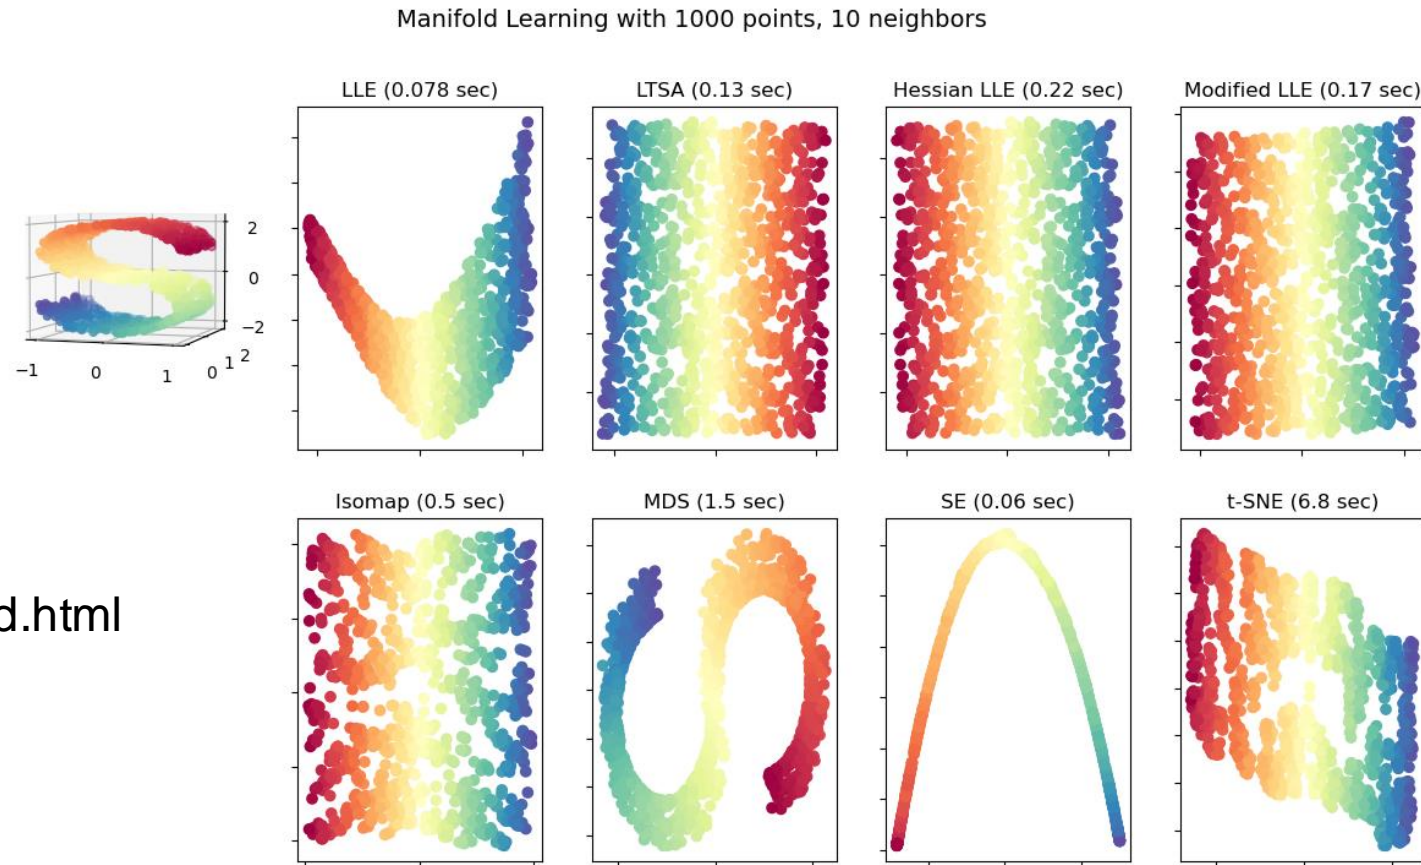
Sklearn.discriminant_analysis

- Linear Discriminant Analysis


Sklearn.manifold

<https://scikit-learn.org/stable/modules/manifold.html>




- Isomap
- Locally Linear Embedding
- Spectral Embedding
- TSNE
- Multidimensional Scaling



Example on Alanine Dipeptide



1. Perform a molecular dynamics simulation (300 K)
2. Get coordinates of atoms as a function of time
3. From these coordinates compute the distance matrix amongst all atoms as a function of time
4. Use DR techniques to identify collective variables that provide a simplified description of the molecular motions



```
#Now we can normalize all of the distances in the feature space by
subtracting the mean and dividing by std of each feature.
scaler = StandardScaler()
tdflat_norm = scaler.fit_transform(tdflat)
```

```
#Then we can do a PCA and keep the two largest components
pca = PCA(n_components=2)
pca.fit(tdflat_norm)
```

```
#Apply the PCA transformation matrix to the data
tdflat_trans = pca.transform(tdflat_norm)
```

```
#Do the same thing using Isomap
Isomap_embedding = Isomap(n_components=2)
tdflat_trans_isomap = Isomap_embedding.fit_transform(tdflat_norm)
```

```
#Do the same thing using t-SNE
TSNE_embedding = TSNE(n_components=2)
tdflat_trans_tsne = TSNE_embedding.fit_transform(tdflat_norm)
```

```
import os
import sys
import numpy as np
import math
from scipy.spatial.distance import cdist
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
from sklearn.manifold import Isomap, TSNE
```


t-Distributed Stochastic Neighborhood Embedding (t-SNE)

Given a set of high dimensional data points $\mathbf{x}_1, \dots, \mathbf{x}_n$ t-SNE computes the conditional probabilities, $p_{j|i}$, that are proportional to the similarity of data points \mathbf{x}_i and \mathbf{x}_j

For $i \neq j$, define

$$p_{j|i} = \frac{\exp(-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|\mathbf{x}_i - \mathbf{x}_k\|^2 / 2\sigma_i^2)}$$

and set $p_{i|i} = 0$. Note that $\sum_j p_{j|i} = 1$ for all i .



Why can we call the similarity a conditional probability?

As explained in the original article:

“The similarity of datapoint \mathbf{x}_j to datapoint \mathbf{x}_i is the conditional probability, $p_{j|i}$, that \mathbf{x}_i would pick \mathbf{x}_j as its neighbor if neighbors were picking in proportion to their probability density under a Gaussian centered at \mathbf{x}_i .”


t-Distributed Stochastic Neighborhood Embedding (t-SNE)

Given a set of high dimensional data points $\mathbf{x}_1, \dots, \mathbf{x}_n$ t-SNE computes the conditional probabilities, $p_{j|i}$, that are proportional to the similarity of data points \mathbf{x}_i and \mathbf{x}_j

For $i \neq j$, define

$$p_{j|i} = \frac{\exp(-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|\mathbf{x}_i - \mathbf{x}_k\|^2 / 2\sigma_i^2)}$$

and set $p_{i|i} = 0$. Note that $\sum_j p_{j|i} = 1$ for all i .



So how the heck do we pick σ ? σ effectively determines the number of nearest neighbors that any given point “feels”.

In practice, this value of sigma is determined in a fairly complicated and mathematically involved fashion...

One defines the target perplexity “**k**” which is computed using:

$$H = - \sum_j p_{j|i} \log_2 p_{j|i} = \log_2 k,$$

One then finds the optimal σ by finding the value that satisfies this prespecified equation. **High k -> high σ .**

t-Distributed Stochastic Neighborhood Embedding (t-SNE)

Given a set of high dimensional data points x_1, \dots, x_n t-SNE computes the conditional probabilities, $p_{j|i}$, that are proportional to the similarity of data points x_i and x_j

Now – we find a lower dimensional space to which we can map the higher dimensional space that preserves these conditional probabilities as best as possible. We refer to the conditional probabilities in the low dimensional space as $q_{j|i}$.

$$q_{j|i} = \frac{e^{-\|y_i - y_j\|^2}}{\sum_{k \neq i} e^{-\|y_i - y_k\|^2}}.$$

How do we actually do this?

We can measure the mismatch between the high dimensional probabilities and the low dimensional probabilities using the Kullback-Leibler divergence as a loss function for each data point.

The Essence of the Algorithm:

$$C = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}.$$

- Randomly place the low dimensional points and then minimize the cost function using gradient descent to find improved points. This is messy to do in practice but simple conceptually.

$$\frac{\delta C}{\delta y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)(1 + \|y_i - y_j\|^2)^{-1}$$

t-Distributed Stochastic Neighborhood Embedding (t-SNE)

This isn't Exactly t-SNE. To Get to Real t-SNE You Modify a Couple of Things:

1. Symmetrize the conditional probabilities (Makes KL-divergence calculation faster)

$$p_{ij} = \frac{p_{j|i} + p_{i|j}}{2N} \quad \leftarrow \text{Total \# of high-dimensional data points.}$$

and note that $p_{ij} = p_{ji}$, $p_{ii} = 0$, and $\sum_{i,j} p_{ij} = 1$.

2. Use a Student's t-distribution instead of a Gaussian for the low-dimensional probability distribution (Helps prevent crowding in the projection).

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$


And then you do the same gradient of the KL divergence, etc.

t-Distributed Stochastic Neighborhood Embedding (t-SNE)

A few considerations/guidelines

Check out <https://distill.pub/2016/misread-tsne/>.

perplexity should be less than the number of points



convergence is something to monitor



don't over-interpret distances/shapes

Another Useful t-SNE Tip:

- Try initializing your low dimensional representation with a PCA-derived projection instead of random points. Usually this helps, and scikit-learn has a flag that will do this for you.

Uniform Manifold Approximation and Projection (UMAP)

The original paper is a pain to understand and read through, but it suffices to say that people like it because it has a more rigorous mathematical justification than t-SNE. In terms of function, it has many similarities to t-SNE.

**Gaussian similarities
between high dimensional
data points**

$$v_{j|i} = e^{-\frac{d(x_i, x_j) - \rho_i}{\sigma_i}}$$

**Symmetrization of
similarities**

$$v_{ij} = (v_{j|i} + v_{i|j}) - v_{j|i}v_{i|j}$$

**T-distribution similarities between
low dimensional data points**

$$w_{ij} = (1 + a\|y_i - y_j\|^{2b})^{-1}$$

A more nuanced definition of “perplexity”

$$\sum_{j=1}^k \exp\left(-\frac{\max(0, d(x_i, x_j) - \rho_i)}{\sigma_i}\right) = \log_2(k)$$

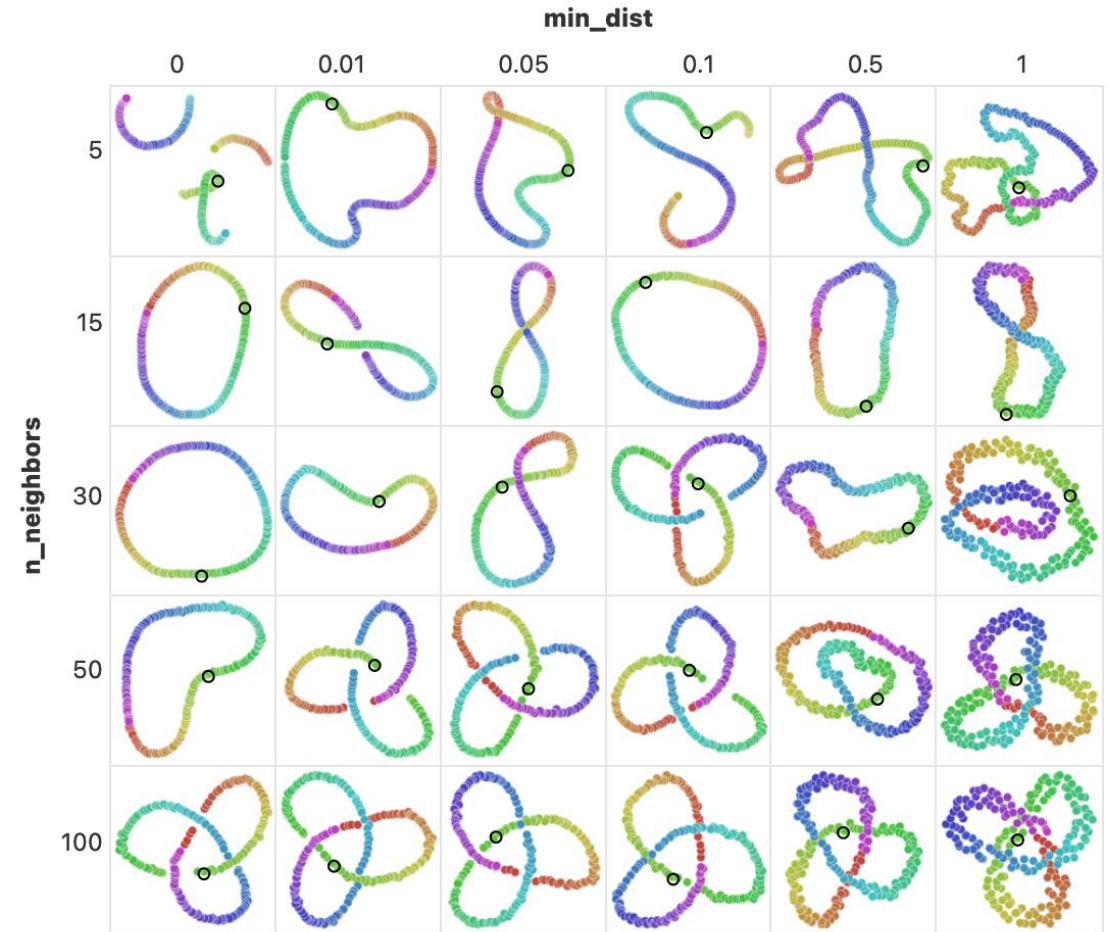

Other Points of Interest About UMAP

- Uses Cross-Entropy instead of KL Divergence.
- Uses Stochastic Gradient Descent instead of plain gradient descent
- Initializes the low dimensional space with a spectral method instead of random or PCA.

<https://umap-learn.readthedocs.io/en/latest/>

Uniform Manifold Approximation and Projection (UMAP)

Hyperparameters are again important to consider in how they dictate the resulting manifold;
in UMAP there are effectively two to control



UMAP vs. t-SNE

If you would like to better understand UMAP I would recommend:

<https://pair-code.github.io/understanding-umap/>

The biggest differences between UMAP and t-SNE are:

- UMAP has a firmer mathematical foundation than t-SNE, and thus many folks like it simply for this reason.
- UMAP uses a few clever tricks to improve numerical stability. It is also faster!!!
- UMAP is often better at preserving global structure than t-SNE - the inter-cluster relationships are potentially more meaningful than t-SNE.
- But both are very widely used!