


and why they may not be machine learning(?)

# Gaussian Processes



Useful references:

*Gaussian Processes for Machine Learning* by Rasmussen and Williams @ [www.GaussianProcess.org/gpml](http://www.GaussianProcess.org/gpml)

# What are Gaussian Processes?

A **Gaussian process (GP)** is a probability distribution over possible functions; any combination of these functions jointly Gaussian distributed

## Key Points:

- GPs can be used as models for both regression and classification
- Since the model is a probability distribution (over a function space that is Gaussian distributed), we can calculate means AND variances
- The **mean function** is calculated from the posterior (recall Baye's) distribution of possible functions → regression predictions
- The **variances** provide native uncertainty measures on these predictions, which are typically absent in other ML models (e.g., ANNs)
- The model, as a posterior, can be updated based on new observations


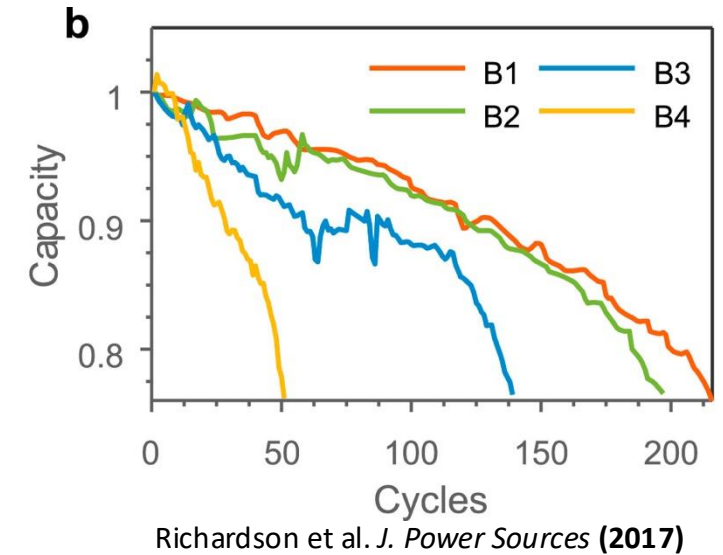

## In “math” speak:

$$f(x) \sim \mathcal{N}(\mu(x), \mathbf{K}(x, x'))$$

covariance/Kernel  
function that dictates  
model “smoothness”

# Why should we use them?

- Our observations may be “**noisy**,” and it is important to capture that behavior
- They allow you to incorporate **prior knowledge**
- They provide a reasonable framework to **regulate model complexity** and more reliably **indicate model uncertainty**




# Some basic intuition underlying Gaussian Processes


Suppose we have a random variable  $X_1$  with distribution

$$P_{X_1}(x_1) \sim \mathcal{N}(\mu, \sigma^2)$$

We can generate samples from this distribution



and then project them onto a new axis




# Some basic intuition underlying Gaussian Processes


Suppose we have a random variable  $X_1$  with distribution

$$P_{X_1}(x_1) \sim \mathcal{N}(\mu, \sigma^2)$$

We can generate samples from this distribution



and then project them onto a new axis



*We can do this for  
another set of samples to  
obtain two independent  
Gaussian vectors*

*Then by connecting random  
points between the two  
vectors, we would get a set  
of linear functions*




# Some basic intuition underlying Gaussian Processes


Suppose we have a random variable  $X_1$  with distribution

$$P_{X_1}(x_1) \sim \mathcal{N}(\mu, \sigma^2)$$

We can generate samples from this distribution



and then project them onto a new axis



*using more and more  
Gaussian vectors, we could  
represent more complex  
functions (showing only ten),  
but the resulting functions are  
inherently noisy!*


**Why?**

# Some basic intuition underlying Gaussian Processes


Suppose we have a random variable  $X_1$  with distribution

$$P_{X_1}(x_1) \sim \mathcal{N}(\mu, \sigma^2)$$

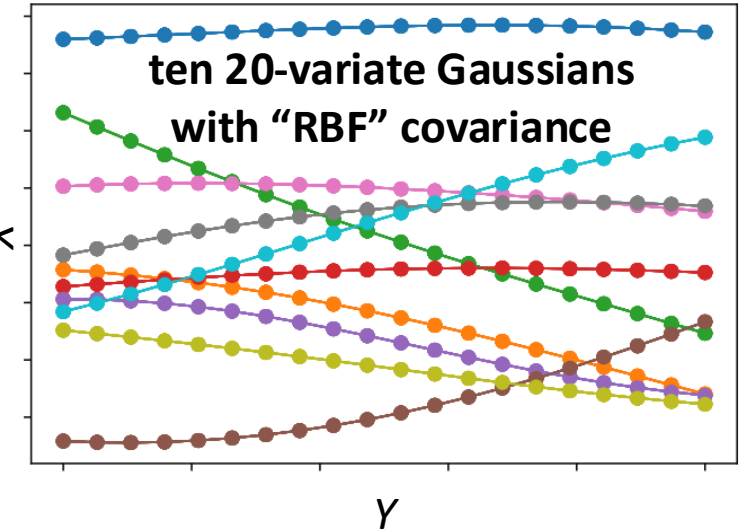
We can generate samples from this distribution



and then project them onto a new axis



*using multivariate Gaussians with true covariance enables more smoothly varying functions*



e.g.,


$$K(x_i, x_j) = \exp \left[ -\frac{(x_i - x_j)^2}{2} \right]$$

*using more and more Gaussian vectors, we could represent more complex functions (showing only ten), but the resulting functions are inherently noisy!*


**...because the Gaussian vectors are independent**




**200 200-variate Gaussians with "RBF" covariance**



# Some basic intuition underlying Gaussian Processes



- If the dimensionality of the ***multivariate Gaussians*** (MVG) becomes infinite  $\rightarrow$  a continuous function
- With an infinite number of such functions, we could make predictions at any point
- Functions as MVGs is our initial assumption (***prior***)
- The functions shown provide expected outputs as a function of inputs without having observed any data



- As data is collected, we can downselect from infinite functions to only the functions that describe the data/observations  $\rightarrow$  ***posterior*** (kind of like wavefunctions in quantum mechanics, if that helps)
- As we add more data, the current posterior can be used as the prior to obtain a new posterior



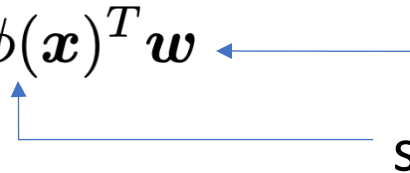
# Theoretical Development of GPs for Regression

A **Gaussian process (GP)** is a probability distribution over possible functions; any combination of these functions jointly Gaussian distributed

$f(\mathbf{x}) \sim \mathcal{GP}(\mu(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$  this is a stochastic function, i.e., a collection of random variables

$\mu(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$  its distribution is characterized by its mean and covariance function

$$k(\mathbf{x}, \mathbf{x}') = \text{Cov}[f(\mathbf{x}), f(\mathbf{x}')] = \mathbb{E}[(f(\mathbf{x}) - \mu(\mathbf{x}))(f(\mathbf{x}') - \mu(\mathbf{x}'))]$$

A simple example  $f(\mathbf{x}) = \phi(\mathbf{x})^T \mathbf{w}$   weighting coefficients BUT  $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \Sigma_p)$   
some vector field, e.g.,  $\phi(\mathbf{x})^T = (\sum_i x_i, \sum_i x_i^2, \dots, \sum_i x_i^m)$

for such a process, it is easy to show that  $\mu(\mathbf{x}) = 0; k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \Sigma_p \phi(\mathbf{x}')$

that is just to say that  $f(\mathbf{x})$  and  $f(\mathbf{x}')$  are indeed jointly Gaussian with the above mean and covariance  $\rightarrow$  any number of input points will be jointly Gaussian

# Theoretical Development of GPs for Regression

Let's suppose we have a dataset now  $\mathcal{D} := \{(\mathbf{x}_i, f_i = f(\mathbf{x}_i)) \text{ for } i = 1, \dots, n\}$

$\mathbf{f}, \mathbf{x}$   
*training*

$\mathbf{f}_*, \mathbf{x}_*$   
*test*

These two sets should be jointly distributed according to our Gaussian prior as

$$\begin{bmatrix} \mathbf{f} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N} \left( \mathbf{0}, \begin{bmatrix} k(\mathbf{x}, \mathbf{x}) & k(\mathbf{x}, \mathbf{x}_*) \\ k(\mathbf{x}_*, \mathbf{x}) & k(\mathbf{x}_*, \mathbf{x}_*) \end{bmatrix} \right) \quad \text{prior distribution}$$

*Now, to get the posterior distribution, we must restrict our prior distribution to only the set of functions that “agree” with our observations/training data (i.e., we should make our prior distribution conditional on our observations)*

# Theoretical Development of GPs for Regression

Let's suppose we have a dataset now  $\mathcal{D} := \{(\mathbf{x}_i, f_i = f(\mathbf{x}_i)) \text{ for } i = 1, \dots, n\}$

$\mathbf{f}, \mathbf{x}$   
*training*

$\mathbf{f}_*, \mathbf{x}_*$   
*test*

These two sets should be jointly distributed according to our Gaussian prior as

$$\begin{bmatrix} \mathbf{f} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N} \left( \mathbf{0}, \begin{bmatrix} k(\mathbf{x}, \mathbf{x}) & k(\mathbf{x}, \mathbf{x}_*) \\ k(\mathbf{x}_*, \mathbf{x}) & k(\mathbf{x}_*, \mathbf{x}_*) \end{bmatrix} \right) \quad \text{prior distribution}$$

*from math review*

$$p(\mathbf{x}, \mathbf{y}) = \mathcal{N} \left( \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}, \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix} \right)$$

$$p(\mathbf{x} | \mathbf{y}) = \mathcal{N}(\mu_{x|y}, \Sigma_{x|y})$$

$$\mu_{x|y} = \mu_x + \Sigma_{xy} \Sigma_{yy}^{-1} (\mathbf{y} - \mu_y)$$

$$\Sigma_{x|y} = \Sigma_{xx} - \Sigma_{xy} \Sigma_{yy}^{-1} \Sigma_{yx}$$

$$\mathbf{f}_* | \mathbf{f}, \mathbf{x}, \mathbf{x}_* \sim \mathcal{N} \left( \begin{aligned} &k(\mathbf{x}_*, \mathbf{x}) k(\mathbf{x}, \mathbf{x})^{-1} \mathbf{f}, \\ &k(\mathbf{x}_*, \mathbf{x}_*) - k(\mathbf{x}_*, \mathbf{x}) k(\mathbf{x}, \mathbf{x})^{-1} k(\mathbf{x}, \mathbf{x}_*) \end{aligned} \right)$$

*easy to just sample function values from the posterior by computing the mean and covariance in the above!*

# Theoretical Development of GPs for Regression

Let's suppose we have a dataset now  $\mathcal{D} := \{(\mathbf{x}_i, f_i = f(\mathbf{x}_i)) \text{ for } i = 1, \dots, n\}$

$\mathbf{f}, \mathbf{x}$   
*training*

$\mathbf{f}_*, \mathbf{x}_*$   
*test*

In the prior example, our observations were treated as Noise-free; otherwise...

$$y = f(\mathbf{x}) + \varepsilon \quad \varepsilon \sim \mathcal{N}(0, \sigma_n^2) \quad \text{cov}(\mathbf{y}, \mathbf{y}) = k(\mathbf{x}, \mathbf{x}) + \sigma_n^2 \mathbf{I}$$

This would slightly modify our joint distribution as

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N} \left( \mathbf{0}, \begin{bmatrix} k(\mathbf{x}, \mathbf{x}) + \sigma_n^2 \mathbf{I} & k(\mathbf{x}, \mathbf{x}_*) \\ k(\mathbf{x}_*, \mathbf{x}) & k(\mathbf{x}_*, \mathbf{x}_*) \end{bmatrix} \right)$$

Then, the posterior is obtained by again conditioning on our observations

*these are governing equations*

$$\begin{aligned} \mathbf{f}_* | \mathbf{y}, \mathbf{x}, \mathbf{x}_* &\sim \mathcal{N}(\bar{\mathbf{f}}_*, \text{cov}(\mathbf{f}_*, \mathbf{f}_*)) \\ \bar{\mathbf{f}}_* &:= \mathbb{E}[\mathbf{f}_* | \mathbf{y}, \mathbf{x}, \mathbf{x}_*] = k(\mathbf{x}_*, \mathbf{x}) [k(\mathbf{x}, \mathbf{x}) + \sigma_n^2 \mathbf{I}]^{-1} \mathbf{y} \\ \mathbb{V}[\mathbf{f}_*] &= \text{cov}(\mathbf{f}_*, \mathbf{f}_*) = k(\mathbf{x}_*, \mathbf{x}_*) - k(\mathbf{x}_*, \mathbf{x}) [k(\mathbf{x}, \mathbf{x}) + \sigma_n^2 \mathbf{I}]^{-1} k(\mathbf{x}, \mathbf{x}_*) \end{aligned}$$



# Log Marginal Likelihood

*“marginal likelihood” or “evidence”*

$$p(\mathbf{y}|\mathbf{x}) = \int p(\mathbf{y}|\mathbf{f}, \mathbf{x})p(\mathbf{f}|\mathbf{x})d\mathbf{f}$$

*→ how likely would we observe the data that we have given the data originates from  $\mathbf{f}(\mathbf{x})$*

*Assuming a Gaussian process, we can show that*

$$\log p(\mathbf{y}|\mathbf{x}) = -\frac{1}{2}\mathbf{y}^T (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} - \frac{1}{2} \log |\mathbf{K} + \sigma_n^2 \mathbf{I}| - \frac{n}{2} \log 2\pi$$

*“log marginal likelihood”*

- All the terms necessary to evaluate this are also used to obtain the mean function
- Useful metric for evaluating different models
- Hyperparameter optimization based on minimizing log likelihood

# Basic Algorithm Outline

1. Define data/models:  $\mathbf{x}, \mathbf{y}, k, \sigma_n^2, \mathbf{x}_*$
2. Perform Cholesky decomposition  $\mathbf{L} := \text{cholesky}(k(\mathbf{x}, \mathbf{x}) + \sigma_n^2 \mathbf{I})$
3. Compute predictive mean function  $\boldsymbol{\alpha} := \mathbf{L}^T \setminus (\mathbf{L} \setminus \mathbf{y})$   
 $\bar{f}_* := k(\mathbf{x}, \mathbf{x}_*)^T \boldsymbol{\alpha}$
4. Compute predictive variance  $\mathbf{v} := \mathbf{L} \setminus k(\mathbf{x}, \mathbf{x}_*)$   
 $\mathbb{V}[f_*] = k(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{v}^T \mathbf{v}$
5. Compute log marginal likelihood  
$$\log p(\mathbf{y}|\mathbf{x}) = -\frac{1}{2} \mathbf{y}^T \boldsymbol{\alpha} - \sum_i \log L_{ii} - \frac{n}{2} \log 2\pi$$
6. Return mean, variance, log marginal likelihood

# Basic Algorithm Outline

1. Define data/models:  $\mathbf{x}, \mathbf{y}, k, \sigma_n^2, \mathbf{x}_*$
2. Perform Cholesky decomposition  $\mathbf{L} := \text{cholesky}(k(\mathbf{x}, \mathbf{x}) + \sigma_n^2 \mathbf{I})$
3. Compute the inverse of the lower triangular matrix  $\mathbf{Q} := \mathbf{L}^T \setminus (\mathbf{L} \setminus \mathbf{y})$

***Go to Jupyter Notebook***

5. Compute log marginal likelihood

$$\log p(\mathbf{y}|\mathbf{x}) = -\frac{1}{2}\mathbf{y}^T \boldsymbol{\alpha} - \sum_i \log L_{ii} - \frac{n}{2} \log 2\pi$$

6. Return mean, variance, log marginal likelihood

# Gaussian Distribution

**Gaussian distribution** is the most well-studied distribution in science, engineering, and ML

- *ubiquity (often stemming from **central limit theorem**)*
- *computationally convenient*

univariate

$$p(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

multivariate

$$p(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{m/2} |\boldsymbol{\Sigma}|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

$$p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) \quad \text{for short}$$

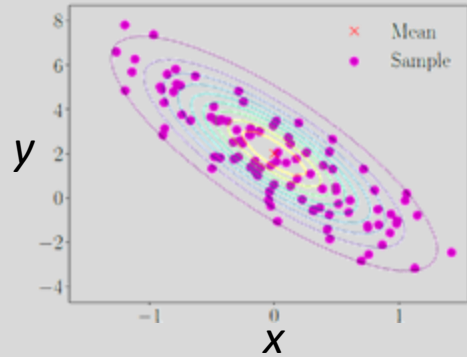
standard normal distribution  $\rightarrow \mathbf{0}$  mean and  $\mathbf{I}_n$  as the covariance matrix

Basic tenant of Gaussian distributions: *doing things with Gaussians results in Gaussians*



# Gaussian Distribution


Consider the joint distribution over  $X$  and  $Y$ :




$$p(\mathbf{x}, \mathbf{y}) = \mathcal{N}\left(\begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}, \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}\right)$$

- *Marginals over Gaussians are Gaussians*

$$p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{y}) d\mathbf{y} = \mathcal{N}(\mathbf{x} | \mu_x, \Sigma_{xx})$$



- *Conditionals of Gaussians are Gaussians*



$$p(\mathbf{x} | \mathbf{y}) = \mathcal{N}(\mu_{x|y}, \Sigma_{x|y})$$

$$\mu_{x|y} = \mu_x + \Sigma_{xy} \Sigma_{yy}^{-1} (\mathbf{y} - \mu_y)$$

$$\Sigma_{x|y} = \Sigma_{xx} - \Sigma_{xy} \Sigma_{yy}^{-1} \Sigma_{yx}$$

- *The product of Gaussians is Gaussian*

$$\mathcal{N}(\mathbf{x} | \mathbf{a}, \mathbf{A}) \mathcal{N}(\mathbf{x} | \mathbf{b}, \mathbf{B}) \rightarrow c \mathcal{N}(\mathbf{x} | \mathbf{c}, \mathbf{C})$$

$$\mathbf{C} = (\mathbf{A}^{-1} + \mathbf{B}^{-1})^{-1}$$

$$\mathbf{c} = \mathbf{C}(\mathbf{A}^{-1} \mathbf{a} + \mathbf{B}^{-1} \mathbf{b}) \quad c = \mathcal{N}(\mathbf{a} | \mathbf{b}, \mathbf{A} + \mathbf{B})$$

- *The sum over independent Gaussian variables is... Gaussian*

$$p(\mathbf{x}, \mathbf{y}) = p(\mathbf{x}) p(\mathbf{y})$$

$$p(\mathbf{x} + \mathbf{y}) = \mathcal{N}(\mu_x + \mu_y, \Sigma_x + \Sigma_y)$$

- *Any linear transformation of Gaussians is Gaussian*  $p(a\mathbf{x} + b\mathbf{y}) = \mathcal{N}(a\mu_x + b\mu_y, a^2\Sigma_x + b^2\Sigma_y)$